Intel® FPGA SDK for OpenCL

Programming Guide

Subscribe

Send Feedback

Contents

1 Intel® FPGA SDK for OpenCL™ Programming Guide

The *Intel® FPGA SDK for OpenCL™ Programming Guide* provides descriptions, recommendations and usage information on the Intel Software Development Kit (SDK) for OpenCL compiler and tools. The Intel FPGA SDK for OpenCL[1] is an OpenCL[2]-based heterogeneous parallel programming environment for Intel FPGA products.

1.1 Intel FPGA SDK for OpenCL Programming Guide Prerequisites

The *Intel FPGA SDK for OpenCL Programming Guide* assumes that you are knowledgeable in OpenCL concepts and application programming interfaces (APIs). It also assumes that you have experience creating OpenCL applications and are familiar with the OpenCL Specification version 1.0.

Before using the Intel FPGA SDK for OpenCL or the Intel FPGA Runtime Environment (RTE) for OpenCL to program your device, familiarize yourself with the respective getting started guides. This document assumes that you have performed the following tasks:

- For developing and deploying OpenCL kernels, download the tar file and run the installers to install the SDK, the Quartus® Prime software, and device support.

- For deployment of OpenCL kernels, download and install the RTE.

- If you want to use the SDK or the RTE to program a Cyclone® V SoC Development Kit, you also have to download and install the SoC Embedded Design Suite (EDS).

- Install and set up your FPGA board.

- Program your device with the device-compatible version of the hello_world example OpenCL application

If you have not performed the tasks described above, refer to the SDK's getting starting guides for more information.

Prior to creating an OpenCL design and programming your FPGA board, review the Intel FPGA SDK for OpenCL allocation limits.

Related Links

- [Intel FPGA SDK for OpenCL Allocation Limits](#) on page 157

- [OpenCL References Pages](#)

---

1 The Intel FPGA SDK for OpenCL is based on a published Khronos Specification, and has passed the Khronos Conformance Testing Process. Current conformance status can be found at www.khronos.org/conformance.

2 OpenCL and the OpenCL logo are trademarks of Apple Inc. used by permission of the Khronos Group™.

- OpenCL Specification version 1.0
- Intel FPGA SDK for OpenCL Getting Started Guide
- Intel FPGA RTE for OpenCL Getting Started Guide
- Intel FPGA SDK for OpenCL Cyclone V SoC Getting Started Guide

1.2 Intel FPGA SDK for OpenCL FPGA Programming Flow

The Intel FPGA SDK for OpenCL programs an FPGA with an OpenCL application in a two-step process. The Intel FPGA SDK for OpenCL Offline Compiler first compiles your OpenCL kernels. The host-side C compiler compiles your host application and then links the compiled OpenCL kernels to it.

Figure 1.   **Schematic Diagram of the Intel FPGA SDK for OpenCL Programming Model**
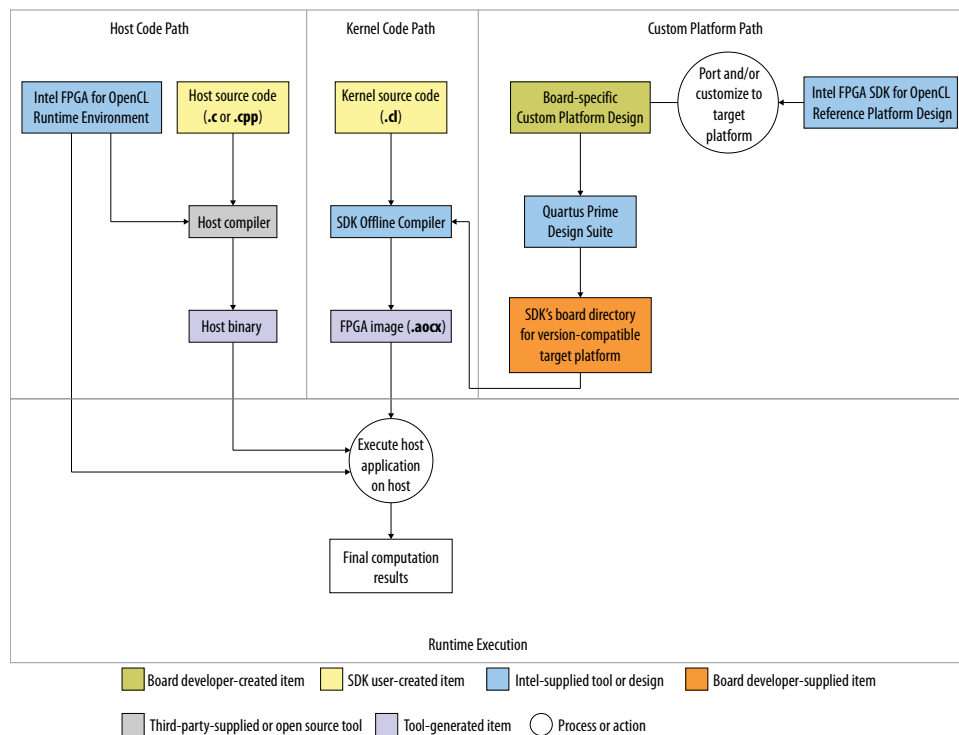
Three main parts in the SDK's programming model:

- The host application and the host compiler
- The OpenCL kernel and the offline compiler
- The Custom Platform

The Custom Platform provides the board design. The offline compiler targets the board design when compiling the OpenCL kernel to generate the hardware image. The host then runs the host application to execute the hardware image onto the FPGA.

1.3 Intel FPGA SDK for OpenCL Offline Compiler Kernel Compilation Flows

The Intel FPGA SDK for OpenCL Offline Compiler can create your FPGA hardware configuration file in a one-step or a multistep process. The complexity of your kernel dictates the compilation option you implement.

Figure 2.   The Intel FPGA SDK for OpenCL FPGA Programming Flow
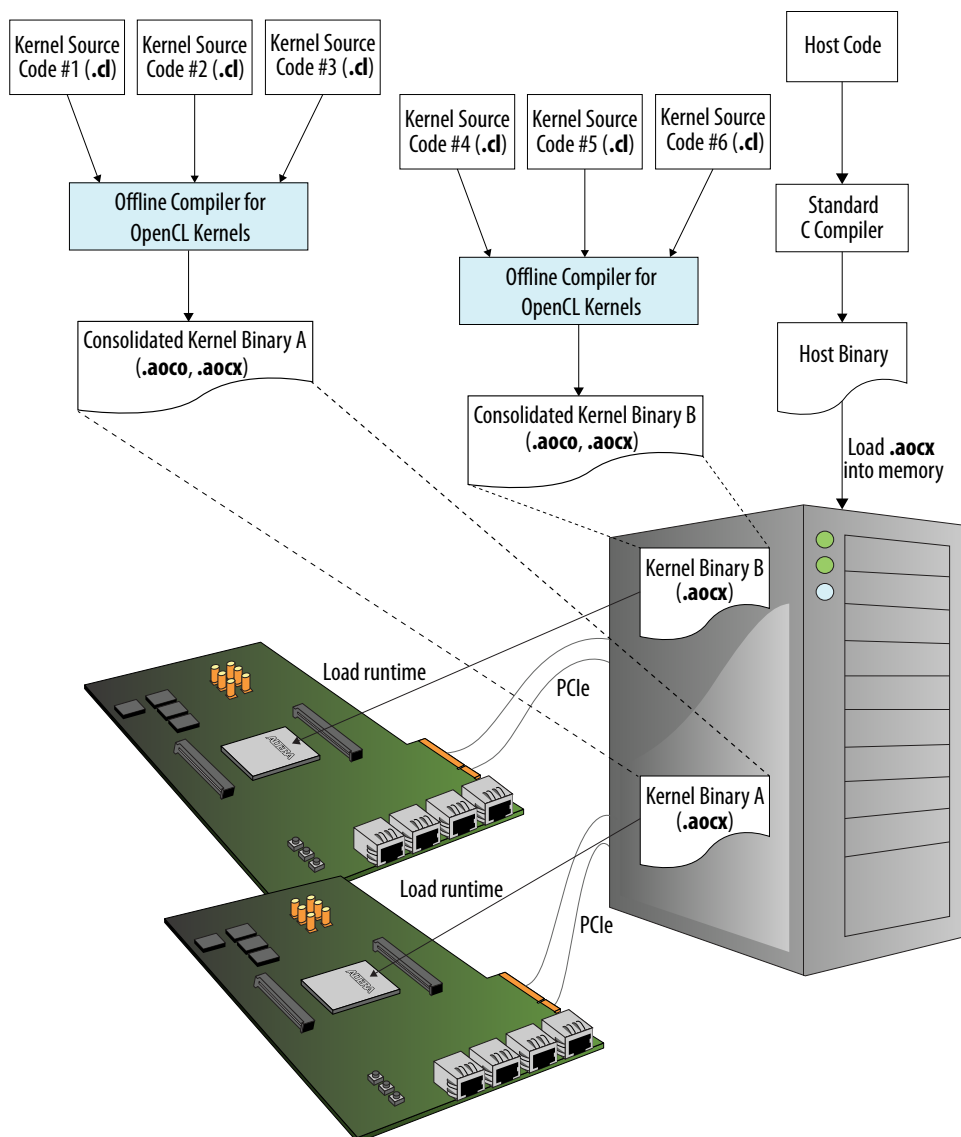
An OpenCL kernel source file (`.cl`) contains your OpenCL source code. The offline compiler groups one or more kernels into a temporary file and then compiles this file to generate the following files and folders:

- A `.aoco` object file is an intermediate object file that contains information for later stages of the compilation.

- A `.aocx` executable file is the hardware configuration file and contains information necessary at runtime.

- The `<your_kernel_filename>` folder or subdirectory, which contains data necessary to create the `.aocx` file.

The offline compiler creates the `.aocx` file from the contents of the `<your_kernel_filename>` folder or subdirectory. It also incorporates information from the `.aoco` file into the `.aocx` file during hardware compilation. The `.aocx` file contains data that the host application uses to create program objects for the target FPGA. The host application loads these program objects into memory. The host runtime then calls these program objects from memory and programs the target FPGA as required.

1.3.1 One-Step Compilation for Simple Kernels

By default, the Intel FPGA SDK for OpenCL Offline Compiler compiles your OpenCL kernel and creates the hardware configuration file in a single step. Choose this compilation option only if your OpenCL application requires minimal optimizations.

The following figure illustrates the OpenCL kernel design flow that has a single compilation step.

Figure 3.   One-Step OpenCL Kernel Compilation Flow
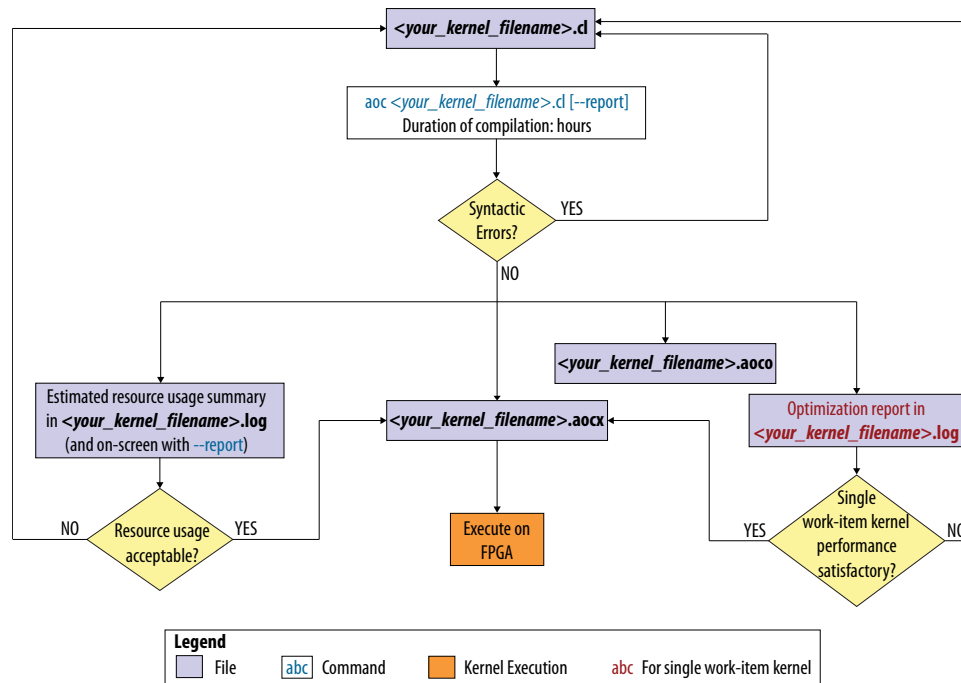
A successful compilation results in the following files and reports:

- A `.aoco` file

- A `.aocx` file

- In the *`<your_kernel_filename>/<your_kernel_filename>`*`.log` file, the estimated resource usage summary provides a preliminary assessment of area usage. If you have a single work-item kernel, the optimization report identifies performance bottlenecks.

Attention: It is very time consuming to iterate on your design using the one-step compilation flow. For each iteration, you must perform a full compilation, which takes hours. Then you must execute the kernel on the FPGA before you can assess its performance.

Related Links

Compiling Your Kernel to Create Hardware Configuration File on page 95
> You can compile an OpenCL kernel and create the hardware configuration file (that is, the `.aocx` file) in a single step.

1.3.2 Multistep Intel FPGA SDK for OpenCL Design Flow

Choose the multistep Intel FPGA SDK for OpenCL design flow if you want to iterate on your OpenCL kernel design to implement performance-improving optimizations .

The figure below outlines the stages in the SDK's design flow. The steps in the design flow serve as checkpoints for identifying functional errors and performance bottlenecks. They allow you to modify your OpenCL kernel code without performing a full compilation after each iteration.

Figure 4.   The Multistep Intel FPGA SDK for OpenCL Design Flow

<your_kernel_filename>.cl

Intermediate Compilation
aoc -c *<your_kernel_filename>* .cl [--report]
Duration of compilation: minutes

Syntactic errors? — YES / NO

<your_kernel_filename>.aoco

Estimated resource usage summary in *<your_kernel_filename>*.log (and on-screen with --report)

<your_kernel_filename>.aoco

Optimization Repot in *<your_kernel_filename>*.log

Resource Usage Acceptable? — NO / YES

Emulation
aoc -march=emulator *<your_kernel_filename>* .cl
Duration of compilation: seconds

<your_kernel_filename>.aocx

Execute on emulation device

Emulation successful? — NO / YES

Single work-item kernel? — YES / NO

Kernel performance satisfactory? — YES / NO

Review HTML Report
<your_kernel_filename>/reports/report.html

Estimated kernel performance data acceptable? — NO / YES

Profiling
aoc --profile *<your_kernel_filename>*.cl
Duration of compilation: hours

<your_kernel_filename>.aocx

<your_kernel_filename>.source

Execute kernel on FPGA

profile.mon

aocl report *<your_kernel_filename>*.aocx profile.mon [*<your_kernel_filename>*.source]

Profiler GUI

Kernel performance satisfactory? — NO / YES

Full Deployment
aoc *<your_kernel_filename>*.cl
Duration of compilation: hours

<your_kernel_filename>.aocx

Execute kernel on FPGA

Legend
- File
- abc Command
- Kernel Execution
- GUI
- abc Single work-item kernel step

The SDK's design flow includes the following steps:

1. Intermediate compilation

   The intermediate compilation step checks for syntatic errors. It then generates a `.aoco` file without building the hardware configuration file. The estimated resource usage summary in the `<your_kernel_filename>/<your_kernel_filename>.log` file can provide insight into the type of kernel optimizations you can perform. .

2. Emulation

   Assess the functionality of your OpenCL kernel by executing it on one or multiple emulation devices on an x86-64 host. For Linux systems, the Emulator offers symbolic debug support. Symbolic debug allows you to locate the origins of functional errors in your kernel code.

3. Review HTML Report

   Review the `<your_kernel_filename>/reports/report.html` file of your OpenCL application to determine whether the estimated kernel performace data is acceptable. The HTML report also provides suggestions on how you can modify your kernel to increase performance.

4. Profiling

   Instruct the Intel FPGA SDK for OpenCL Offline Compiler to instrument performance counters in the Verilog code in the `.aocx` file. During execution, the performance counters collect performance information which you can then review in the Profiler GUI.

5. Full deployment

   If you are satisfied with the performance of your OpenCL kernel throughout the design flow, perform a full compilation. You can then execute the `.aocx` file on the FPGA.

Related Links

- Reviewing Your Kernel's report.html File on page 112
  To launch the HTML report, open the `report.html` file in the `<your_kernel_filename>/reports` directory.

- Compiling Your OpenCL Kernel on page 95
  The Intel FPGA SDK for OpenCL offers a list of compiler options that allows you to customize the kernel compilation process.

- Emulating and Debugging Your OpenCL Kernel on page 105
  The Intel FPGA SDK for OpenCL Emulator assesses the functionality of your kernel.

- Profiling Your OpenCL Kernel on page 112
  The Intel FPGA SDK for OpenCL Profiler measures and reports performance data collected during OpenCL kernel execution on the FPGA.

1.4 Obtaining General Information on Software, Compiler, and Custom Platform

The Intel FPGA SDK for OpenCL includes two sets of command options: the SDK utility commands (`aocl <command_option>`) and the offline compiler commands (`aoc <command_option>`). Each set of commands includes options you can invoke to obtain general information on the software, the compiler, and the Custom Platform.

Displaying the Software Version (version) on page 12
    To display the version of the Intel FPGA SDK for OpenCL, invoke the `version` utility command.

Displaying the Compiler Version (--version) on page 12

To display the version of the Intel FPGA SDK for OpenCL Offline Compiler, invoke the `--version` compiler command.

To display information on the Intel FPGA SDK for OpenCL utility command options, invoke the `help` utility command.

To display information on the Intel FPGA SDK for OpenCL Offline Compiler command options, invoke the compiler command without an argument, or invoke the compiler command with the `--help` or `-h` command option.

To list the FPGA boards available in your Custom Platform, include the `--list-boards` option in the `aoc` command.

To display the Intel FPGA SDK for OpenCL Offline Compiler's input arguments and the environment for a compiled OpenCL design, invoke the `env` utility command.

1.4.1 Displaying the Software Version (version)

To display the version of the Intel FPGA SDK for OpenCL, invoke the `version` utility command.

- At the command prompt, invoke the `aocl version` command.
  Example output:

  `aocl <version>.<build> (Intel(R) FPGA SDK for OpenCL(TM), Version <version> Build <build>, Copyright (C) <year> Intel Corporation)`

1.4.2 Displaying the Compiler Version (--version)

To display the version of the Intel FPGA SDK for OpenCL Offline Compiler, invoke the `--version` compiler command.

- At a command prompt, invoke the `aoc --version` command.
  Example output:

  ```
  Intel(R) FPGA SDK for OpenCL(TM), 64-Bit Offline Compiler
  Version <version> Build <build>
  Copyright (C) <year> Intel Corporation
  ```

1.4.3 Listing the Intel FPGA SDK for OpenCL Utility Command Options (help)

To display information on the Intel FPGA SDK for OpenCL utility command options, invoke the `help` utility command.

- At a command prompt, invoke the `aocl help` command.
  The SDK categorizes the utility command options based on their functions. It also provides a description for each option.

1.4.3.1 Displaying Information on an Intel FPGA SDK for OpenCL Utility Command Option (help <command_option>)

To display information on a specific Intel FPGA SDK for OpenCL utility command option, include the command option as an argument of the `help` utility command.

- At a command prompt, invoke the `aocl help <command_option>` command.

   For example, to obtain more information on the `install` utility command option, invoke the `aocl help install` command.

   Example output:

   ```
   aocl install - Installs a board onto your host system.

   Usage: aocl install

   Description:
   This command installs a board's drivers and other necessary software for the
   host operating system to communicate with the board.
   For example this might install PCIe drivers.
   ```

1.4.4 Listing the Intel FPGA SDK for OpenCL Offline Compiler Command Options (no argument, --help, or -h)

To display information on the Intel FPGA SDK for OpenCL Offline Compiler command options, invoke the compiler command without an argument, or invoke the compiler command with the `--help` or `-h` command option.

- At a command prompt, invoke one of the following commands:
  - `aoc`
  - `aoc --help`
  - `aoc -h`

   The SDK categorizes the offline compiler command options based on their functions. It also provides a description for each option.

1.4.5 Listing the Available FPGA Boards in Your Custom Platform (--list-boards)

To list the FPGA boards available in your Custom Platform, include the `--list-boards` option in the `aoc` command.

To view the list of available boards in your Custom Platform, you must first set the environment variable *AOCL_BOARD_PACKAGE_ROOT* to point to the location of your Custom Platform.

- At a command prompt, invoke the `aoc --list-boards` command.

  The Intel FPGA SDK for OpenCL Offline Compiler generates an output that resembles the following:

  ```
  Board list:
      <board_name_1>
      <board_name_2>
  ...
  ```

  Where *<board_name_N>* is the board name you use in your `aoc` command to target a specific FPGA board.

1.4.6 Displaying the Compilation Environment of an OpenCL Binary (env)

To display the Intel FPGA SDK for OpenCL Offline Compiler's input arguments and the environment for a compiled OpenCL design, invoke the `env` utility command.

- At the command prompt, invoke the `aocl env <object_file_name>` or the `aocl env <executable_file_name>` command,

  where *<object_file_name>* is the name of the `.aoco` file of your OpenCL kernel, and the *<executable_file_name>* is the name of the `.aocx` file of your kernel.

  Output for the example command `aocl env vector_add.aocx`:

  ```
  INPUT_ARGS=-march=emulator -v device/vector_add.cl -o bin/vector_add.aocx
  BUILD_NUMBER=90
  ACL_VERSION=16.1.0
  OPERATING_SYSTEM=linux
  PLATFORM_TYPE=s5_net
  ```

1.5 Managing an FPGA Board

The Intel FPGA SDK for OpenCL includes utility commands you can invoke to install, uninstall, diagnose, and program your FPGA board.

You can install multiple Custom Platforms simultaneously on the same system. To use the SDK utilities, such as `aocl diagnose` with multiple Custom Platforms, you must set the *AOCL_BOARD_PACKAGE_ROOT* environment variable to point to the location of the Custom Platform subdirectory of the board on which you wish to run the utility. The Custom Platform subdirectory contains the `board_env.xml` file. To run the SDK utilities on a different Custom Platform, you must update the *AOCL_BOARD_PACKAGE_ROOT* environment variable to point to the location of the Custom Platform subdirectory of that specific board.

In a system with multiple Custom Platforms, ensure that the host program uses the FPGA Client Driver (FCD), formerly Altera Client Driver (ACD), to discover the boards rather than linking to the Custom Platforms' memory-mapped device (MMD) libraries directly. As long as FCD is correctly set up for Custom Platform, FCD finds all the installed boards at runtime.

Installing an FPGA Board (install) on page 15
   To install your board into the host system, invoke the `install` utility command.

Uninstalling the FPGA Board (uninstall) on page 17

To uninstall an FPGA board, invoke the `uninstall` utility command, uninstall the Custom Platform, and unset the relevant environment variables.

Querying the Device Name of Your FPGA Board (diagnose) on page 17
When you query a list of accelerator boards, the OpenCL software produces a list of installed devices on your machine in the order of their device names.

Running a Board Diagnostic Test (diagnose <device_name>) on page 18
To perform a detailed diagnosis on a specific FPGA board, include *<device_name>* as an argument of the `diagnose` utility command.

Programming the FPGA Offline or without a Host (program <device_name>) on page 18
To program an FPGA device offline or without a host, invoke the `program` utility command.

Programming the Flash Memory (flash <device_name>) on page 18
If supported, invoke the `flash` utility command to initialize the FPGA with a specified startup configuration.

Related Links

- Installing an FPGA Board (install) on page 15
  To install your board into the host system, invoke the `install` utility command.

- Linking Your Host Application to the Khronos ICD Loader Library on page 86
  The Intel FPGA SDK for OpenCL supports the OpenCL ICD extension from the Khronos Group.

1.5.1 Installing an FPGA Board (install)

Before creating an OpenCL application for an FPGA board, you must first download and install the Custom Platform from your board vendor. Most Custom Platform installers require administrator privileges. To install your board into the host system, invoke the `install` utility command.

The steps below outline the board installation procedure. Some Custom Platforms require additional installation tasks. Consult your board vendor's documentation for further information on board installation.

Attention: If you are installing the Cyclone V SoC Development Kit for use with the Cyclone V SoC Development Kit Reference Platform (c5soc), refer to *Installing the Cyclone V SoC Development Kit* in the *Intel FPGA SDK for OpenCL Cyclone V SoC Getting Started Guide* for more information.

1. Follow your board vendor's instructions to connect the FPGA board to your system.

2. Download the Custom Platform for your FPGA board from your board vendor's website. To download an Intel FPGA SDK for OpenCL Reference Platform (for example, the Stratix® V Network Reference Platform (s5_net)), refer to the Intel FPGA SDK for OpenCL FPGA Platforms page.

3. Install the Custom Platform in a directory that you own (that is, not a system directory).

4. Set the user environment variable *AOCL_BOARD_PACKAGE_ROOT* to point to the location of the Custom Platform subdirectory containing the `board_env.xml` file.

For example, for s5_net, set *AOCL_BOARD_PACKAGE_ROOT* to point to the
`<path_to_s5_net>`/`s5_net` directory.

5. Set the *QUARTUS_ROOTDIR_OVERRIDE* user environment variable to point to the
correct Quartus Prime software installation directory.

If you have an Arria® 10 device, set *QUARTUS_ROOTDIR_OVERRIDE* to point to
the installation directory of the Quartus Prime Pro Edition software. Otherwise, set
QUARTUS_ROOTDIR_OVERRIDE to point to the installation directory of the
Quartus Prime Standard Edition software.

6. Add the paths to the Custom Platform libraries (for example, the memory-mapped
(MMD) library) to the *PATH* (Windows) or *LD_LIBRARY_PATH* (Linux) environment
variable setting.

For example, if you use s5_net, the Windows *PATH* environment variable setting is
`%AOCL_BOARD_PACKAGE_ROOT%\windows64\bin`. The Linux *LD_LIBRARY_PATH*
setting is `$AOCL_BOARD_PACKAGE_ROOT/linux64/lib`.

The *Intel FPGA SDK for OpenCL Getting Started Guide* contains more information
on the `init_opencl` script. For information on setting user environment variables
and running the `init_opencl` script, refer to the *Setting the Intel FPGA SDK for
OpenCL User Environment Variables* section.

7. *Remember:* You need administrative rights to install a board. To run a Windows
command prompt as an administrator, click **Start ➤ All Programs ➤
Accessories**. Under **Accessories**, right click **Command Prompt**, In
the right-click menu, click **Run as Administrator**.

Invoke the command `aocl install` at a command prompt.

Invoking `aocl install` also installs a board driver that allows communication
between host applications and hardware kernel programs.

8. To query a list of FPGA devices installed in your machine, invoke the `aocl
diagnose` command.
The software generates an output that includes the *<device_name>*, which is an
acl number that ranges from acl0 to acl31.

For more information on querying the *<device_name>* of your accelerator board,
refer to the *Querying the Device Name of Your FPGA Board* section.

9. To verify the successful installation of the FPGA board, invoke the command `aocl
diagnose` *<device_name>* to run any board vendor-recommended diagnostic
test.

Related Links

- [Installing the Cyclone V SoC Development Kit](#)
- [Querying the Device Name of Your FPGA Board (diagnose)](#) on page 17
  When you query a list of accelerator boards, the OpenCL software produces a
  list of installed devices on your machine in the order of their device names.
- [Setting the Intel FPGA SDK for OpenCL User Environment Variables (Windows)](#)
- [Setting the Intel FPGA SDK for OpenCL User Environment Variables (Linux)](#)
- [Intel FPGA SDK for OpenCL FPGA Platforms page](#)

1.5.2 Uninstalling the FPGA Board (uninstall)

To uninstall an FPGA board, invoke the `uninstall` utility command, uninstall the Custom Platform, and unset the relevant environment variables. You must uninstall the existing FPGA board if you migrate your OpenCL application to another FPGA board that belongs to a different Custom Platform.

To uninstall your FPGA board, perform the following tasks:

1. Following your board vendor's instructions to disconnect the board from your machine.

2. Invoke the `aocl uninstall` utility command to remove the current host computer drivers (for example, PCIe® drivers). The Intel FPGA SDK for OpenCL uses these drivers to communicate with the FPGA board.

3. Uninstall the Custom Platform.

4. Unset the *LD_LIBRARY_PATH* (for Linux) or *PATH* (for Windows) environment variable.

5. Unset the *AOCL_BOARD_PACKAGE_ROOT* environment variable.

1.5.3 Querying the Device Name of Your FPGA Board (diagnose)

Some OpenCL software utility commands require you to specify the device name (*<device_name>*). The *<device_name>* refers to the acl number (e.g. acl0 to acl31) that corresponds to the FPGA device. When you query a list of accelerator boards, the OpenCL software produces a list of installed devices on your machine in the order of their device names.

- To query a list of installed devices on your machine, type `aocl diagnose` at a command prompt.
  The software generates an output that resembles the example shown below:

```
aocl diagnose: Running diagnostic from ALTERAOCLSDKROOT/board/<board_name>/
<platform>/libexec

Verified that the kernel mode driver is installed on the host machine.

Using board package from vendor: <board_vendor_name>
Querying information for all supported devices that are installed on the host
machine ...

device_name   Status   Information

acl0          Passed   <descriptive_board_name>
                       PCIe dev_id = <device_ID>, bus:slot.func = 02:00.00,
                         at Gen 2 with 8 lanes.
                       FPGA temperature=43.0 degrees C.

acl1          Passed   <descriptive_board_name>
                       PCIe dev_id = <device_ID>, bus:slot.func = 03:00.00,
                         at Gen 2 with 8 lanes.
                       FPGA temperature = 35.0 degrees C.

Found 2 active device(s) installed on the host machine, to perform a full
diagnostic on a specific device, please run aocl diagnose <device_name>

DIAGNOSTIC_PASSED
```

Related Links

    The host must identify the number of OpenCL FPGA devices installed into the system.

1.5.4 Running a Board Diagnostic Test (diagnose <device_name>)

To perform a detailed diagnosis on a specific FPGA board, include *<device_name>* as an argument of the `diagnose` utility command.

- At a command prompt, invoke the `aocl diagnose` *<device_name>* command, where *<device_name>* is the acl number (for example, acl0 to acl31) that corresponds to your FPGA device.

  You can identify the *<device_name>* when you query the list of installed boards in your system.

Consult your board vendor's documentation for more board-specific information on using the `diagnose` utility command to run diagnostic tests on multiple FPGA boards.

1.5.5 Programming the FPGA Offline or without a Host (program <device_name>)

To program an FPGA device offline or without a host, invoke the `program` utility command.

- At a command prompt, invoke the `aocl program` *<device_name>* *<your_kernel_filename>*`.aocx` command

  where:

  <device_name> refers to the acl number (for example, acl0 to acl31) that corresponds to your FPGA device, and

  <your_kernel_filename>`.aocx` is the executable file you use to program the hardware.

Note: To program an SoC such as the Cyclone V SoC, you must specify the full path of the device when invoking the `program` utility command. For example, `aocl program /dev/<device_name> <your_kernel_filename>.aocx`.

1.5.6 Programming the Flash Memory (flash <device_name>)

If supported, invoke the `flash` utility command to initialize the FPGA with a specified startup configuration.

Note: For instructions on programming the micro SD flash card of the Cyclone V SoC Development Kit, refer to the *Writing an SD Card Image onto the Micro SD Flash Card* section of the *Intel FPGA SDK for OpenCL Cyclone V SoC Getting Started Guide*.

- At a command prompt, invoke the `aocl flash <device_name>`
  `<your_kernel_filename>.aocx` command

  where:

  <device_name> refers to the acl number (for example, acl0 to acl31) that corresponds to your FPGA device, and

  <your_kernel_filename>`.aocx` is the executable file you use to program the hardware.

Related Links

- Writing an SD Card Image onto the Micro SD Flash Card on Windows
- Writing an SD Card Image onto the Micro SD Flash Card on Linux

1.6 Structuring Your OpenCL Kernel

Intel offers recommendations on how to structure your OpenCL kernel code. Consider implementing these programming recommendations when you create a kernel or modify a kernel written originally to target another architecture.

Guidelines for Naming the Kernel on page 20
  Intel recommends that you include only alphanumeric characters in your file names.

Programming Strategies for Optimizing Data Processing Efficiency on page 21
  Optimize the data processing efficiency of your kernel by implementing strategies such as unrolling loops, setting work-group sizes, and specifying compute units and work-items.

Programming Strategies for Optimizing Local Memory Efficiency on page 25
  You can optimize the memory access efficiency of your kernel by implementing strategies such as specifying local memory pointer size. Specifying the local memory pointer size affects the size of the local memory size, which affects the local memory hardware footprint.

Implementing the Intel FPGA SDK for OpenCL Channels Extension on page 26
  The Intel FPGA SDK for OpenCL channels extension provides a mechanism for passing data to kernels and synchronizing kernels with high efficiency and low latency.

Implementing OpenCL Pipes on page 45
  The Intel FPGA SDK for OpenCL provides preliminary support for OpenCL pipe functions.

Implementing Arbitrary Precision Integers on page 61
  Use the Intel FPGA SDK for OpenCL arbitrary precision integer extension to define integers with a custom bit-width. You can define integer custom bit-widths up to and including 64 bits.

Using Predefined Preprocessor Macros in Conditional Compilation on page 62
  You may take advantage of predefined preprocessor macros that allow you to conditionally compile portions of your kernel code.

Declaring __constant Address Space Qualifiers on page 63
  There are several limitations and workarounds you must consider when you include `__constant` address space qualifiers in your kernel.

Including Structure Data Types as Arguments in OpenCL Kernels on page 64

Convert each structure parameter (`struct`) to a pointer that points to a structure.

Inferring a Register on page 68
In general, the offline compiler chooses registers if the access to a variable is fixed and does not require any dynamic indexes.

Enabling Double Precision Floating-Point Operations on page 69
The Intel FPGA SDK for OpenCL offers preliminary support for all double precision floating-point functions.

Single-Cycle Floating-Point Accumulator for Single Work-Item Kernels on page 70
Single work-item kernels that perform accumulation in a loop can leverage the single-cycle floating-point accumulator feature of the Intel FPGA SDK for OpenCL Offline Compiler.

1.6.1 Guidelines for Naming the Kernel

Intel recommends that you include only alphanumeric characters in your file names.

- Begin a file name with an alphanumeric character.

  If the file name of your OpenCL application begins with a nonalphanumeric character, compilation fails with the following error message:

  ```
  Error: Quartus compilation FAILED
  See quartus_sh_compile.log for the output log.
  ```

- Do not differentiate file names using nonalphanumeric characters.

  The Intel FPGA SDK for OpenCL Offline Compiler translates any nonalphanumeric character into an underscore ("_"). If you differentiate two file names by ending them with different nonalphanumeric characters only (for example, `myKernel#.cl` and `myKernel&.cl`), the offline compiler translates both file names to *<your_kernel_filename>*_.cl (for example, `myKernel_.cl`).

- For Windows system, ensure that the combined length of the kernel file name and its file path does not exceed 260 characters.

  64-bit Windows 7 and Windows 8.1 has a 260-character limit on the length of a file path. If the combined length of the kernel file name and its file path exceeds 260 characters, the offline compiler generates the following error message:

  ```
  The filename or extension is too long.
  The system cannot find the path specified.
  ```

  In addition to the compiler error message, the following error message appears in the *<your_kernel_filename>*/quartus_sh_compile.log file:
  Error: Can't copy *<file_type>* files: Can't open *<your_kernel_filename>* for write: No such file or directory

  For Windows 10, you can remove the 260-character limit. For more information, see your Windows 10 documentation.

- Do not name your .cl OpenCL kernel source file "kernel". Naming the source file kernel.cl causes the offline compiler to generate intermediate design files that have the same names as certain internal files, which leads to an compilation error.

1.6.2 Programming Strategies for Optimizing Data Processing Efficiency

Optimize the data processing efficiency of your kernel by implementing strategies such as unrolling loops, setting work-group sizes, and specifying compute units and work-items.

1.6.2.1 Unrolling a Loop

The Intel FPGA SDK for OpenCL Offline Compiler might unroll simple loops even if they are not annotated by a pragma.
To direct the offline compiler to unroll a loop, insert an `unroll` kernel pragma in the kernel code preceding a loop you wish to unroll.

Attention:

- Provide an unroll factor whenever possible. To specify an unroll factor *N*, insert the `#pragma unroll <N>` directive before a loop in your kernel code.

  The offline compiler attempts to unroll the loop at most *<N>* times.

  Consider the code fragment below. By assigning a value of 2 as an argument to `#pragma unroll`, you direct the offline compiler to unroll the loop twice.

  ```
  #pragma unroll 2
  for(size_t k = 0; k < 4; k++)
  {
      mac += data_in[(gid * 4) + k] * coeff[k];
  }
  ```

- To unroll a loop fully, you may omit the unroll factor by simply inserting the `#pragma unroll` directive before a loop in your kernel code.

  The offline compiler attempts to unroll the loop fully if it understands the trip count. The offline compiler issues a warning if it cannot execute the unroll request.

1.6.2.2 Coalescing Nested Loops

Use the `loop_coalesce` pragma to direct the Intel FPGA SDK for OpenCL Offline Compiler to coalesce nested loops into a single loop without affecting the loop functionality. Coalescing loops can help reduce your kernel area usage by reducing the overhead needed for the loops.

For NDRange kernels, the compiler automatically attempts to coalesce loops even if they are not annotated by the `loop_coalesce` pragma. Coalescing loops in NDRange kernels improves throughtput as well as reducing kernel area usage. You can use the `loop_coalesce` pragma to prevent the automatic coalescing of loops in NDRange kernels.
To coalesce nested loops, specify the pragma as follows:

```
#pragma loop_coalesce <loop_nesting_level>
```

The *<loop_nesting_level>* parameter is optional and is an integer that specifies how many nested loop levels that you want the compiler to attempt to coalesce. If you do not specify the *<loop_nesting_level>* parameter, the compiler attempts to coalesce all of the nested loops.

For example, consider a set of nested loops like this:

```
for (A)
  for (B)
    for (C)
      for (D)
    for (E)
```

If you put the pragma before loop (A), the loop nesting level for these loops is as follows:

- Loop (A) has a loop nesting level of 1.
- Loop (B) has a loop nesting level of 2.
- Loop (C) has a loop nesting level of 3.
- Loop (D) has a loop nesting level of 4.
- Loop (E) has a loop nesting level of 3.

Depending on the loop nesting level that you specify, the compiler attempts to coalesce loops differently:

- If you specify `#pragma loop_coalesce 1`, the compiler does not attempt to coalesce any of the nested loops.
- If you specify `#pragma loop_coalesce 2`, the compiler attempts to coalesce loops (A) and (B).
- If you specify `#pragma loop_coalesce 3`, the compiler attempts to coalesce loops (A), (B), (C), and (E).
- If you specify `#pragma loop_coalesce 4`, the compiler attempts to coalesce all of the loops [loop (A) - loop (E)].

Important: If you specify `#pragma loop_coalesce 1` for a loop in an NDRange kernel, you prevent the automatic loop coalescing for that loop.

Example

The following simple example shows how the compiler coalesces two loops into a single loop.

Consider a simple nested loop written as follows:

```
#pragma loop_coalesce
for(int i = 0; i < N; i++)
  for(int j = 0; j < M; j++)
    sum[i][j] += i+j;
```

The compiler coalesces the two loops together so that they run as if they were a single loop written as follows:

```
int i = 0;
int j = 0;
while(i < N){

  sum[i][j] += i+j;
  j++;

  if(j == M){
    j=0;
```

```
        i++;
    }
}
```

1.6.2.3 Specifying a loop initiation interval (II)

Use the `ii` pragma to direct the Intel FPGA SDK for OpenCL Offline Compiler to attempt to set the loop initiation interval (II) for the loop that follows the pragma declaration. If the compiler cannot achieve the specified II for the loop, then your compile fails. The iteration interval, or II, is the number of clock cycles between launching successive loop iterations. The higher the II value, the longer the wait before the next loop iteration.

For some loops in your kernel, specifying a higher II with the `ii` pragma than the compiler might choose can increase the maximum operating frequency (f_{max}) of your kernel without losing throughput.

A loop is a good candidate to have the `ii` pragma applied to it if the loop meets the following conditions:

- The loop is not critical to the throughput of your kernel.

- The running time of the loop is small compared to other loops it might contain.

To specify a loop initiation interval for a loop, specify the pragma before the loop as follows:

```
#pragma ii <desired_initiation_interval>
```

The *<desired_initiation_interval>* parameter is required and is an integer that specifies the number of clock cycles to wait between successive loop iterations.

Example

Consider a case where your kernel has two distinct, pipelineable loops: a short-running initialization loop that has a loop-carried dependence and a long-running loop that does the bulk of your processing. In this case, the compiler does not know that the initialization loop has a much smaller impact on the overall throughput of your design. If possible, the compiler attempts to pipeline both loops with an II of 1.

Because the initialization loop has a loop-carried dependence, it will have a feedback path in the generated hardware. To achieve an II with such a feedback path, some clock frequency might be sacrificed. Depending on the feedback path in the main loop, the rest of your design could have run at a higher frequency.

If you specify `#pragma ii 2` on the initialization loop, you tell the compiler that it can be less aggressive in optimizing II for this loop. Less aggressive optimization allows the compiler to pipeline the path limiting the maximum operating frequency (f_{max}) and could allow your overall kernel design to achieve a higher f_{max}.

The initialization loop takes longer to run with its new II. However, the decrease in the running time of the long-running loop due to higher f_{max} compensates for the increased length in running time of the initialization loop.

1.6.2.4 Specifying Work-Group Sizes

Specify a maximum or required work-group size whenever possible. The Intel FPGA SDK for OpenCL Offline Compiler relies on this specification to optimize hardware usage of the OpenCL kernel without involving excess logic.

If you do not specify a `max_work_group_size` or a `reqd_work_group_size` attribute in your kernel, the work-group size assumes a default value depending on compilation time and runtime constraints.

- If your kernel contains a barrier, the offline compiler sets a default maximum work-group size of 256 work-items.

- If your kernel contains a barrier or refers to the local work-item ID, or if you query the work-group size in your host code, the runtime defaults the work-group size to one work-item.

- If your kernel does not contain a barrier or refer to the local work-item ID, or if your host code does not query the work-group size, the runtime defaults the work-group size to the global NDRange size.

To specify the work-group size, modify your kernel code in the following manner:

- To specify the maximum number of work-items that the offline compiler may allocate to a work-group in a kernel, insert the `max_work_group_size(N)` attribute in your kernel source code.

  For example:

  ```
  __attribute__((max_work_group_size(512)))
  __kernel void sum (__global const float * restrict a,
                     __global const float * restrict b,
                     __global float * restrict answer)
  {
      size_t gid = get_global_id(0);
      answer[gid] = a[gid] + b[gid];
  }
  ```

- To specify the required number of work-items that the offline compiler allocates to a work-group in a kernel, insert the `reqd_work_group_size(X, Y, Z)` attribute to your kernel source code.

  For example:

  ```
  __attribute__((reqd_work_group_size(64,1,1)))
  __kernel void sum (__global const float * restrict a,
                     __global const float * restrict b,
                     __global float * restrict answer)
  {
      size_t gid = get_global_id(0);
      answer[gid] = a[gid] + b[gid];
  }
  ```

  The offline compiler allocates the exact amount of hardware resources to manage the work-items in a work-group.

1.6.2.5 Specifying Number of Compute Units

To increase the data-processing efficiency of an OpenCL kernel, you can instruct the Intel FPGA SDK for OpenCL Offline Compiler to generate multiple kernel compute units. Each compute unit is capable of executing multiple work-groups simultaneously.

Caution: Multiplying the number of kernel compute units increases data throughput at the expense of global memory bandwidth contention among compute units.

- To specify the number of compute units for a kernel, insert the `num_compute_units(N)` attribute in the kernel source code.

  For example, the code fragment below directs the offline compiler to instantiate two compute units in a kernel:

  ```
  __attribute__((num_compute_units(2)))
  __kernel void test(__global const float * restrict a,
                     __global const float * restrict b,
                     __global float * restrict answer)
  {
     size_t gid = get_global_id(0);
     answer[gid] = a[gid] + b[gid];
  }
  ```

  The offline compiler distributes work-groups across the specified number of compute units.

1.6.2.6 Specifying Number of SIMD Work-Items

To increase the data-processing efficiency of an OpenCL kernel, specify the number of work-items within a work-group that the Intel FPGA SDK for OpenCL Offline Compiler executes in a single instruction multiple data (SIMD) manner.

Important: Introduce the `num_simd_work_items` attribute in conjunction with the `reqd_work_group_size` attribute. The `num_simd_work_items` attribute you specify must evenly divide the work-group size you specify for the `reqd_work_group_size` attribute.

- To specify the number of SIMD work-items in a work-group, insert the `num_simd_work_item(N)` attribute in the kernel source code.

  For example, the code fragment below assigns a fixed work-group size of 64 work-items to a kernel. It then consolidates the work-items within each work-group into four SIMD vector lanes:

  ```
  __attribute__((num_simd_work_items(4)))
  __attribute__((reqd_work_group_size(64,1,1)))
  __kernel void test(__global const float * restrict a,
                     __global const float * restrict b,
                     __global float * restrict answer)
  {
     size_t gid = get_global_id(0);
     answer[gid] = a[gid] + b[gid];
  }
  ```

  The offline compiler replicates the kernel datapath according to the value you specify for `num_simd_work_items` whenever possible.

1.6.3 Programming Strategies for Optimizing Local Memory Efficiency

You can optimize the memory access efficiency of your kernel by implementing strategies such as specifying local memory pointer size. Specifying the local memory pointer size affects the size of the local memory size, which affects the local memory hardware footprint.

- To specify a pointer size other than the default size of 16 kilobytes (kB), include the `local_mem_size(N)` attribute in the pointer declaration within your kernel source code.

  For example:

  ```
  __kernel void myLocalMemoryPointer(
                      __local float * A,
                      __attribute__((local_mem_size(1024))) __local float * B,
                      __attribute__((local_mem_size(32768))) __local float * C)
                      {
                      //statements
                      }
  ```

  In the `myLocalMemoryPointer` kernel, 16 kB of local memory (default) is allocated to pointer `A`, 1 kB is allocated to pointer `B`, and 32 kB is allocated to pointer `C`.

1.6.4 Implementing the Intel FPGA SDK for OpenCL Channels Extension

The Intel FPGA SDK for OpenCL channels extension provides a mechanism for passing data to kernels and synchronizing kernels with high efficiency and low latency.

Attention: If you want to leverage the capabilities of channels but have the ability to run your kernel program using other SDKs, implement OpenCL pipes instead.

Related Links

Implementing OpenCL Pipes on page 45
> The Intel FPGA SDK for OpenCL provides preliminary support for OpenCL pipe functions.

1.6.4.1 Overview of the Intel FPGA SDK for OpenCL Channels Extension

The Intel FPGA SDK for OpenCL channels extension allows kernels to communicate directly with each other via FIFO buffers.

Implementation of channels decouples kernel execution from the host processor. Unlike the typical OpenCL execution model, the host does not need to coordinate data movement across kernels.

Figure 5. **Overview of Channels Implementation**
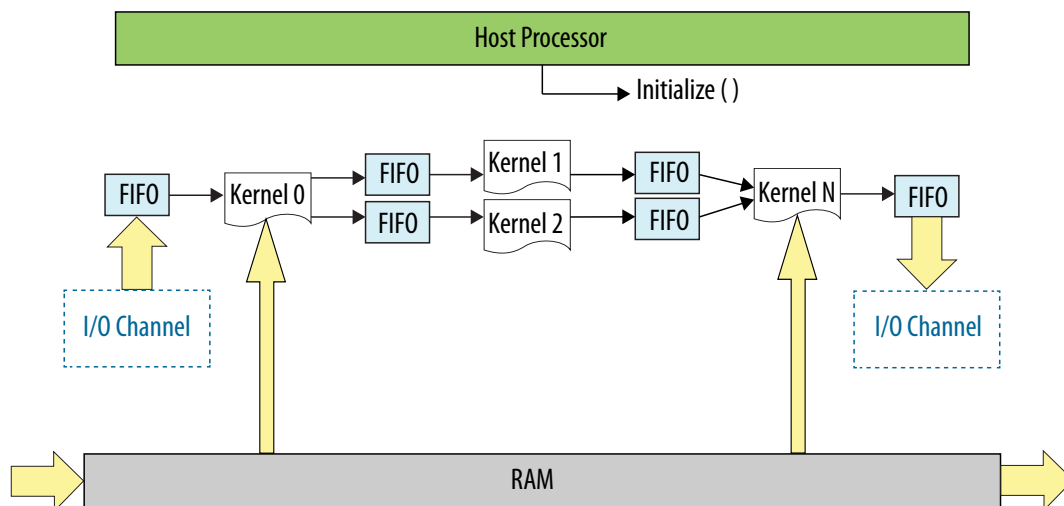
1.6.4.2 Channel Data Behavior

Data written to a channel remains in a channel as long as the kernel program remains loaded on the FPGA device. In other words, data written to a channel persists across multiple work-groups and NDRange invocations. However, data is not persistent across multiple or different invocations of kernel programs.

Consider the following code example:

```
channel int c0;

__kernel void producer()
{
    for(int i=0; i < 10; i++)
    {
        write_channel_intel(c0, i);
    }
}

__kernel void consumer( __global uint * restrict dst )
{
    for(int i=0; i < 5; i++)
    {
        dst[i] = read_channel_intel(c0);
    }
}
```

Figure 6. **Channel Data FIFO Ordering**

The kernel `producer` writes ten elements ([0, 9]) to the channel. The kernel `consumer` reads five elements from the channel per NDRange invocation. During the first invocation, the kernel `consumer` reads values 0 to 4 from the channel. Because the data persists across NDRange invocations, the second time you execute the kernel `consumer`, it reads values 5 to 9.

For this example, to avoid a deadlock from occurring, you need to invoke the kernel `consumer` twice for every invocation of the kernel `producer`. If you call `consumer` less than twice, `producer` stalls because the channel becomes full. If you call `consumer` more than twice, `consumer` stalls because there is insufficient data in the channel.

1.6.4.3 Multiple Work-Item Ordering for Channels

The OpenCL specification does not define a work-item ordering. The Intel FPGA SDK for OpenCL enforces a work-item order to maintain the consistency in channel read and write operations.

Multiple work-item accesses to a channel can be useful in some scenarios. For example, they are useful when data words in the channel are independent, or when the channel is implemented for control logic. The main concern regarding multiple work-item accesses to a channel is the order in which the kernel writes data to and reads data from the channel. If possible, the SDK's channels extension processes work-items read and write operations to the channel in a deterministic order. As such, the read and write operations remain consistent across kernel invocations.

Requirements for Deterministic Multiple Work-Item Ordering

To guarantee deterministic ordering, the SDK checks that the channel call is work-item invariant based on the following characteristics:

- All paths through the kernel must execute the channel call.

- If the first requirement is not satisfied, none of the branch conditions that reach the channel call should execute in a work-item-dependent manner.

If the SDK cannot guarantee deterministic ordering of multiple work-item accesses to a channel, it warns you that the channels might not have well-defined ordering with nondeterministic execution. Primarily, the SDK fails to provide deterministic ordering if you have work-item-variant code on loop executions with channel calls, as illustrated below:

```
__kernel void ordering( __global int * restrict check,
                        __global int * restrict data )
{
    int condition = check[get_global_id(0)];

    if(condition)
    {
        for(int i=0; i < N, i++)
        {
            process(data);
            write_channel_altera(req, data[i]);
        }
    }
    else
    {
```

```
          process(data);
      }
}
```

1.6.4.3.1 Work-Item Serial Execution of Channels

Work-item serial execution refers to an ordered execution behavior where work-item sequential IDs determine their execution order in the compute unit.

When you implement channels in a kernel, the Intel FPGA SDK for OpenCL Offline Compiler enforces that kernel behavior is equivalent to having at most one work-group in flight. The compiler also ensures that the kernel executes channels in work-item serial execution, where the kernel executes work-items with smaller IDs first. A work-item has the identifier (x, y, z, group), where x, y, z are the local 3D identifiers, and group is the work-group identifier.

The work-item ID (x0, y0, z0, group0) is considered to be smaller than the ID (x1, y1, z1, group1) if one of the following conditions is true:

* group0 < group1
* group0 = group1 and z0 < z1
* group0 = group1 and z0 = z1 and y0 < y1
* group0 = group1 and z0 = z1 and y0 = y1 and x0 < x1

Work-items with incremental IDs execute in a sequential order. For example, the work-item with an ID (x0, y0, z0, group0) executes the write channel call first. Then, the work-item with an ID (x1, y0, z0, group0) executes the call, and so on. Defining this order ensures that the system is verifiable with external models.

Channel Execution in Loop with Multiple Work-Items

When channels exist in the body of a loop with multiple work-items, as shown below, each loop iteration executes prior to subsequent iterations. This implies that loop iteration 0 of each work-item in a work-group executes before iteration 1 of each work-item in a work-group, and so on.

```
__kernel void ordering( __global int * data, int X )
{
  int n=0;
  while (n<X)
  {
    write_channel_intel(req, data[get_global_id(0)]);
    n++;
  }
}
```

1.6.4.4 Restrictions in the Implementation of Intel FPGA SDK for OpenCL Channels Extension

There are certain design restrictions to the implementation of channels in your OpenCL application.

Single Call Site

Because the channel read and write operations do not function deterministically, for a given kernel, you can only assign one call site per channel ID. For example, the Intel FPGA SDK for OpenCL Offline Compiler cannot compile the following code example:

```
in_data1 = read_channel_intel(channel1);
in_data2 = read_channel_intel(channel2);
in_data3 = read_channel_intel(channel1);
```

The second `read_channel_intel` call to `channel1` causes compilation failure because it creates a second call site to `channel1`.

To gather multiple data from a given channel, divide the channel into multiple channels, as shown below:

```
in_data1 = read_channel_intel(channel1);
in_data2 = read_channel_intel(channel2);
in_data3 = read_channel_intel(channel3);
```

Because you can only assign a single call site per channel ID, you cannot unroll loops containing channels. Consider the following code:

```
#pragma unroll 4
for (int i=0; i < 4; i++)
{
    in_data = read_channel_intel(channel1);
}
```

The offline compiler issues the following warning message during compilation:

```
Compiler Warning: Unroll is required but the loop cannot be
unrolled.
```

Feedback and Feed-Forward Channels

Channels within a kernel can be either `read_only` or `write_only`. Performance of a kernel that reads and writes to the same channel is poor.

Static Indexing

The Intel FPGA SDK for OpenCL channels extension does not support dynamic indexing into arrays of channel IDs.

Consider the following example:

```
channel int ch[WORKGROUP_SIZE];

__kernel void consumer()
{

    int gid = get_global_id(0);
    int value = read_channel_intel(ch[gid]);

    //statements
}
```

Compilation of this example kernel fails with the following error message:

```
Compiler Error: Indexing into channel array ch could not be resolved to all
constant
```

To avoid this compilation error, index into arrays of channel IDs statically, as shown below:

```
channel int ch[WORKGROUP_SIZE];

__kernel void consumer()
{

    int gid = get_global_id(0);
    int value;

    switch(gid)
    {
        case 0: value = read_channel_intel(ch[0]); break;
        case 1: value = read_channel_intel(ch[1]); break;
        case 2: value = read_channel_intel(ch[2]); break;
        case 3: value = read_channel_intel(ch[3]); break;
        //statements

        case WORKGROUP_SIZE-1:read_channel_intel(ch[WORKGROUP_SIZE-1]); break;
    }

    //statements
}
```

Kernel Vectorization Support

You cannot vectorize kernels that use channels; that is, do not include the `num_simd_work_items` kernel attribute in your kernel code. Vectorizing a kernel that uses channels creates multiple channel masters and requires arbitration, which the SDK's channels extension does not support.

Instruction-Level Parallelism on read_channel_intel and write_channel_intel Calls

If no data dependencies exist between `read_channel_intel` and `write_channel_intel` calls, the offline compiler attempts to execute these instructions in parallel. As a result, the offline compiler might execute these `read_channel_intel` and `write_channel_intel` calls in an order that does not follow the sequence expressed in the OpenCL kernel code.

Consider the following code sequence:

```
in_data1 = read_channel_intel(channel1);
in_data2 = read_channel_intel(channel2);
in_data3 = read_channel_intel(channel3);
```

Because there are no data dependencies between the `read_channel_intel` calls, the offline compiler can execute them in any order.

1.6.4.5 Enabling the Intel FPGA SDK for OpenCL Channels for OpenCL Kernel

To implement the Intel FPGA SDK for OpenCL channels extension, modify your OpenCL kernels to include channels-specific pragma and API calls.

Channel declarations are unique within a given OpenCL kernel program. Also, channel instances are unique for every OpenCL kernel program device pair. If the runtime loads a single OpenCL kernel program onto multiple devices, each device will have a single copy of the channel. However, these channel copies are independent and do not share data across the devices.

1.6.4.5.1 Declaring the Channel Handle

Use the channel variable to define the connectivity between kernels or between kernels and I/O.

To read from and write to a channel, the kernel must pass the channel variable to each of the corresponding API call.

- Declare the channel handle as a file scope variable in the kernel source code in the following convention: `channel <type> <variable_name>`

  For example: `channel int c;`

- The Intel FPGA SDK for OpenCL channel extension supports simultaneous channel accesses by multiple variables declared in a data structure. Declare a `struct` data structure for a channel in the following manner:

```
typedef struct type_ {
    int a;
    int b;
} type_t;

channel type_t foo;
```

1.6.4.5.2 Implementing Blocking Channel Write Extensions

The `write_channel_intel` API call allows you to send data across a channel.

Note: The write channel calls support single-call sites only. For a given channel, only one write channel call to it can exist in the entire kernel program.

- To implement a blocking channel write, include the following `write_channel_intel` function signature:

  `void write_channel_intel (channel <type> channel_id, const <type> data);`

  Where:

  `channel_id` identifies the buffer to which the channel connects, and it must match the `channel_id` of the corresponding read channel (`read_channel_intel`).

  `data` is the data that the channel write operation writes to the channel. Data *<type>* must match the *<type>* of the `channel_id`.

  <type> defines a channel data width, which cannot be a constant. Follow the OpenCL conversion rules to ensure that data the kernel writes to a channel is convertible to *<type>*.

The following code snippet demonstrates the implementation of the `write_channel_intel` API call:

```
//Defines chan, the kernel file-scope channel variable.
channel long chan;

/*Defines the kernel which reads eight bytes (size of long) from global memory,
and passes this data to the channel.*/
__kernel void kernel_write_channel( __global const long * src )
{
    for(int i=0; i < N; i++)
    {
        //Writes the eight bytes to the channel.
        write_channel_intel(chan, src[i]);
    }
}
```

Caution: When you send data across a write channel using the `write_channel_intel` API call, keep in mind that if the channel is full (that is, if the FIFO buffer is full of data), your kernel will stall. Use the Intel FPGA SDK for OpenCL Profiler to check for channel stalls.

Related Links

Profiling Your OpenCL Kernel on page 112
> The Intel FPGA SDK for OpenCL Profiler measures and reports performance data collected during OpenCL kernel execution on the FPGA.

Implementing Nonblocking Channel Write Extensions

Perform nonblocking channel writes to facilitate applications where data write operations might not occur. A nonblocking channel write extension returns a Boolean value that indicates whether data is written to the channel.

Consider a scenario where your application has one data producer with two identical workers. Assume the time each worker takes to process a message varies depending on the contents of the data. In this case, there might be situations where one worker is busy while the other is free. A nonblocking write can facilitate work distribution such that both workers are busy.

- To implement a nonblocking channel write, include the following `write_channel_nb_altera` function signature:

    bool write_channel_nb_altera(channel *<type>* channel_id, const *<type>* data);

The following code snippet of the kernel `producer` facilitates work distribution using the nonblocking channel write extension:

```
channel long worker0, worker1;
__kernel void producer( __global const long * src )
{
    for(int i=0; i < N; i++)
    {

        bool success = false;
        do
        {
            success = write_channel_nb_altera(worker0, src[i]);
            if(!success)
            {
                success = write_channel_nb_altera(worker1, src[i]);
```

```
            }
        }
        while(!success);
    }
}
```

1.6.4.5.3 Implementing Blocking Channel Read Extensions

The `read_channel_intel` API call allows you to receive data across a channel.

Note: The read channel calls support single-call sites only. For a given channel, only one read channel call to it can exist in the entire kernel program.

- To implement a blocking channel read, include the following `read_channel_intel` function signature:

  <type> read_channel_intel(channel *<type>* channel_id);

  Where:

  `channel_id` identifies the buffer to which the channel connects, and it must match the `channel_id` of the corresponding write channel (`write_channel_intel`).

  <type> defines a channel data width, which cannot be a constant. Ensure that the variable the kernel assigns to read the channel data is convertible from *<type>*.

The following code snippet demonstrates the implementation of the `read_channel_intel` API call:

```
//Defines chan, the kernel file-scope channel variable.
channel long chan;

/*Defines the kernel, which reads eight bytes (size of long) from the channel and
writes it back to global memory.*/
__kernel void kernel_read_channel( __global long * dst );
{
    for(int i=0; i < N; i++)
    {
        //Reads the eight bytes from the channel.
        dst[i] = read_channel_intel(chan);
    }
}
```

Caution: If the channel is empty (that is, if the FIFO buffer is empty), you cannot receive data across a read channel using the `read_channel_intel` API call. Doing so causes your kernel to stall.

Implementing Nonblocking Channel Read Extensions

Perform nonblocking reads to facilitate applications where data is not always available. The nonblocking reads signature is similar to blocking reads. However, it returns an integer value that indicates whether a read operation takes place successfully.

- To implement a blocking channel write, include the following `read_channel_nb_intel` function signature:

  <type> read_channel_nb_intel(channel *<type>* channel_id, bool * valid);

The following code snippet demonstrates the use of the nonblocking channel read extension:

```
channel long chan;

__kernel void kernel_read_channel( __global long * dst )
{
    int i=0;
    while(i < N)
    {
        bool valid0, valid1;
        long data0 = read_channel_nb_intel(chan, &valid0);
        long data1 = read_channel_nb_intel(chan, &valid1);
        if (valid0)
        {
            process(data0);
        }
        if (valid1)
        {
            process(data1);
        }
    }
}
```

1.6.4.5.4 Implementing I/O Channels Using the io Channels Attribute

Include an `io` attribute in your channel declaration to declare a special I/O channel to interface with input or output features of an FPGA board.
These features might include network interfaces, PCIe, cameras, or other data capture or processing devices or protocols.

The `io("chan_id")` attribute specifies the I/O feature of an accelerator board with which a channel interfaces, where *chan_id* is the name of the I/O interface listed in the `board_spec.xml` file of your Custom Platform.

Because peripheral interface usage might differ for each device type, consult your board vendor's documentation when you implement I/O channels in your kernel program. Your OpenCL kernel code must be compatible with the type of data generated by the peripheral interfaces.

Caution:
- Implicit data dependencies might exist for channels that connect to the board directly and communicate with peripheral devices via I/O channels. These implicit data dependencies might lead to compilation issues because the Intel FPGA SDK for OpenCL Offline Compiler cannot identify these dependencies.

- External I/O channels communicating with the same peripherals do not obey any sequential ordering. Ensure that the external device does not require sequential ordering because unexpected behavior might occur.

1. Consult the `board_spec.xml` file in your Custom Platform to identify the input and output features available on your FPGA board.

   For example, a `board_spec.xml` file might include the following information on I/O features:

```
<channels>
  <interface name="udp_0" port="udp0_out"  type="streamsource" width="256"
   chan_id="eth0_in"/>
  <interface name="udp_0" port="udp0_in"  type="streamsink" width="256"
   chan_id="eth0_out"/>
  <interface name="udp_0" port="udp1_out"  type="streamsource" width="256"
```

```
  chan_id="eth1_in"/>
 <interface name="udp_0" port="udp1_in"  type="streamsink" width="256"
  chan_id="eth1_out"/>
</channels>
```

The `width` attribute of an `interface` element specifies the width, in bits, of the data type used by that channel. For the example above, both the `uint` and `float` data types are 32 bits wide. Other bigger or vectorized data types must match the appropriate bit width specified in the `board_spec.xml` file.

2. Implement the `io` channel attribute as demonstrated in the following code example. The `io` channel attribute names must match those of the I/O channels (`chan_id`) specified in the `board_spec.xml` file.

```
channel QUDPWord udp_in_IO __attribute__((depth(0)))
                           __attribute__((io("eth0_in")));
channel QUDPWord udp_out_IO __attribute__((depth(0)))
                           __attribute__((io("eth0_out")));

__kernel void io_in_kernel( __global ulong4 *mem_read,
                            uchar read_from,
                            int size )
{
  int index = 0;
  ulong4 data;
  int half_size = size >> 1;
  while (index < half_size)
  {
    if (read_from & 0x01)
    {
      data = read_channel_intel(udp_in_IO);
    }
    else
    {
      data = mem_read[index];
    }
    write_channel_intel(udp_in, data);
    index++;
  }
}

__kernel void io_out_kernel( __global ulong2 *mem_write,
                             uchar write_to,
                             int size )
{
  int index = 0;
  ulong4 data;
  int half_size = size >> 1;
  while (index < half_size)
  {
    ulong4 data = read_channel_intel(udp_out);
    if (write_to & 0x01)
    {
      write_channel_intel(udp_out_IO, data);
    }
    else
    {
      //only write data portion
      ulong2 udp_data;
      udp_data.s0 = data.s0;
      udp_data.s1 = data.s1;
      mem_write[index] = udp_data;
    }
    index++;
  }
}
```

> ***Attention:*** Declare a unique `io("chan_id")` handle for each I/O channel specified in the channels eXtensible Markup Language (XML) element within the `board_spec.xml` file.

1.6.4.5.5 Emulating I/O Channels

When you emulate a kernel that has a channel declared with the `io` attribute, I/O channel input is emulated by reading from a file, and channel output is emulated by writing to a file.

When you run your emulation, the file name used for reading or writing matches the name in the `io` attribute. For example, if you have a channel declaration as follows, your emulation would read or write (but not both) to a file called `myIOChannel`.

```
channel uint chanA __attribute__((io("myIOChannel")));
```

Channels are unidirectional. You can either read from a channel or write to a channel, but not both. However, you can have separate read channels and write channels with the same `io` attribute value.

```
channel uint readChannel __attribute__((io("myIOChannel")));
channel uint writeChannel __attribute__((io("myIOChannel")));
```

Emulating Reading from an I/O Channel

If a read call issued to a channel with an `io` attribute called `myfile`, a read attempt is made on a file on the disk called `myfile`. If the `myfile` file does not exist or there is insufficient data to read from the file, the behavior depends on the call type:

Non-blocking read call	If the file does not exist or there is insufficient data, the read attempt returns with a failure message.
Blocking read call	If the file does not exist or there is insufficient data, the read attempt blocks your program until the file is created on the disk, or the file contains sufficient data.

Emulating Writing to an I/O Channel

If a write call is issued to a channel with an `io` attribute called `myfile`, a write attempt is made to a file on the disk called `myfile`. If the `myfile` file does not exist, a regular file is created and written to. If the `myfile` file exists, it is overwritten. If the write fails, the behavior depends on the call type:

Non-blocking write call	If the write attempt fails, an error is returned.
Blocking write call	If the write attempt fails, further write attempts are made.

Emulating Communication Between a Kernel and a Host or Other Process (Linux only)

Host channels, which enable host-kernel communication, are not enabled in the emulator. However, on Linux systems, you can use the `io` attribute along with named pipes (also known as FIFOs) to emulate host-kernel communication or communication between the kernel and another process.

You must create the FIFO manually using the `mkfifo` command before you start your kernel emulation. The FIFO name must match the `io` attribute name of the channel.

1.6.4.5.6 Use Models of Intel FPGA SDK for OpenCL Channels Implementation

Concurrent execution can improve the effectiveness of channels implementation in your OpenCL kernels.
During concurrent execution, the host launches the kernels in parallel. The kernels share memory and can communicate with each other through channels where applicable.

The use models provide an overview on how to exploit concurrent execution safely and efficiently.

Feed-Forward Design Model

Implement the feed-forward design model to send data from one kernel to the next without creating any cycles between them. Consider the following code example:

```
__kernel void producer( __global const uint * src,
                        const uint iterations )
{
    for(int i=0; i < iterations; i++)
    {
        write_channel_intel(c0, src[2*i]);
        mem_fence(CLK_CHANNEL_MEM_FENCE);
        write_channel_intel(c1, src[2*i+1]);
    }
}

__kernel void consumer( __global uint * dst,
                        const uint iterations )
{
    for (int i=0;i<iterations;i++)
    {
        dst[2*i] = read_channel_intel(c0);
        mem_fence(CLK_CHANNEL_MEM_FENCE);
        dst[2*i+1] = read_channel_intel(c1);
    }
}
```

The `producer` kernel writes data to channels `c0` and `c1`. The `consumer` kernel reads data from `c0` and `c1`. The figure below illustrates the feed-forward data flow between the two kernels:

Figure 7.   Feed-Forward Data Flow

Buffer Management

In the feed-forward design model, data traverses between the `producer` and `consumer` kernels one word at a time. To facilitate the transfer of large data messages consisting of several words, you can implement a ping-pong buffer, which is a common design pattern found in applications for communication. The figure below illustrates the interactions between kernels and a ping-pong buffer:

Figure 8.   Feed-Forward Design Model with Buffer Management
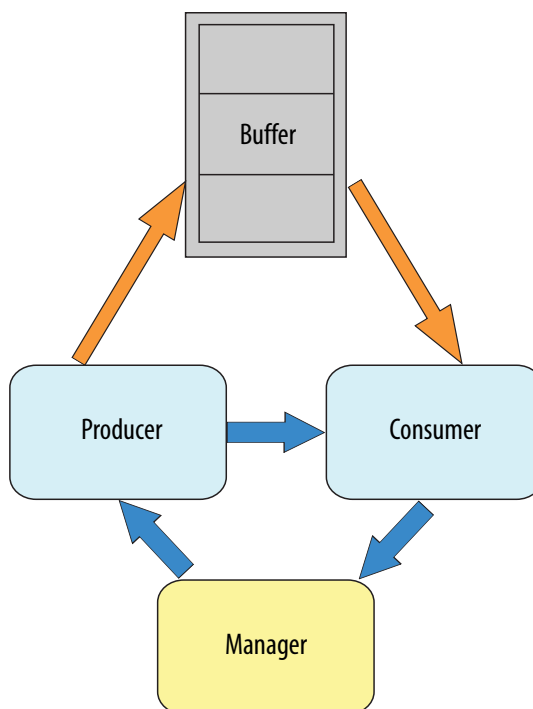
The `manager` kernel manages circular buffer allocation and deallocation between the `producer` and `consumer` kernels. After the `consumer` kernel processes data, the `manager` receives memory regions that the `consumer` frees up and sends them to the `producer` for reuse. The `manager` also sends to the `producer` kernel the initial set of free locations, or tokens, to which the `producer` can write data.

The following figure illustrates the sequence of events that take place during buffer management:

Figure 9.   Kernels Interaction during Buffer Management
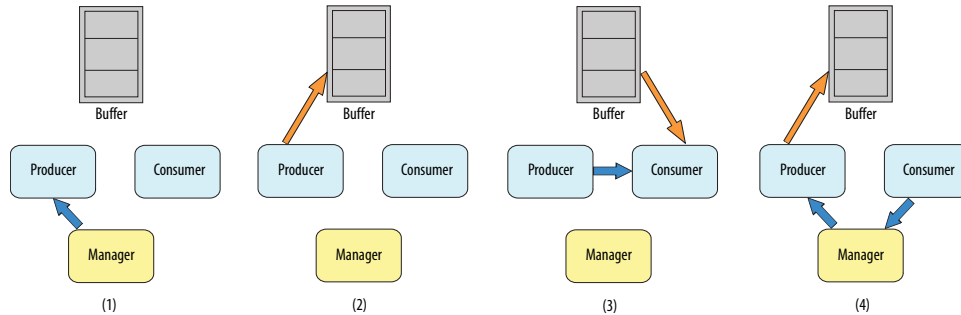
1. The `manager` kernel sends a set of tokens to the `producer` kernel to indicate initially which regions in memory are free for `producer` to use.

2. After `manager` allocates the memory region, `producer` writes data to that region of the ping-pong buffer.

3. After `producer` completes the write operation, it sends a synchronization token to the `consumer` kernel to indicate what memory region contains data for processing. The `consumer` kernel then reads data from that region of the ping-pong buffer.

   Note: When `consumer` is performing the read operation, `producer` can write to other free memory locations for processing because of the concurrent execution of the `producer`, `consumer`, and `manager` kernels.

4. After `consumer` completes the read operation, it releases the memory region and sends a token back to the `manager` kernel. The `manager` kernel then recycles that region for `producer` to use.

Implementation of Buffer Management for OpenCL Kernels

To ensure that the SDK implements buffer management properly, the ordering of channel read and write operations is important. Consider the following kernel example:

```
__kernel void producer( __global const uint * restrict src,
                        __global volatile uint * restrict shared_mem,
                        const uint iterations )
{
    int base_offset;

    for (uint gID = 0; gID < iterations; gID++)
    {
        // Assume each block of memory is 256 words
        uint lID = 0x0ff & gID;

        if(lID == 0)
        {
            base_offset = read_channel_intel(req);
        }

        shared_mem[base_offset + lID] = src[gID];

        // Make sure all memory operations are committed before
```

```
        // sending token to the consumer
        mem_fence(CLK_GLOBAL_MEM_FENCE | CLK_CHANNEL_MEM_FENCE);

        if (lID == 255)
        {
            write_channel_intel(c, base_offset);
        }
    }
}
```

In this kernel, because the following lines of code are independent, the Intel FPGA SDK for OpenCL Offline Compiler can schedule them to execute concurrently:

```
shared_mem[base_offset + lID] = src[gID];
```

and

```
write_channel_intel(c, base_offset);
```

Writing data to `base_offset` and then writing `base_offset` to a channel might be much faster than writing data to global memory. The `consumer` kernel might then read `base_offset` from the channel and use it as an index to read from global memory. Without synchronization, `consumer` might read data from `producer` before `shared_mem[base_offset + lID] = src[gID];` finishes executing. As a result, `consumer` reads in invalid data. To avoid this scenario, the synchronization token must occur after the `producer` kernel commits data to memory. In other words, a `consumer` kernel cannot consume data from the `producer` kernel until `producer` stores its data in global memory successfully.

To preserve this ordering, include an OpenCL `mem_fence` token in your kernels. The `mem_fence` construct takes two flags: `CLK_GLOBAL_MEM_FENCE` and `CLK_CHANNEL_MEM_FENCE`. The `mem_fence` effectively creates a control flow dependence between operations that occur before and after the `mem_fence` call. The `CLK_GLOBAL_MEM_FENCE` flag indicates that global memory operations must obey the control flow. The `CLK_CHANNEL_MEM_FENCE` indicates that channel operations must obey the control flow. As a result, the `write_channel_intel` call in the example cannot start until the global memory operation is committed to the shared memory buffer.

1.6.4.5.7 Implementing Buffered Channels Using the depth Channels Attribute

You may have buffered or unbuffered channels in your kernel program. If there are imbalances in channel read and write operations, create buffered channels to prevent kernel stalls by including the `depth` attribute in your channel declaration. Buffered channels decouple the operation of concurrent work-items executing in different kernels.

You may use a buffered channel to control data traffic, such as limiting throughput or synchronizing accesses to shared memory. In an unbuffered channel, a write operation cannot proceed until the read operation reads a data value. In a buffered channel, a write operation cannot proceed until the data value is copied to the buffer. If the buffer is full, the operation cannot proceed until the read operation reads a piece of data and removes it from the channel.

- If you expect any temporary mismatch between the consumption rate and the production rate to the channel, set the buffer size using the `depth` channel attribute.

  The following example demonstrates the use of the `depth` channel attribute in kernel code that implements the Intel FPGA SDK for OpenCL channels extension. The `depth(N)` attribute specifies the minimum depth of a buffered channel, where *N* is the number of data values.

```
channel int c __attribute__((depth(10)));

__kernel void producer( __global int * in_data )
{
    for(int i=0; i < N; i++)
    {
        if(in_data[i])
        {
            write_channel_intel(c, in_data[i]);
        }
    }
}

__kernel void consumer( __global int * restrict check_data,
                        __global int * restrict out_data )
{
    int last_val = 0;

    for(int i=0; i< N, i++)
    {
        if(check_data[i])
        {
            last_val = read_channel_intel(c);
        }
        out_data[i] = last_val;
    }
}
```

  In this example, the write operation can write ten data values to the channel without blocking. Once the channel is full, the write operation cannot proceed until an associated read operation to the channel occurs.

  Because the channel read and write calls are conditional statements, the channel might experience an imbalance between read and write calls. You may add a buffer capacity to the channel to ensure that the `producer` and `consumer` kernels are decoupled. This step is particularly important if the `producer` kernel is writing data to the channel when the `consumer` kernel is not reading from it.

1.6.4.5.8 Enforcing the Order of Channel Calls

To enforce the order of channel calls, introduce memory fence or barrier functions in your kernel program to control memory accesses. A memory fence function is necessary to create a control flow dependence between the channel synchronization calls before and after the fence.

When the Intel FPGA SDK for OpenCL Offline Compiler generates a compute unit, it does not create instruction-level parallelism on all instructions that are independent of each other. As a result, channel read and write operations might not execute independently of each other even if there is no control or data dependence between them. When channel calls interact with each other, or when channels write data to external devices, deadlocks might occur.

For example, the code snippet below consists of a `producer` kernel and a `consumer` kernel. Channels `c0` and `c1` are unbuffered channels. The schedule of the channel read operations from `c0` and `c1` might occur in the reversed order as the channel write operations to `c0` and `c1`. That is, the `producer` kernel writes to `c0` but the `consumer` kernel might read from `c1` first. This rescheduling of channel calls might cause a deadlock because the `consumer` kernel is reading from an empty channel.

```
__kernel void producer( __global const uint * src,
                        const uint iterations )
{
    for(int i=0; i < iterations; i++)
    {
        write_channel_intel(c0, src[2*i]);
        write_channel_intel(c1, src[2*i+1]);
    }
}

__kernel void consumer( __global uint * dst,
                        const uint iterations )
{
    for(int i=0; i < iterations; i++)
    {
        /*During compilation, the AOC might reorder the way the consumer kernel
        writes to memory to optimize memory access. Therefore, c1 might be read
        before c0, which is the reverse of what appears in code.*/

        dst[2*i+1] = read_channel_intel(c0);
        dst[2*i] = read_channel_intel(c1);
    }
}
```

- To prevent deadlocks from occurring by enforcing the order of channel calls, include memory fence functions (`mem_fence`) in your kernel.

  Inserting the `mem_fence` call with each kernel's channel flag forces the sequential ordering of the write and read channel calls. The code snippet below shows the modified `producer` and `consumer` kernels:

  ```
  channel uint c0 __attribute__((depth(0)));
  channel uint c1 __attribute__((depth(0)));

  __kernel void producer( __global const uint * src,
                          const uint iterations )
  {
      for(int i=0; i < iterations; i++)
      {
          write_channel_intel(c0, src[2*i]);
          mem_fence(CLK_CHANNEL_MEM_FENCE);
          write_channel_intel(c1, src[2*i+1]);
      }
  }

  __kernel void consumer( __global uint * dst;
                          const uint iterations )
  {
      for(int i=0; i < iterations; i++)
      {
          dst[2*i+1] = read_channel_intel(c0);
          mem_fence(CLK_CHANNEL_MEM_FENCE);
          dst[2*i] = read_channel_intel(c1);
      }
  }
  ```

  In this example, `mem_fence` in the `producer` kernel ensures that the channel write operation to `c0` occurs before that to `c1`. Similarly, `mem_fence` in the `consumer` kernel ensures that the channel read operation from `c0` occurs before that from `c1`.

Defining Memory Consistency Across Kernels When Using Channels

According to the OpenCL Specification version 1.0, memory behavior is undefined unless a kernel completes execution. A kernel must finish executing before other kernels can visualize any changes in memory behavior. However, kernels that use channels can share data through common global memory buffers and synchronized memory accesses. To ensure that data written to a channel is visible to the read channel after execution passes a memory fence, define memory consistency across kernels with respect to memory fences.

- To create a control flow dependency between the channel synchronization calls and the memory operations, add the `CLK_GLOBAL_MEM_FENCE` flag to the `mem_fence` call.

  For example:

```
__kernel void producer( __global const uint * src,
                        const uint iterations )
{
    for(int i=0; i < iterations; i++)
    {
        write_channel_intel(c0, src[2*i]);
        mem_fence(CLK_CHANNEL_MEM_FENCE | CLK_GLOBAL_MEM_FENCE);
        write_channel_intel(c1, src[2*i+1]);
    }
}
```

  In this kernel, the `mem_fence` function ensures that the write operation to `c0` and memory access to `src[2*i]` occur before the write operation to `c1` and memory access to `src[2*i+1]`. This allows data written to `c0` to be visible to the read channel before data is written to `c1`.

1.6.5 Implementing OpenCL Pipes

The Intel FPGA SDK for OpenCL provides preliminary support for OpenCL pipe functions.
OpenCL pipes are part of the OpenCL Specification version 2.0. They provide a mechanism for passing data to kernels and synchronizing kernels with high efficiency and low latency.

Implement pipes if it is important that your OpenCL kernel is compatible with other SDKs.

Refer to the *OpenCL Specification version 2.0* for OpenCL C programming language specification and general information about pipes.

The Intel FPGA SDK for OpenCL implementation of pipes does not encompass the entire pipes specification. As such, it is not fully conformant to the OpenCL Specification version 2.0. The goal of the SDK's pipes implementation is to provide a solution that works seamlessly on a different OpenCL 2.0-conformant device. To enable pipes for Intel FPGA products, your design must satisfy certain additional requirements.

Related Links

OpenCL Specification version 2.0 (API)

1.6.5.1 Overview of the OpenCL Pipe Functions

OpenCL pipes allow kernels to communicate directly with each other via FIFO buffers.

Figure 10. Overview of a Pipe Network Implementation
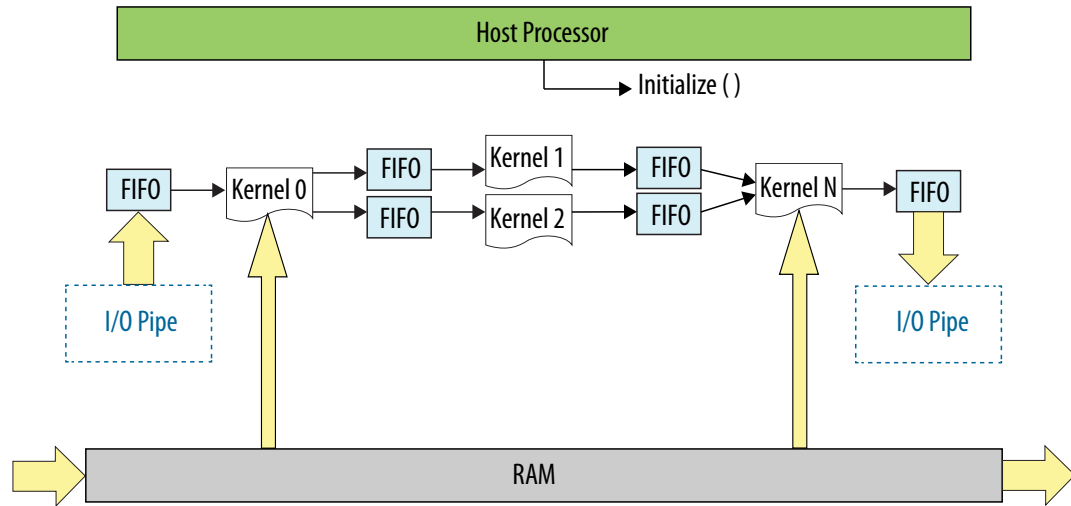
Implementation of pipes decouples kernel execution from the host processor. The foundation of the Intel FPGA SDK for OpenCL pipes support is the SDK's channels extension. However, the syntax for pipe functions differs from the channels syntax.

Important: Unlike channels, pipes have a default nonblocking behavior.

For more information on blocking and nonblocking functions, refer to the corresponding documentation on channels.

Related Links

- Implementing Blocking Channel Write Extensions on page 32
    The `write_channel_intel` API call allows you to send data across a channel.

- Implementing Nonblocking Channel Write Extensions on page 33
    Perform nonblocking channel writes to facilitate applications where data write operations might not occur.

- Implementing Nonblocking Channel Read Extensions on page 34
    Perform nonblocking reads to facilitate applications where data is not always available.

- Implementing Blocking Channel Read Extensions on page 34
    The `read_channel_intel` API call allows you to receive data across a channel.

1.6.5.2 Pipe Data Behavior

Data written to a pipe remains in a pipe as long as the kernel program remains loaded on the FPGA device. In other words, data written to a pipe persists across multiple work-groups and NDRange invocations. However, data is not persistent across multiple or different invocations of kernel programs.

Consider the following code example:

```
__kernel void
producer (write_only pipe uint __attribute__((blocking)) c0)
{
    for (uint i=0;i<10;i++)
    {
        write_pipe( c0, &i );
    }
}

__kernel void
consumer (__global uint * restrict dst,
          read_only pipe uint __attribute__((blocking))
                              __attribute__((depth(10))) c0)
{
    for (int i=0;i<5;i++)
    {
        read_pipe( c0, &dst[i] );
    }
}
```

A read operation to a pipe reads the *least* recent piece of data written to the pipe first. Pipes data maintains their FIFO ordering within the pipe.

Figure 11.   Pipe Data FIFO Ordering

The kernel `producer` writes ten elements ([0, 9]) to the pipe. The kernel `consumer` reads five elements from the pipe per NDRange invocation. During the first invocation, the kernel `consumer` reads values 0 to 4 from the pipe. Because the data persists across NDRange invocations, the second time you execute the kernel `consumer`, it reads values 5 to 9.

For this example, to avoid a deadlock from occurring, you need to invoke the kernel `consumer` twice for every invocation of the kernel `producer`. If you call `consumer` less than twice, `producer` stalls because the pipe becomes full. If you call `consumer` more than twice, `consumer` stalls because there is insufficient data in the pipe.

1.6.5.3 Multiple Work-Item Ordering for Pipes

The OpenCL specification does not define a work-item ordering. The Intel FPGA SDK for OpenCL enforces a work-item order to maintain the consistency in pipe read and write operations.

Multiple work-item accesses to a pipe can be useful in some scenarios. For example, they are useful when data words in the pipe are independent, or when the pipe is implemented for control logic. The main concern regarding multiple work-item accesses to a pipe is the order in which the kernel writes data to and reads data from the pipe. If possible, the OpenCL pipes process work-items read and write operations to a pipe in a deterministic order. As such, the read and write operations remain consistent across kernel invocations.

Requirements for Deterministic Multiple Work-Item Ordering

To guarantee deterministic ordering, the SDK checks that the pipe call is work-item invariant based on the following characteristics:

- All paths through the kernel must execute the pipe call.

- If the first requirement is not satisfied, none of the branch conditions that reach the pipe call should execute in a work-item-dependent manner.

If the SDK cannot guarantee deterministic ordering of multiple work-item accesses to a pipe, it warns you that the pipes might not have well-defined ordering with nondeterministic execution. Primarily, the SDK fails to provide deterministic ordering if you have work-item-variant code on loop executions with pipe calls, as illustrated below:

```
__kernel void
ordering (__global int * check, global int * data,
          write_only pipe int __attribute__((blocking)) req)
{
    int condition = check[get_global_id(0)];

    if (condition)
    {
        for (int i=0;i<N;i++)
        {
            process(data);
            write_pipe( req, &data[i] );
        }
    }
    else
    {
        process(data);
    }
}
```

1.6.5.3.1 Work-Item Serial Execution of Pipes

Work-item serial execution refers to an ordered execution behavior where work-item sequential IDs determine their execution order in the compute unit.

When you implement pipes in a kernel, the Intel FPGA SDK for OpenCL Offline Compiler enforces that kernel behavior is equivalent to having at most one work-group in flight. The offline compiler also ensures that the kernel executes pipes in work-item serial execution, where the kernel executes work-items with smaller IDs first. A work-item has the identifier $(x, y, z, group)$, where x, y, z are the local 3D identifiers, and $group$ is the work-group identifier.

The work-item ID $(x0, y0, z0, group0)$ is considered to be smaller than the ID $(x1, y1, z1, group1)$ if one of the following conditions is true:

- $group0 < group1$

- $group0 = group1$ and $z0 < z1$

- $group0 = group1$ and $z0 = z1$ and $y0 < y1$

- $group0 = group1$ and $z0 = z1$ and $y0 = y1$ and $x0 < x1$

Work-items with incremental IDs execute in a sequential order. For example, the work-item with an ID (`x0, y0, z0, group0`) executes the write channel call first. Then, the work-item with an ID (`x1, y0, z0, group0`) executes the call, and so on. Defining this order ensures that the system is verifiable with external models.

Pipe Execution in Loop with Multiple Work-Items

When pipes exist in the body of a loop with multiple work-items, as shown below, each loop iteration executes prior to subsequent iterations. This implies that loop iteration 0 of each work-item in a work-group executes before iteration 1 of each work-item in a work-group, and so on.

```
__kernel void
ordering (__global int * data,
          write_only pipe int __attribute__((blocking)) req)
{
    write_pipe( req, &data[get_global_id(0)] );
}
```

1.6.5.4 Restrictions in OpenCL Pipes Implementation

There are certain design restrictions to the implementation of pipes in your OpenCL application.

Default Behavior

By default, pipes exhibit nonblocking behavior. If you want the pipes in your kernel to exhibit blocking behavior, specify the blocking attribute (`__attribute__((blocking))`) when you declare the read and write pipes.

Emulation Support

The Intel FPGA SDK for OpenCL Emulator supports emulation of kernels that contain pipes. The level of Emulator support aligns with the subset of OpenCL pipes support that is implemented for the FPGA hardware.

Pipes API Support

Currently, the SDK's implementation of pipes does not support all the built-in pipe functions in the OpenCL Specification version 2.0. For a list of supported and unsupported pipe APIs, refer to *OpenCL 2.0 C Programming Language Restrictions for Pipes*.

Single Call Site

Because the pipe read and write operations do not function deterministically, for a given kernel, you can only assign one call site per pipe ID. For example, the Intel FPGA SDK for OpenCL Offline Compiler cannot compile the following code example:

```
read_pipe(pipe1, &in_data1);
read_pipe(pipe2, &in_data2);
read_pipe(pipe1, &in_data3);
```

The second `read_pipe` call to `pipe1` causes compilation failure because it creates a second call site to `pipe1`.

To gather multiple data from a given pipe, divide the pipe into multiple pipes, as shown below:

```
read_pipe(pipe1, &in_data1);
read_pipe(pipe2, &in_data2);
read_pipe(pipe3, &in_data3);
```

Because you can only assign a single call site per pipe ID, you cannot unroll loops containing pipes. Consider the following code:

```
#pragma unroll 4
for (int i=0; i < 4; i++)
{
    read_pipe(pipe1, &in_data1);
}
```

The offline compiler issues the following warning message during compilation:

```
Compiler Warning: Unroll is required but the loop cannot be
unrolled.
```

Feedback and Feed-Forward Pipes

Pipes within a kernel can be either `read_only` or `write_only`. Performance of a kernel that reads and writes to the same pipe is poor.

Kernel Vectorization Support

You cannot vectorize kernels that use pipes; that is, do not include the `num_simd_work_items` kernel attribute in your kernel code. Vectorizing a kernel that uses pipes creates multiple pipe masters and requires arbitration, which OpenCL pipes specification does not support.

Instruction-Level Parallelism on read_pipe and write_pipe Calls

If no data dependencies exist between `read_pipe` and `write_pipe` calls, the offline compiler attempts to execute these instructions in parallel. As a result, the offline compiler might execute these `read_pipe` and `write_pipe` calls in an order that does not follow the sequence expressed in the OpenCL kernel code.

Consider the following code sequence:

```
in_data1 = read_pipe(pipe1);
in_data2 = read_pipe(pipe2);
in_data3 = read_pipe(pipe3);
```

Because there are no data dependencies between the `read_pipe` calls, the offline compiler can execute them in any order.

Related Links

OpenCL 2.0 C Programming Language Restrictions for Pipes on page 156
   The Intel FPGA SDK for OpenCL offers preliminary support of OpenCL pipes.

1.6.5.5 Enabling OpenCL Pipes for Kernels

To implement pipes, modify your OpenCL kernels to include pipes-specific API calls.

Pipes declarations are unique within a given OpenCL kernel program. Also, pipe instances are unique for every OpenCL kernel program-device pair. If the runtime loads a single OpenCL kernel program onto multiple devices, each device will have a single copy of each pipe. However, these pipe copies are independent and do not share data across the devices.

1.6.5.5.1 Ensuring Compatibility with Other OpenCL SDKs

Currently, Intel's implementation of OpenCL pipes is partially conformant to the OpenCL Specification version 2.0. If you port a kernel that implements pipes from another OpenCL SDK to the Intel FPGA SDK for OpenCL, you must modify the host code and the kernel code. The modifications do not affect subsequent portability of your application to other OpenCL SDKs.

Host Code Modification

Below is an example of a modified host application:

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include "CL/opencl.h"
#define SIZE 1000

const char *kernel_source = "__kernel void pipe_writer(__global int *in,"
                            "                          write_only pipe int p_in)
\n"
                            "{\n"
                            "    int gid = get_global_id(0);\n"
                            "    write_pipe(p_in, &in[gid]);\n"
                            "}\n"
                            "__kernel void pipe_reader(__global int *out,"
                            "                          read_only pipe int p_out)
\n"
                            "{\n"
                            "    int gid = get_global_id(0);\n"
                            "    read_pipe(p_out, &out[gid]);\n"
                            "}\n";

int main()
{
    int *input = (int *)malloc(sizeof(int) * SIZE);
    int *output = (int *)malloc(sizeof(int) * SIZE);
    memset(output, 0, sizeof(int) * SIZE);
    for (int i = 0; i != SIZE; ++i)
    {
        input[i] = rand();
    }

    cl_int status;
    cl_platform_id platform;
    cl_uint num_platforms;
    status = clGetPlatformIDs(1, &platform, &num_platforms);

    cl_device_id device;
    cl_uint num_devices;
    status = clGetDeviceIDs(platform,
                            CL_DEVICE_TYPE_ALL,
                            1,
                            &device,
                            &num_devices);

    cl_context context = clCreateContext(0, 1, &device, NULL, NULL, &status);

    cl_command_queue queue = clCreateCommandQueue(context, device, 0, &status);
```

```
        size_t len = strlen(kernel_source);
        cl_program program = clCreateProgramWithSource(context,
                                                       1,
                                                       (const char **)&kernel_source,
                                                       &len,
                                                       &status);

        status = clBuildProgram(program, num_devices, &device, "", NULL, NULL);

        cl_kernel pipe_writer = clCreateKernel(program, "pipe_writer", &status);
        cl_kernel pipe_reader = clCreateKernel(program, "pipe_reader", &status);

        cl_mem in_buffer = clCreateBuffer(context,
                                          CL_MEM_READ_ONLY | CL_MEM_COPY_HOST_PTR,
                                          sizeof(int) * SIZE,
                                          input,
                                          &status);
        cl_mem out_buffer = clCreateBuffer(context,
                                           CL_MEM_WRITE_ONLY,
                                           sizeof(int) * SIZE,
                                           NULL,
                                           &status);

        cl_mem pipe = clCreatePipe(context, 0, sizeof(cl_int), SIZE, NULL, &status);

        status = clSetKernelArg(pipe_writer, 0, sizeof(cl_mem), &in_buffer);
        status = clSetKernelArg(pipe_writer, 1, sizeof(cl_mem), &pipe);
        status = clSetKernelArg(pipe_reader, 0, sizeof(cl_mem), &out_buffer);
        status = clSetKernelArg(pipe_reader, 1, sizeof(cl_mem), &pipe);

        size_t size = SIZE;
        cl_event sync;
        status = clEnqueueNDRangeKernel(queue,
                                        pipe_writer,
                                        1,
                                        NULL,
                                        &size,
                                        &size,
                                        0,
                                        NULL,
                                        &sync);
        status = clEnqueueNDRangeKernel(queue,
                                        pipe_reader,
                                        1,
                                        NULL,
                                        &size,
                                        &size,
                                        1,
                                        &sync,
                                        NULL);
        status = clFinish(queue);

        status = clEnqueueReadBuffer(queue,
                                     out_buffer,
                                     CL_TRUE,
                                     0,
                                     sizeof(int) * SIZE,
                                     output,
                                     0,
                                     NULL,
                                     NULL);

        int golden = 0, result = 0;
        for (int i = 0; i != SIZE; ++i)
        {
          golden += input[i];
          result += output[i];
        }

        int ret = 0;
        if (golden != result)
```

```
    {
        printf("FAILED!");
        ret = 1;
    } else
    {
        printf("PASSED!");
    }
    printf("\n");

    return ret;
}
```

Kernel Code Modification

If your kernel code runs on OpenCL SDKs that conforms to the OpenCL Specification version 2.0, you must modify it before running it on the Intel FPGA SDK for OpenCL. To modify the kernel code, perform the following modifications:

- Rename the pipe arguments so that they are the same in both kernels. For example, rename `p_in` and `p_out` to `p`.

- Specify the `depth` attribute for the pipe arguments. Assign a `depth` attribute value that equals to the maximum number of packets that the pipe creates to hold in the host.

- Execute the kernel program in the offline compilation mode because the Intel FPGA SDK for OpenCL has an offline compiler.

The modified kernel code appears as follows:

```
#define SIZE 1000

__kernel void pipe_writer(__global int *in,
                          write_only pipe int __attribute__((depth(SIZE))) p)
{
    int gid = get_global_id(0);
    write_pipe(p, &in[gid]);
}

__kernel void pipe_reader(__global int *out,
                          read_only pipe int __attribute__((depth(SIZE))) p)
{
    int gid = get_global_id(0);
    read_pipe(p, &out[gid]);
}
```

1.6.5.5.2 Declaring the Pipe Handle

Use the `pipe` variable to define the static pipe connectivity between kernels or between kernels and I/O.

To read from and write to a pipe, the kernel must pass the pipe variable to each of the corresponding API call.

- Declare the pipe handle as a file scope variable in the kernel source code in the following convention: *<access qualifier>* pipe *<type> <variable_name>*

  The *<type>* of the pipe may be any OpenCL built-in scalar or vector data type with a scalar size of 1024 bits or less. It may also be any user-defined type that is comprised of scalar or vector data type with a scalar size of 1024 bits or less.

  Consider the following pipe handle declarations:

  `__kernel void first (pipe int c)`

  `__kernel void second (write_only pipe int c)`

  The first example declares a read-only pipe handle of type `int` in the kernel `first`. The second example declares a write-only pipe in the kernel `second`. The kernel `first` may only read from pipe `c`, and the kernel `second` may only write to pipe `c`.

  Important: The Intel FPGA SDK for OpenCL Offline Compiler statically infers the connectivity of pipes in your system by matching the names of the pipe arguments. In the example above, the kernel `first` is connected to the kernel `second` by the pipe `c`.

In an Intel OpenCL system, only one kernel may read to a pipe. Similarly, only one kernel may write to a pipe. If a non-I/O pipe does not have at least one corresponding reading operation and one writing operation, the offline compiler issues an error.

For more information in the Intel FPGA SDK for OpenCL I/O pipe implementation, refer to *Implementing I/O Pipes Using the io Attribute*.

Related Links

Implementing I/O Pipes Using the io Attribute on page 57
> Include an `io` attribute in your OpenCL pipe declaration to declare a special I/O pipe to interface with input or output features of an FPGA board.

1.6.5.5.3 Implementing Pipe Writes

The `write_pipe` API call allows you to send data across a pipe.

Intel only supports the convenience version of the `write_pipe` function. By default, `write_pipe` calls are nonblocking. Pipe write operations are successful only if there is capacity in the pipe to hold the incoming packet.

Attention: The write pipe calls support single-call sites only. For a given pipe, only one write pipe call to it can exist in the entire kernel program.

- To implement a pipe write, include the following `write_pipe` function signature:

  `int write_pipe (write_only pipe <type> pipe_id, const <type> *data);`

  Where:

  `pipe_id` identifies the buffer to which the pipe connects, and it must match the `pipe_id` of the corresponding read pipe (`read_pipe`).

  `data` is the data that the pipe write operation writes to the pipe. It is a pointer to the packet type of the pipe. Note that writing to the pipe might lead to a global or local memory load, depending on the source address space of the data pointer.

  <type> defines a pipe data width. The return value indicates whether the pipe write operation is successful. If successful, the return value is 0. If pipe write is unsuccessful, the return value is -1.

The following code snippet demonstrates the implementation of the `write_pipe` API call:

```
/*Declares the writable nonblocking pipe, p, which contains packets of type int*/
__kernel void kernel_write_pipe (__global const long *src,
                                 write_only pipe int p)
{
    for (int i=0; i < N; i++)
    {
        //Performs the actual writing
        //Emulates blocking behavior via the use of a while loop
        while (write_pipe(p, &src[i]) < 0) { }
    }
}
```

The `while` loop is unnecessary if you specify a *blocking attribute*. To facilitate better hardware implementations, Intel provides facility for blocking `write_pipe` calls by specifying the blocking attribute (that is, `__attribute__((blocking))`) on the pipe arugment declaration for the kernel. Blocking `write_pipe` calls always return success.

Caution: When you send data across a blocking write pipe using the `write_pipe` API call, keep in mind that if the pipe is full (that is, if the FIFO buffer is full of data), your kernel will stall. Use the Intel FPGA SDK for OpenCL Profiler to check for pipe stalls.

Related Links

Profiling Your OpenCL Kernel on page 112
    The Intel FPGA SDK for OpenCL Profiler measures and reports performance data collected during OpenCL kernel execution on the FPGA.

1.6.5.5.4 Implementing Pipe Reads

The `read_pipe` API call allows you to receive data across a pipe.

Intel only supports the convenience version of the `read_pipe` function. By default, `read_pipe` calls are nonblocking.

Note: The read pipe calls support single-call sites only. For a given pipe, only one read pipe call to it can exist in the entire kernel program.

- To implement a pipe read, include the following `read_pipe` function signature:

  `int read_pipe (read_only_pipe <type> pipe_id, <type> *data);`

  Where:

  `pipe_id` identifies the buffer to which the pipe connects, and it must match the `pipe_id` of the corresponding pipe write operation (`write_pipe`).

  `data` is the data that the pipe read operation reads from the pipe. It is a pointer to the location of the data. Note that `write_pipe` call might lead to a global or local memory load, depending on the source address space of the data pointer.

  <type> defines the packet size of the data.

The following code snippet demonstrates the implementation of the `read_pipe` API call:

```
/*Declares the read_only_pipe that contains packets
of type long.*/
/*Declares that read_pipe calls within the kernel will exhibit
blocking behavior*/
__kernel void kernel_read_pipe (__global long *dst,
                                read_only pipe long __attribute__((blocking)) p)
{
    for (int i=0; i < N; i++)
    {
        /*Reads from a long from the pipe and stores it
        into global memory at the specified location*/
        read_pipe(p, &dst[i]);
    }
}
```

To facilitate better hardware implementations, Intel provides facility for blocking `write_pipe` calls by specifying the blocking attribute (that is, `__attribute__((blocking))`) on the pipe arugment declaration for the kernel. Blocking `write_pipe` calls always return success.

Caution: If the pipe is empty (that is, if the FIFO buffer is empty), you cannot receive data across a blocking read pipe using the `read_pipe` API call. Doing so causes your kernel to stall.

1.6.5.5.5 Implementing Buffered Pipes Using the depth Attribute

You may have buffered or unbuffered pipes in your kernel program. If there are imbalances in pipe read and write operations, create buffered pipes to prevent kernel stalls by including the `depth` attribute in your pipe declaration. Buffered pipes decouple the operation of concurrent work-items executing in different kernels.

You may use a buffered pipe to control data traffic, such as limiting throughput or synchronizing accesses to shared memory. In an unbuffered pipe, a write operation can only proceed when the read operation is expecting to read data. Use unbuffered pipes in conjunction with blocking read and write behaviors in kernels that execute concurrently. The unbuffered pipes provide self-synchronizing data transfers efficiently.

In a buffered pipe, a write operation can only proceed if there is capacity in the pipe to hold the incoming packet. A read operation can only proceed if there is at least one packet in the pipe.

Use buffered pipes if pipe calls are predicated differently in the writer and reader kernels, and the kernels do not execute concurrently.

- If you expect any temporary mismatch between the consumption rate and the production rate to the pipe, set the buffer size using the `depth` attribute.

The following example demonstrates the use of the `depth` attribute in kernel code that implements the OpenCL pipes. The `depth(N)` attribute specifies the minimum depth of a buffered pipe, where *N* is the number of data values. If the read and write kernels specify different depths for a given buffered pipe, the Intel FPGA SDK for OpenCL Offline Compiler will use the larger depth of the two.

```
__kernel void
producer (__global int *in_data,
          write_only pipe int __attribute__((blocking))
                               __attribute__((depth(10))) c)
{
    for (i=0; i < N; i++)
    {
        if (in_data[i])
        {
            write_pipe( c, &in_data[i] );
        }
    }
}

__kernel void
consumer (__global int *check_data,
          __global int *out_data,
          read_only pipe int __attribute__((blocking)) c )
{
    int last_val = 0;
    for (i=0; i < N; i++)
    {
        if (check_data[i])
        {
            read_pipe( c, &last_val );
        }
        out_data[i] = last_val;
    }
}
```

In this example, the write operation can write ten data values to the pipe successfully. After the pipe is full, the write kernel returns failure until a read kernel consumes some of the data in the pipe.

Because the pipe read and write calls are conditional statements, the pipe might experience an imbalance between read and write calls. You may add a buffer capacity to the pipe to ensure that the `producer` and `consumer` kernels are decoupled. This step is particularly important if the `producer` kernel is writing data to the pipe when the `consumer` kernel is not reading from it.

1.6.5.5.6 Implementing I/O Pipes Using the io Attribute

Include an `io` attribute in your OpenCL pipe declaration to declare a special I/O pipe to interface with input or output features of an FPGA board.
These features might include network interfaces, PCIe, cameras, or other data capture or processing devices or protocols.

In the Intel FPGA SDK for OpenCL channels extension, the `io("chan_id")` attribute specifies the I/O feature of an accelerator board with which a channel interfaces. The *chan_id* argument is the name of the I/O interface listed in the `board_spec.xml` file of your Custom Platform. The same I/O features can be used to identify I/O pipes.

Because peripheral interface usage might differ for each device type, consult your board vendor's documentation when you implement I/O pipes in your kernel program. Your OpenCL kernel code must be compatible with the type of data generated by the peripheral interfaces. If there is a difference in the byte ordering between the external I/O pipes and the kernel, the Intel FPGA SDK for OpenCL Offline Compiler converts the byte ordering seamlessly upon entry and exit.

Caution:
- Implicit data dependencies might exist for pipes that connect to the board directly and communicate with peripheral devices via I/O pipes. These implicit data dependencies might lead to compilation issues because the offline compiler cannot identify these dependencies.

- External I/O pipes communicating with the same peripherals do not obey any sequential ordering. Ensure that the external device does not require sequential ordering because unexpected behavior might occur.

1. Consult the `board_spec.xml` file in your Custom Platform to identify the input and output features available on your FPGA board.

   For example, a `board_spec.xml` file might include the following information on I/O features:

   ```
   <channels>
     <interface name="udp_0" port="udp0_out"  type="streamsource" width="256"
      chan_id="eth0_in"/>
     <interface name="udp_0" port="udp0_in"  type="streamsink" width="256"
      chan_id="eth0_out"/>
     <interface name="udp_0" port="udp1_out"  type="streamsource" width="256"
      chan_id="eth1_in"/>
     <interface name="udp_0" port="udp1_in"  type="streamsink" width="256"
      chan_id="eth1_out"/>
   </channels>
   ```

   The `width` attribute of an `interface` element specifies the width, in bits, of the data type used by that pipe. For the example above, both the `uint` and `float` data types are 32 bits wide. Other bigger or vectorized data types must match the appropriate bit width specified in the `board_spec.xml` file.

2. Implement the `io` attribute as demonstrated in the following code example. The `io` attribute names must match those of the I/O channels (`chan_id`) specified in the `board_spec.xml` file.

   ```
   __kernel void test (pipe uint pkt    __attribute__((io("enet"))),;
                       pipe float data __attribute__((io("pcie"))));
   ```

   Attention: Declare a unique `io("chan_id")` handle for each I/O pipe specified in the channels XML element within the `board_spec.xml` file.

1.6.5.5.7 Enforcing the Order of Pipe Calls

To enforce the order of pipe calls, introduce memory fence or barrier functions in your kernel program to control memory accesses. A memory fence function is necessary to create a control flow dependence between the pipe synchronization calls before and after the fence.

When the Intel FPGA SDK for OpenCL Offline Compiler generates a compute unit, it does not create instruction-level parallelism on all instructions that are independent of each other. As a result, pipe read and write operations might not execute independently of each other even if there is no control or data dependence between them. When pipe calls interact with each other, or when pipes write data to external devices, deadlocks might occur.

For example, the code snippet below consists of a `producer` kernel and a `consumer` kernel. Pipes `c0` and `c1` are unbuffered pipes. The schedule of the pipe read operations from `c0` and `c1` might occur in the reversed order as the pipe write operations to `c0` and `c1`. That is, the `producer` kernel writes to `c0` but the `consumer` kernel might read from `c1` first. This rescheduling of pipe calls might cause a deadlock because the `consumer` kernel is reading from an empty pipe.

```
__kernel void
producer (__global const uint * restrict src,
          const uint iterations,
          write_only pipe uint __attribute__((blocking)) c0,
          write_only pipe uint __attribute__((blocking)) c1)
{
    for (int i=0; i < iterations; i++) {
        write_pipe( c0, &src[2*i  ] );
        write_pipe( c1, &src[2*i+1] ); }
}

__kernel void
consumer (__global uint * restrict dst,
          const uint iterations,
          read_only pipe uint __attribute__((blocking)) c0,
          read_only pipe uint __attribute__((blocking)) c1)
{
    for (int i=0; i < iterations; i++) {
        read_pipe( c0, &dst[2*i+1] );
        read_pipe( c1, &dst[2*i] ); }
}
```

- To prevent deadlocks from occurring by enforcing the order of pipe calls, include memory fence functions (`mem_fence`) in your kernel.

  Inserting the `mem_fence` call with each kernel's pipe flag forces the sequential ordering of the write and read pipe calls. The code snippet below shows the modified `producer` and `consumer` kernels:

  ```
  __kernel void
  producer (__global const uint * src,
            const uint iterations,
            write_only_pipe uint __attribute__((blocking)) c0,
            write_only_pipe uint __attribute__((blocking)) c1)
  {
      for(int i=0; i < iterations; i++)
      {
          write_pipe(c0, &src[2*i  ]);
          mem_fence(CLK_CHANNEL_MEM_FENCE);
          write_pipe(c1, &src[2*i+1]);
      }
  }

  __kernel void
  consumer (__global uint * dst;
            const uint iterations,
            read_only_pipe uint __attribute__((blocking)) c0,
            read_only_pipe uint __attribute__((blocking)) c1)
  {
      for(int i=0; i < iterations; i++)
      {
          read_pipe(c0, &dst[2*i  ]);
          mem_fence(CLK_CHANNEL_MEM_FENCE);
          read_pipe(c1, &dst[2*i+1]);
      }
  }
  ```

  In this example, `mem_fence` in the `producer` kernel ensures that the pipe write operation to `c0` occurs before that to `c1`. Similarly, `mem_fence` in the `consumer` kernel ensures that the pipe read operation from `c0` occurs before that from `c1`.

Defining Memory Consistency Across Kernels When Using Pipes

According to the OpenCL Specification version 2.0, memory behavior is undefined unless a kernel completes execution. A kernel must finish executing before other kernels can visualize any changes in memory behavior. However, kernels that use pipes can share data through common global memory buffers and synchronized memory accesses. To ensure that data written to a pipe is visible to the read pipe after execution passes a memory fence, define memory consistency across kernels with respect to memory fences.

- To create a control flow dependency between the pipe synchronization calls and the memory operations, add the `CLK_GLOBAL_MEM_FENCE` flag to the `mem_fence` call.

  For example:

```
__kernel void
producer (__global const uint * restrict src,
          const uint iterations,
          write_only pipe uint __attribute__((blocking)) c0,
          write_only pipe uint __attribute__((blocking)) c1)
{
    for (int i=0;i<iterations;i++)
    {
        write_pipe( c0, &src[2*i  ] );
        mem_fence( CLK_CHANNEL_MEM_FENCE | CLK_GLOBAL_MEM_FENCE );
        write_pipe( c1, &src[2*i+1] );
    }
}
```

  In this kernel, the `mem_fence` function ensures that the write operation to `c0` and memory access to `src[2*i]` occur before the write operation to `c1` and memory access to `src[2*i+1]`. This allows data written to `c0` to be visible to the read pipe before data is written to `c1`.

1.6.6 Implementing Arbitrary Precision Integers

Use the Intel FPGA SDK for OpenCL arbitrary precision integer extension to define integers with a custom bit-width. You can define integer custom bit-widths up to and including 64 bits.

To use the arbitrary precision integer extension, include the following line in your list of header files in your kernel code:

```
#include "ihc_apint.h"
```

When you compile a kernel that includes the `ihc_apint.h` header file, you must include the `-l $ALTERAOCLSDKROOT/include/kernel_headers` option with the `aoc` command. For example:

```
aoc <other command options> -l $ALTERAOCLSDKROOT/include/kernel_headers
<my_kernel_file>
```

The header enables the arbitrary precision integers extension, and has macros that define C-style declarations for signed and unsigned arbitrary precision integers as follows:

```
#define ap_int<d> intd_t
#define ap_uint<d> uintd_t
```

For example, you can declare a 10-bit signed and unsigned arbitrary precision integers as follows:

```
int10_t x_signed;
uint10_t x_unsigned;
```

You can declare arbitrary precision integers with widths up to 64 bits.

To use arbitrary precision integers without using the header files, enable the extension with the following pragma directive:

```
#pragma OPENCL EXTENSION cl_intel_arbitrary_precision_integers : enable
```

After the pragma declaration, you can declare your arbitrary precision integers as follows:

```
ap_int<d> intd_t my_signed_integer
ap_uint<d> uintd_t my_unsigned_integer
```

If you do operations where the bit width of the result is larger than the bit widths of the arguments, you must explicitly cast one of the arguments to the resulting bit width.

For example, if you had the following operation, the result overflows the declared size of the integer:

```
int10_t a;
int10_t b;
int20_t res;

res = a * b;
```

In the example, the compiler attempts to instantiate a multiplier that multiplies two 10-bit integers and put the results into another 10-bit integer. The result is then sign extended or zero extended up to 20-bits.

To prevent the overflow, explicitly cast one of the arguments to the resulting bit width as follows:

```
res = ((int20_t)a) * b
```

Remember: When you compile a program for x86-64 platforms, the bit widths for arbitrary precisions integers are rounded up to either 32 bits or 64 bits. When you compile a kernel for an FPGA platform, the bit widths are not rounded up and the arbitrary precision integers remain at their declared bit width.

As a result, an operation that appears to work correctly in an x86-64 program can overflow and lose precision when you compile that same operation in an FPGA kernel. The additional precision provided by bit-width rounding on x86-64 platforms masks possible overflow and precision-loss problems you might encounter when your compile your FPGA kernel.

1.6.7 Using Predefined Preprocessor Macros in Conditional Compilation

You may take advantage of predefined preprocessor macros that allow you to conditionally compile portions of your kernel code.
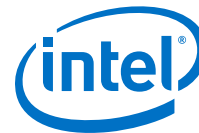
- To include device-specific (for example, FPGA_board_1) code in your kernel program, structure your kernel program in the following manner:

```
#if defined(AOCL_BOARD_FPGA_board_1)
    //FPGA_board_1-specific statements
#else
    //FPGA_board_2-specific statements
#endif
```

When you target your kernel compilation to a specific board, it sets the predefined preprocessor macro AOCL_BOARD_*<board_name>* to 1. If *<board_name>* is FPGA_board_1, the Intel FPGA SDK for OpenCL Offline Compiler will compile the FPGA_board_1-specific parameters and features.

- To introduce Intel FPGA SDK for OpenCL Offline Compiler-specific compiler features and optimizations, structure your kernel program in the following manner:

```
#if defined(ALTERA_CL)
    //statements
#else
    //statements
#endif
```

Where ALTERA_CL is the Intel predefined preprocessor macro for the offline compiler.

Related Links

Defining Preprocessor Macros to Specify Kernel Parameters (-D <macro_name>) on page 99
> The Intel FPGA SDK for OpenCL Offline Compiler supports preprocessor macros that allow you to pass macro definitions and compile code on a conditional basis.

1.6.8 Declaring __constant Address Space Qualifiers

There are several limitations and workarounds you must consider when you include __constant address space qualifiers in your kernel.

Function Scope __constant Variables

The Intel FPGA SDK for OpenCL Offline Compiler does not support function scope __constant variables. Replace function scope __constant variables with file scope constant variables. You can also replace function scope __constant variables with __constant buffers that the host passes to the kernel.

File Scope __constant Variables

If the host always passes the same constant data to your kernel, consider declaring that data as a constant preinitialized file scope array within the kernel file. Declaration of a constant preinitialized file scope array creates a ROM directly in the hardware to store the data. This ROM is available to all work-items in the NDRange.

The offline compiler supports only scalar file scope constant data. For example, you may set the `__constant` address space qualifier as follows:

```
__constant int my_array[8] = {0x0, 0x1, 0x2, 0x3, 0x4, 0x5, 0x6, 0x7};

__kernel void my_kernel (__global int * my_buffer)
{
    size_t gid = get_global_id(0);
    my_buffer[gid] += my_array[gid % 8];
}
```

In this case, the offline compiler sets the values for `my_array` in a ROM because the file scope constant data does not change between kernel invocations.

Warning: Do not set your file scope `__constant` variables in the following manner because the offline compiler does not support vector type `__constant` arrays declared at the file scope:

```
__constant int2 my_array[4] = {(0x0, 0x1), (0x2, 0x3); (0x4,
0x5), (0x6, 0x7)};
```

Pointers to __constant Parameters from the Host

You can replace file scope constant data with a pointer to a `__constant` parameter in your kernel code. You must then modify your host application in the following manner:

1. Create `cl_mem` memory objects associated with the pointers in global memory.

2. Load constant data into `cl_mem` objects with `clEnqueueWriteBuffer` prior to kernel execution.

3. Pass the `cl_mem` objects to the kernel as arguments with the `clSetKernelArg` function.

For simplicity, if a constant variable is of a complex type, use a `typedef` argument, as shown in the table below:

Table 1.   Replacing File Scope __constant Variable with Pointer to __constant Parameter

If your source code is structured as follows:	Rewrite your code to resemble the following syntax:

<pre>__constant int Payoff[2][2] = {{ 1, 3}, {5, 3}};
__kernel void original(__global int * A)
{
    *A = Payoff[1][2];
    // and so on
}</pre>
	<pre>__kernel void modified(__global int * A,
	__constant Payoff_type * PayoffPtr)
	{
	    *A = (PayoffPtr)[1][2];
	    // and so on
	}</pre>

Attention: Use the same type definition in both your host application and your kernel.

1.6.9 Including Structure Data Types as Arguments in OpenCL Kernels

Convert each structure parameter (`struct`) to a pointer that points to a structure.

The table below describes how you can convert structure parameters:

Table 2.   Converting Structure Parameters to Pointers that Point to Structures

If your source code is structured as follows:	Rewrite your code to resemble the following syntax:
| ```
struct Context
{
   float param1;
   float param2;
   int param3;
   uint param4;
};

__kernel void algorithm(__global float * A,
struct Context c)
{
   if ( c.param3 )
   {
      // statements
   }
}
``` | ```
struct Context
{
   float param1;
   float param2;
   int param3;
   uint param4;
};

__kernel void algorithm(__global float * A,
__global struct Context * restrict c)
{
   if ( c->param3 )
   {
      // Dereference through a
      // pointer and so on
   }
}
``` |

Attention: The `__global struct` declaration creates a new buffer to store the structure. To prevent pointer aliasing, include a `restrict` qualifier in the declaration of the pointer to the structure.

1.6.9.1 Matching Data Layouts of Host and Kernel Structure Data Types

If you use structure data types (`struct`) as arguments in OpenCL kernels, match the member data types and align the data members between the host application and the kernel code.

To match member data types, use the `cl_` version of the data type in your host application that corresponds to the data type in the kernel code. The `cl_` version of the data type is available in the `opencl.h` header file. For example, if you have a data member of type `float4` in your kernel code, the corresponding data member you declare in the host application is `cl_float4`.

Align the structures and align the `struct` data members between the host and kernel applications. Manage the alignments carefully because of the variability among different host compilers.

For example, if you have `float 4` OpenCL data types in the struct, the alignments of these data items must satisfy the OpenCL specification (that is, 16-byte alignment for `float4`).

The following rules apply when the Intel FPGA SDK for OpenCL Offline Compiler compiles your OpenCL kernels:

1. Alignment of built-in scalar and vector types follow the rules outlined in Section 6.1.5 of the *OpenCL Specification version 1.0*.

   The offline compiler usually aligns a data type based on its size. However, the compiler aligns a value of a three-element vector the same way it aligns a four-element vector.

2. An array has the same alignment as one of its elements.

3. A `struct` (or a `union`) has the same alignment as the maximum alignment necessary for any of its data members.

Consider the following example:

```
struct my_struct
{
    char data[3];
    float4 f4;
    int index;
};
```

The offline compiler aligns the `struct` elements above at 16-byte boundaries because of the `float4` data type. As a result, both `data` and `index` also have 16-byte alignment boundaries.

4. The offline compiler does not reorder data members of a `struct`.

5. Normally, the offline compiler inserts a minimum amount of data structure padding between data members of a `struct` to satisfy the alignment requirements for each data member.

   a. In your OpenCL kernel code, you may specify data packing (that is, no insertion of data structure padding) by applying the `packed` attribute to the `struct` declaration. If you impose data packing, ensure that the alignment of data members satisfies the OpenCL alignment requirements. The Intel FPGA SDK for OpenCL does not enforce these alignment requirements. Ensure that your host compiler respects the kernel attribute and sets the appropriate alignments.

   b. In your OpenCL kernel code, you may specify the amount of data structure padding by applying the `aligned(N)` attribute to a data member, where *N* is the amount of padding. The SDK does not enforce these alignment requirements. Ensure that your host compiler respects the kernel attribute and sets the appropriate alignments.

   For Windows systems, some versions of the Microsoft Visual Studio compiler pack structure data types by default. If you do not want to apply data packing, specify an amount of data structure padding as shown below:

```
struct my_struct
{
    __declspec(align(16)) char data[3];

    /*Note that cl_float4 is the only known float4 definition on the
host*/
    __declspec(align(16)) cl_float4 f4;

    __declspec(align(16)) int index;
};
```

   Tip: An alternative way of adding data structure padding is to insert dummy `struct` members of type `char` or array of `char`.

Related Links

- Modifying Host Program for Structure Parameter Conversion on page 81
  If you convert any structure parameters to pointers-to-constant structures in your OpenCL kernel, you must modify your host application accordingly.

- OpenCL Specification version 1.0

1.6.9.2 Disabling Insertion of Data Structure Padding

You may instruct the Intel FPGA SDK for OpenCL Offline Compiler to disable automatic padding insertion between members of a `struct` data structure.

- To disable automatic padding insertion, insert the `packed` attribute prior to the kernel source code for a `struct` data structure.

    For example:

```
struct __attribute__((packed)) Context
{
    float param1;
    float param2;
    int param3;
    uint param4;
};
__kernel void algorithm(__global float * restrict A,
                        __global struct Context * restrict c)
{
    if ( c->param3 )
    {
        // Dereference through a pointer and so on
    }
}
```

For more information, refer to the *Align a Struct with or without Padding* section of the *Intel FPGA SDK for OpenCL Best Practices Guide*.

Related Links

Align a Struct with or without Padding

1.6.9.3 Specifying the Alignment of a Struct

You may instruct the Intel FPGA SDK for OpenCL Offline Compiler to set a specific alignment of a `struct` data structure.

- To specify the struct alignment, insert the `aligned(N)` attribute prior to the kernel source code for a `struct` data structure.

    For example:

```
struct __attribute__((aligned(2))) Context
{
    float param1;
    float param2;
    int param3;
    uint param4;
};
__kernel void algorithm(__global float * A,
                        __global struct Context * restrict c)
{
    if ( c->param3 )
    {
        // Dereference through a pointer and so on
    }
}
```

For more information, refer to the *Align a Struct with or without Padding* section of the *Intel FPGA SDK for OpenCL Best Practices Guide*.

Related Links

Align a Struct with or without Padding

1.6.10 Inferring a Register

The Intel FPGA SDK for OpenCL Offline Compiler can implement data that is in the private address space in registers or in block RAMs. In general, the offline compiler chooses registers if the access to a variable is fixed and does not require any dynamic indexes. Accessing an array with a variable index usually forces the array into block RAMs. Implementing private data as registers is beneficial for data access that occurs in a single cycle (for example, feedback in a single work-item loop).

The offline compiler infers private arrays as registers either as single values or in a piecewise fashion. Piecewise implementation results in very efficient hardware; however, the offline compiler must be able to determine data accesses statically. To facilitate piecewise implementation, hardcode the access points into the array. You can also facilitate register inference by unrolling loops that access the array.

If array accesses are not inferable statically, the offline compiler might infer the array as registers. However, the offline compiler limits the size of these arrays to 64 bytes in length for single work-item kernels. There is effectively no size limit for kernels with multiple work-items.

Consider the following code example:

```
int array[SIZE];
for (int j = 0; j < N; ++j)
{
    for (int i = 0; i < SIZE - 1; ++i)
    {
        array[i] = array[i + 1];
    }
}
```

The indexing into `array[i]` is not inferable statically because the loop is not unrolled. If the size of `array[SIZE]` is less than or equal to 64 bytes for single work-item kernels, the offline compiler implements `array[SIZE]` into registers as a single value. If the size of `array[SIZE]` is greater than 64 bytes for single work-item kernels, the offline compiler implements the entire array in block RAMs. For multiple work-item kernels, the offline compiler implements `array[SIZE]` into registers as a single value as long as its size is less than 1 kilobyte (KB).

1.6.10.1 Inferring a Shift Register

The shift register design pattern is a very important design pattern for many applications. However, the implementation of a shift register design pattern might seem counterintuitive at first.

Consider the following code example:

```
channel int in, out;

#define SIZE 512
//Shift register size must be statically determinable

__kernel void foo()
{
```

```
        int shift_reg[SIZE];
           //The key is that the array size is a compile time constant

        // Initialization loop
        #pragma unroll
        for (int i=0; i < SIZE; i++)
        {
            //All elements of the array should be initialized to the same value
            shift_reg[i] = 0;
        }

        while(1)
        {
            // Fully unrolling the shifting loop produces constant accesses
            #pragma unroll
            for (int j=0; j < SIZE-1; j++)
            {
                shift_reg[j] = shift_reg[j + 1];
            }
            shift_reg[SIZE - 1] = read_channel_intel(in);

            // Using fixed access points of the shift register
            int res = (shift_reg[0] + shift_reg[1]) / 2;

            // 'out' channel will have running average of the input channel
            write_channel_intel(out, res);
        }
}
```

In each clock cycle, the kernel shifts a new value into the array. By placing this shift register into a block RAM, the Intel FPGA SDK for OpenCL Offline Compiler can efficiently handle multiple access points into the array. The shift register design pattern is ideal for implementing filters (for example, image filters like a Sobel filter or time-delay filters like a finite impulse response (FIR) filter).

When implementing a shift register in your kernel code, keep in mind the following key points:

1. Unroll the shifting loop so that it can access every element of the array.

2. All access points must have constant data accesses. For example, if you write a calculation in nested loops using multiple access points, unroll these loops to establish the constant access points.

3. Initialize all elements of the array to the same value. Alternatively, you may leave the elements uninitialized if you do not require a specific initial value.

4. If some accesses to a large array are not inferable statically, they force the offline compiler to create inefficient hardware. If these accesses are necessary, use __local memory instead of __private memory.

5. Do not shift a large shift register conditionally. The shifting must occur in very loop iteration that contains the shifting code to avoid creating inefficient hardware.

1.6.11 Enabling Double Precision Floating-Point Operations

The Intel FPGA SDK for OpenCL offers preliminary support for all double precision floating-point functions.

Before declaring any double precision floating-point data type in your OpenCL kernel, include the following OPENCL EXTENSION pragma in your kernel code:

```
#pragma OPENCL EXTENSION cl_khr_fp64 : enable
```

1.6.12 Single-Cycle Floating-Point Accumulator for Single Work-Item Kernels

Single work-item kernels that perform accumulation in a loop can leverage the single-cycle floating-point accumulator feature of the Intel FPGA SDK for OpenCL Offline Compiler. The offline compiler searches for these kernel instances and attempts to map an accumulation that executes in a loop into the accumulator structure.

The offline compiler supports an accumulator that adds or subtracts a value. To leverage this feature, describe the accumulation in a way that allows the offline compiler to infer the accumulator.

Attention:
- The accumulator is only available on Arria 10 devices.

- The accumulator must be part of a loop.

- The accumulator must have an initial value of 0.

- The accumulator cannot be conditional.

Below are examples of a description that results in the correct inference of the accumulator by the offline compiler.

```
channel float4 RANDOM_STREAM;

__kernel void acc_test(__global float *a, int k) {
    // Simplest example of an accumulator.
    // In this loop, the accumulator acc is incremented by 5.
    int i;
    float acc = 0.0f;
    for (i = 0; i < k; i++) {
        acc+=5;
    }
    a[0] = acc;
}

__kernel void acc_test2(__global float *a, int k) {
    // Extended example showing that an accumulator can be
    // conditionally incremented. The key here is to describe the increment
    // as conditional, not the accumulation itself.
    int i;
    float acc = 0.0f;
    for (i = 0; i < k; i++) {
        acc += ((i < 30) ? 5 : 0);
    }
    a[0] = acc;
}

__kernel void acc_test3(__global float *a, int k) {
    // A more complex case where the accumulator is fed
    // by a dot product.
    int i;
    float acc = 0.0f;
    for (i = 0; i < k; i++ ){
        float4 v = read_channel_intel(RANDOM_STREAM);
        float x1 = v.x;
        float x2 = v.y;
        float y1 = v.z;
        float y2 = v.w;

        acc += (x1*y1+x2*y2);
    }
    a[0] = acc;
}

__kernel void loader(__global float *a, int k) {
```

```
    int i;
    float4 my_val = 0;
    for(i = 0; i < k; i++) {
        if ((i%4) == 0)
            write_channel_intel(RANDOM_STREAM, my_val);
        if ((i%4) == 0) my_val.x = a[i];
        if ((i%4) == 1) my_val.y = a[i];
        if ((i%4) == 2) my_val.z = a[i];
        if ((i%4) == 3) my_val.w = a[i];
    }
}
```

1.6.12.1 Programming Strategies for Inferring the Accumulator

To leverage the single cycle floating-point accumulator feature, you can modify the accumulator description in your kernel code to improve efficiency or work around programming restrictions.

Describing an Accumulator Using Multiple Loops

Consider a case where you want to describe an accumulator using multiple loops, with some of the loops being unrolled:

```
float acc = 0.0f;
for (i = 0; i < k; i++) {
    #pragma unroll
    for(j=0;j < 16; j++)
        acc += (x[i+j]*y[i+j]);
}
```

In this situation, it is important to compile the kernel with the --fp-relaxed Intel FPGA SDK for OpenCL Offline Compiler command option to enable the offline compiler to rearrange the operations in a way that exposes the accumulation. If you do not compile the kernel with `--fp-relaxed`, the resulting accumulator structure will have a high initiation interval (II). II is the number of cycles between launching successive loop iterations. The higher the II value, the longer the accumulator structure must wait before it can process the next loop iteration.

Modifying a Multi-Loop Accumulator Description

In cases where you cannot compile an accumulator description using the `--fp-relaxed` offline compiler command option, rewrite the code to expose the accumulation.

For the code example above, rewrite it in the following manner:

```
float acc = 0.0f;
for (i = 0; i < k; i++) {
    float my_dot = 0.0f;
    #pragma unroll
    for(j=0;j < 16; j++)
        my_dot += (x[i+j]*y[i+j]);
    acc += my_dot;
}
```

Modifying an Accumulator Description Containing a Variable or Non-Zero Initial Value

Consider a situation where you might want to apply an offset to a description of an accumulator that begins with a non-zero value:

```
float acc = array[0];
for (i = 0; i < k; i++) {
    acc += x[i];
}
```

Because the accumulator hardware does not support variable or non-zero initial values in a description, you must rewrite the description.

```
float acc = 0.0f;
for (i = 0; i < k; i++) {
    acc += x[i];
}
acc += array[0];
```

Rewriting the description in the above manner enables the kernel to use an accumulator in a loop. The loop structure is then followed by an increment of `array[0]`.

1.7 Designing Your Host Application

Intel offers guidelines on host requirements and procedures on structuring the host application. If applicable, implement these design strategies when you create or modify a host application for your OpenCL kernels.

Host Programming Requirements on page 73
    When designing your OpenCL host application for use with the Intel FPGA SDK for OpenCL, ensure that the application satisfies the following host programming requirements.

Allocating OpenCL Buffers for Manual Partitioning of Global Memory on page 74
    Manual partitioning of global memory buffers allows you to control memory accesses across buffers to maximize the memory bandwidth. You can partition buffers across interfaces of the same memory type or across interfaces of different memory types.

Collecting Profile Data During Kernel Execution on page 78
    In cases where kernel execution finishes after the host application completes, you can query the FPGA explicitly to collect profile data during kernel execution.

Accessing Custom Platform-Specific Functions on page 80
    To reference Custom Platform-specific user-accessible functions while linking to the FCD, include the `clGetBoardExtensionFunctionAddressIntel` extension in your host application.

Modifying Host Program for Structure Parameter Conversion on page 81
    If you convert any structure parameters to pointers-to-constant structures in your OpenCL kernel, you must modify your host application accordingly.

Managing Host Application on page 82
    The Intel FPGA SDK for OpenCL includes utility commands you can invoke to obtain information on flags and libraries necessary for compiling and linking your host application.

Allocating Shared Memory for OpenCL Kernels Targeting SoCs on page 93
> Intel recommends that OpenCL kernels that run on Intel SoCs access shared
> memory instead of the FPGA DDR memory.

1.7.1 Host Programming Requirements

When designing your OpenCL host application for use with the Intel FPGA SDK for
OpenCL, ensure that the application satisfies the following host programming
requirements.

1.7.1.1 Host Machine Memory Requirements

The machine that runs the host application must have enough host memory to support
several components simultaneously.

The host machine must support the following components:

- The host application and operating system.

- The working set for the host application.

- The maximum amount of OpenCL memory buffers that can be allocated at once.
  Every device-side `cl_mem` buffer is associated with a corresponding storage area
  in the host process. Therefore, the amount of host memory necessary might be as
  large as the amount of external memory supported by the FPGA.

1.7.1.2 Host Binary Requirement

When compiling the host application, target one of these architectures: x86-64 (64-
bit) or ARM® 32-bit ARMV7-A for SoCs. The Intel FPGA SDK for OpenCL host runtime
does not support x86-32 (32-bit) binaries.

1.7.1.3 Multiple Host Threads

The Intel FPGA SDK for OpenCL host library is thread-safe.

All OpenCL APIs are thread safe except the `clSetKernelArg` function.

It is safe to call `clSetKernelArg` from any host thread or as an reentrant as long as
concurrent calls to any combination of `clSetKernelArg` calls operate on different
`cl_kernel` objects.

Related Links

Multi-Threaded Host Application

1.7.1.4 Out-of-Order Command Queues

The OpenCL host runtime command queues do not support out-of-order command
execution.

1.7.1.5 Requirement for Multiple Command Queues in Channels or Pipes Implementation

Although the Intel FPGA SDK for OpenCL channels extension or OpenCL pipes implementation allows multiple kernels to execute in parallel, channels or pipes facilitate this concurrent behavior only when `cl_command_queue` objects are in order. To enable multiple command queues , instantiate a separate command for each kernel you wish to run concurrently.

1.7.2 Allocating OpenCL Buffers for Manual Partitioning of Global Memory

Manual partitioning of global memory buffers allows you to control memory accesses across buffers to maximize the memory bandwidth. You can partition buffers across interfaces of the same memory type or across interfaces of different memory types.

Related Links

- Disabling Burst-Interleaving of Global Memory (--no-interleaving
  <global_memory_type>) on page 103
    You can disable burst-interleaving for all global memory banks of the same
    type and manage them manually by including the `--no-interleaving`
    `<global_memory_type>` option in your `aoc` command.

- Manual Partitioning of Global Memory
- Heterogeneous Memory Buffers

1.7.2.1 Partitioning Buffers Across Multiple Interfaces of the Same Memory Type

Before you partition the memory across multiple interfaces of the same memory type, you must first disable burst-interleaving during OpenCL kernel compilation. Then, in the host application, you must specify the memory bank to which you allocate the OpenCL buffer.

By default, the Intel FPGA SDK for OpenCL Offline Compiler configures each global memory type in a burst-interleaved fashion. Usually, the burst-interleaving configuration leads to the best load balancing between the memory banks. However, there might be situations where it is more efficient to partition the memory into non-interleaved regions.

The figure below illustrates the differences between burst-interleaved and non-interleaved memory partitions.

Burst-Interleaved **Separate Partitions**

To manually partition some or all of the available global memory types, perform the following tasks:

1. Compile your OpenCL kernel using the `--no-interleaving <global_memory_type>` flag to configure the memory bank(s) of the specified memory type as separate addresses.

   For more information on the usage of the `--no-interleaving <global_memory_type>` flag, refer to the *Disabling Burst-Interleaving of Global Memory (--no-interleaving <global_memory_type>)* section.

2. Create an OpenCL buffer in your host application, and allocate the buffer to one of the banks using the `CL_MEM_BANK` flags.

   — Specify `CL_MEM_BANK_1_INTEL` to allocate the buffer to the lowest available memory region.

   — Specify `CL_MEM_BANK_2_INTEL` to allocation memory to the second bank (if available).

   Attention: Allocate each buffer to a single memory bank only.

   If the second bank is not available at runtime, the memory is allocated to the first bank. If no global memory is available, the `clCreateBuffer` call fails with the error message `CL_MEM_OBJECT_ALLOCATION_FAILURE`.

1.7.2.2 Partitioning Buffers Across Different Memory Types (Heterogeneous Memory)

The board support package for your FPGA board can assemble a global memory space consisting of different memory technologies (for example, DRAM or SRAM). The board support package designates one such memory, which might consist of multiple interfaces, as the default memory. All buffers reside there.

To use the heterogeneous memory, modify the code in your `.cl` file as follows:

1. Determine the names of the global memory types available on your FPGA board in one of the following ways:

   — Refer to the board vendor's documentation for more information.

   — Find the names in the `board_spec.xml` file of your board Custom Platform. For each global memory type, the name is the unique string assigned to the `name` attribute of the `global_mem` element.

2. To instruct the host to allocate a buffer to a specific global memory type, insert the `buffer_location("<memory_type>")` attribute, where *<memory_type>* is the name of the global memory type provided by your board vendor.

   For example:

   ```
   __kernel void foo(__global __attribute__((buffer_location("DDR"))) int *x,
                               __global
   __attribute__((buffer_location("QDR"))) int *y)
   ```

   If you do not specify the `buffer_location` attribute, the host allocates the buffer to the default memory type automatically. To determine the default memory type, consult the documentation provided by your board vendor. Alternatively, in the `board_spec.xml` file of your Custom Platform, search for the memory type that is defined first or has the attribute `default=1` assigned to it.

   Intel recommends that you define the `buffer_location` attribute in a preprocessor macro for ease of reuse, as follows:

   ```
   #define QDR\
                               __global\
                               __attribute__((buffer_location("QDR")))

                               #define DDR\
                               __global\
                               __attribute__((buffer_location("DDR")))

                               __kernel void foo (QDR uint * data, DDR uint
   * lup)
                               {
                               //statements
                               }
   ```

> ***Attention:*** If you assign a kernel argument to a non-default memory (for
> example, `QDR uint * data` and `DDR uint * lup` from the code
> above), you cannot declare that argument using the `constant`
> keyword. In addition, you cannot perform atomic operations with
> pointers derived from that argument.

By default, the host allocates buffers into the main memory when you load kernels
into the OpenCL runtime via the `clCreateProgramWithBinary` function. During
kernel invocation, the host automatically relocates heterogeneous memory buffers
that are bound to kernel arguments to the main memory.

3. To avoid the initial allocation of heterogeneous memory buffers in the main
   memory, include the `CL_MEM_HETEROGENEOUS_INTEL` flag when you call the
   `clCreateBuffer` function. Also, bind the `cl_mem` buffer to the argument that
   used the `buffer_location` attribute using `clSetKernelArg` before doing any
   reads or writes from that buffer, as follows:

```
mem = clCreateBuffer(context,
                     flags|CL_MEM_HETEROGENEOUS_INTEL,
                     memSize,
                     NULL,
                     &errNum);

clSetKernelArg(kernel, 0, sizeof(cl_mem), &mem);

clEnqueueWriteBuffer(queue, mem, CL_FALSE, 0, N, 0, NULL, &write_event);

clEnqueueNDRangeKernel(queue, kernel, 1, NULL, global_work_size, NULL, 0,
NULL, &kernel_event);
```

For example, the following `clCreateBuffer` call allocates memory into the
lowest available memory region of a nondefault memory bank:

```
mem = clCreateBuffer(context,
                     (CL_MEM_HETEROGENEOUS_INTEL|CL_MEM_BANK_1_INTEL),
                     memSize,
                     NULL,
                     &errNum);
```

The `clCreateBuffer` call allocates memory into a certain global memory type
based on what you specify in the kernel argument. If a memory (`cl_mem`) object
residing in a memory type is set as a kernel argument that corresponds to a
different memory technology, the host moves the memory object automatically
when it queues the kernel. Do not pass a buffer as kernel arguments that
associate it with multiple memory technologies.

For more information on optimizing heterogeneous global memory accesses, refer to
the *Heterogeneous Memory Buffers* and the *Manual Partitioning of Global Memory*
sections of the *Intel FPGA SDK for OpenCL Best Practices Guide*.

1.7.2.3 Creating a Pipe Object in Your Host Application

To implement OpenCL pipes in your kernel, you must create Intel FPGA SDK for OpenCL-specific pipe objects in your host application.

An SDK-specific pipe object is not a true OpenCL pipe object as described in the OpenCL Specification version 2.0. This implementation allows you to migrate away from Intel FPGA products with a conformant solution. The SDK-specific pipe object is a memory object (`cl_mem`); however, the host does not allocate any memory for the pipe itself.

The following `clCreatePipe` host API creates a pipe object:

```
cl_mem clCreatePipe(cl_context context,
                    cl_mem_flags flags,
                    cl_uint pipe_packet_size,
                    cl_uint pipe_max_packets,
                    const cl_pipe_properties *properties,
                    cl_int *errcode_ret)
```

For more information on the `clCreatePipe` host API function, refer to section 5.4.1 of the *OpenCL Specification version 2.0*.

Below is an example syntax of the `clCreatePipe` host API function:

```
cl_int status;
cl_mem c0_pipe = clCreatePipe(context,
                              0,
                              sizeof(int),
                              1,
                              NULL,
                              &status);
status = clSetKernelArg(kernel, 1, sizeof(cl_mem), &c0_pipe);
```

Caution: The SDK does not support dynamic channel assignment at runtime. The SDK statically links the pipes during compilation.

Related Links

OpenCL Specification version 2.0 (API)

1.7.3 Collecting Profile Data During Kernel Execution

In cases where kernel execution finishes after the host application completes, you can query the FPGA explicitly to collect profile data during kernel execution.

When you profile your OpenCL kernel during compilation, a `profile.mon` file is generated automatically. The profile data is then written to `profile.mon` after kernel execution completes on the FPGA. However, if kernel execution completes after the host application completes, no profiling information for that kernel invocation will be available in the `profile.mon` file. In this case, you can modify your host code to acquire profiling information during kernel execution.

Important: Collecting profile data during kernel execution can add significant overhead to kernel launches by increasing the latency in your kernel.

- To query the FPGA to collect profile data while the kernel is running, call the following host library call:

```
extern CL_API_ENTRY cl_int CL_API_CALL

clGetProfileInfoIntelFPGA(cl_event);
```

where `cl_event` is the kernel event. The kernel event you pass to this host library call must be the same one you pass to the `clEnqueueNDRangeKernel` call.

Important: If kernel execution completes before the invocation of `clGetProfileInfoIntelFPGA`, the function returns an event error message.

Caution: Invoking the `clGetProfileInfoIntelFPGA` function during kernel execution disables the profile counters momentarily so that the Profiler can collect data from the FPGA. As a result, you will lose some profiling information during this interruption. If you call this function at very short intervals, the profile data might not accurately reflect the actual performance behavior of the kernel.

Consider the following example host code:

```
int main()
{    ...
     clEnqueueNDRangeKernel (queue, kernel, ..., NULL);
     ...
     clEnqueueNDRangeKernel (queue, kernel, .. , NULL);
     ...
}
```

This host application runs on the assumption that a kernel launches twice and then completes. In the `profile.mon` file, there will be two sets of profile data, one for each kernel invocation. To collect profile data while the kernel is running, modify the host code in the following manner:

```
int main()
{
     ...
     clEnqueueNDRangeKernel (queue, kernel, ..., &event);

     //Get the profile data before the kernel completes
     clGetProfileInfoIntelFPGA (event);

     //Wait until the kernel completes
     clFinish (queue);

     ...
     clEnqueueNDRangeKernel (queue, kernel, ..., NULL);
     ...
}
```

The call to `clGetProfileInfoIntelFPGA` adds a new entry in the `profile.mon` file. The Profiler GUI then parses this entry in the report.

For more information on the Intel FPGA SDK for OpenCL Profiler, refer to the following sections:

- *Profile Your Kernel to Identify Performance Bottlenecks* in the *Intel FPGA SDK for OpenCL Best Practices Guide*

- *Profiling Your OpenCL Kernel*

Related Links

- Profile Your Kernel to Identify Performance Bottlenecks

- Profiling Your OpenCL Kernel on page 112
  The Intel FPGA SDK for OpenCL Profiler measures and reports performance data collected during OpenCL kernel execution on the FPGA.

1.7.4 Accessing Custom Platform-Specific Functions

You have the option to include in your application user-accessible functions that are available in your Custom Platform. However, when you link your host application to the FPGA Client Driver (FCD), formerly Altera Client Driver (ACD), you cannot directly reference these Custom Platform-specific functions. To reference Custom Platform-specific user-accessible functions while linking to the FCD, include the `clGetBoardExtensionFunctionAddressIntel` extension in your host application.

The `clGetBoardExtensionFunctionAddressIntel` extension specifies an API that retrieves a pointer to a user-accessible function from the Custom Platform.

Attention: For Linux systems, the `clGetBoardExtensionFunctionAddressIntel` function works with or without FCD. For Windows systems, the function only works in conjunction with FCD. Consult with your board vendor to determine if FCD is supported in your Custom Platform.

Definitions of the extension interfaces are available in the *`ALTERAOCLSDKROOT`*`/host/include/CL/cl_ext.h` file.

- To obtain a pointer to a user-accessible function in your Custom Platform, call the following function in your host application:

```
void* clGetBoardExtensionFunctionAddressIntel (
    const char* function_name,
    cl_device_id device
    );
```

Where:

function_name is the name of the user-accessible function that your Custom Platform vendor provides,

and

device is the device ID returned by the `clGetDeviceIDs` function.

After locating the user-accessible function, the `clGetBoardExtensionFunctionAddressIntel` function returns a pointer to the user-accessible function. If the function does not exist in the Custom Platform, `clGetBoardExtensionFunctionAddressIntel` returns `NULL`.

Attention: To access the `clGetBoardExtensionFunctionAddressIntel` API via the Installable Client Driver (ICD), ensure that the ICD extension API `clGetExtensionFunctionAddress` retrieves the pointer to the `clGetBoardExtensionFunctionAddressIntel` API first.

The following code example shows how you can access the Custom Platform-specific function via ICD:

```
typedef void* (*get_board_extension_function_address_fn_t)
              (const char* func_name, cl_device_id device);

get_board_extension_function_address_fn_t
  get_board_extension_function_address =
  (get_board_extension_function_address_fn_t)
  clGetExtensionFunctionAddress
  ("clGetBoardExtensionFunctionAddressIntel");

if (get_board_extension_function_address == NULL ) {
  printf ("Failed to get clGetBoardExtensionFunctionAddressIntel
  \n");
}

void *board_extension_function_ptr =
  get_board_extension_function_address("function_name",device_id);
```

1.7.5 Modifying Host Program for Structure Parameter Conversion

If you convert any structure parameters to pointers-to-constant structures in your OpenCL kernel, you must modify your host application accordingly.

Perform the following changes to your host application:

1. Allocate a `cl_mem` buffer to store the structure contents.

> ***Attention:*** You need a separate `cl_mem` buffer for every kernel that uses a different structure value.

2. Set the structure kernel argument with a pointer to the structure buffer, not with a pointer to the structure contents.

3. Populate the structure buffer contents before queuing the kernel. Perform one of the following steps to ensure that the structure buffer is populated before the kernel launches:

   — Queue the structure buffer on the same command queue as the kernel queue.

   — Synchronize separate kernel queues and structure buffer queues with an event.

4. When your application no longer needs to call a kernel that uses the structure buffer, release the `cl_mem` buffer.

Related Links

- Including Structure Data Types as Arguments in OpenCL Kernels on page 64
  Convert each structure parameter (`struct`) to a pointer that points to a structure.

- Matching Data Layouts of Host and Kernel Structure Data Types on page 65
  If you use structure data types (`struct`) as arguments in OpenCL kernels, match the member data types and align the data members between the host application and the kernel code.

1.7.6 Managing Host Application

The Intel FPGA SDK for OpenCL includes utility commands you can invoke to obtain information on flags and libraries necessary for compiling and linking your host application.

Attention: To cross-compile your host application to an SoC FPGA board, include the `--arm` option in your utility command.

Caution: For Linux systems, if you debug your host application using the GNU Project Debugger (GDB), invoke the following command prior to running the host application:

```
handle SIG44 nostop
```

Without this command, the GDB debugging process terminates with the following error message:

```
Program received signal SIG44, Real-time event 44.
```

1.7.6.1 Displaying Example Makefile Fragments (example-makefile or makefile)

To display example Makefile fragments for compiling and linking a host application against host runtime libraries available with the Intel FPGA SDK for OpenCL, invoke the `example-makefile` or `makefile` utility command.

- At a command prompt, invoke the `aocl example-makefile` or `aocl makefile` utility command.
  The software displays an output similar to the following:

```
The following are example Makefile fragments for compiling and linking
a host program against the host runtime libraries included with the
Intel FPGA SDK for OpenCL.


Example GNU makefile on Linux, with GCC toolchain:

    AOCL_COMPILE_CONFIG=$(shell aocl compile-config)
    AOCL_LINK_CONFIG=$(shell aocl link-config)

    host_prog : host_prog.o
        g++ -o host_prog host_prog.o $(AOCL_LINK_CONFIG)

    host_prog.o : host_prog.cpp
        g++ -c host_prog.cpp $(AOCL_COMPILE_CONFIG)


Example GNU makefile on Windows, with Microsoft Visual C++ command line
compiler:

    AOCL_COMPILE_CONFIG=$(shell aocl compile-config)
    AOCL_LINK_CONFIG=$(shell aocl link-config)

    host_prog.exe : host_prog.obj
        link -nologo /OUT:host_prog.exe host_prog.obj $(AOCL_LINK_CONFIG)

    host_prog.obj : host_prog.cpp
        cl /MD /Fohost_prog.obj -c host_prog.cpp $(AOCL_COMPILE_CONFIG)


Example GNU makefile cross-compiling to ARM SoC from Linux or Windows, with
Linaro GCC cross-compiler toolchain:

    CROSS-COMPILER=arm-linux-gnueabihf-
    AOCL_COMPILE_CONFIG=$(shell aocl compile-config --arm)
    AOCL_LINK_CONFIG=$(shell aocl link-config --arm)

    host_prog : host_prog.o
        $(CROSS-COMPILER)g++ -o host_prog host_prog.o $(AOCL_LINK_CONFIG)

    host_prog.o : host_prog.cpp
        $(CROSS-COMPILER)g++ -c host_prog.cpp $(AOCL_COMPILE_CONFIG)
```

1.7.6.2 Compiling and Linking Your Host Application

The OpenCL host application uses standard OpenCL runtime APIs to manage device configuration, data buffers, kernel launches, and synchronization. The host application also contains functions such as file I/O, or portions of the source code that do not run on an accelerator device. The Intel FPGA SDK for OpenCL includes utility commands you can invoke to obtain information on C header files describing the OpenCL APIs, and board-specific MMD and host runtime libraries with which you must link your host application.

Important: For Windows systems, you must add the `/MD` flag to link the host runtime libraries against the multithreaded dynamic link library (DLL) version of the Microsoft C Runtime library. You must also compile your host application with the `/MD` compilation flag, or use the `/NODEFAULTLIB` linker option to override the selection of runtime library.

Remember: Include the path to the `ALTERAOCLSDKROOT/host/<OS_platform>/bin` folder in your library search path when you run your host application.

Displaying Flags for Compiling Host Application (compile-config) on page 84
   To display a list of flags necessary for compiling a host application, invoke the `compile-config` utility command.

Displaying Paths to OpenCL Host Runtime and MMD Libraries (ldflags) on page 84
   To display the paths necessary for linking a host application to the OpenCL host runtime and MMD libraries, invoke the `ldflags` utility command.

Listing OpenCL Host Runtime and MMD Libraries (ldlibs) on page 85
   To display the names of the OpenCL host runtime and MMD libraries necessary for linking a host application, invoke the `ldlibs` utility command.

Displaying Information on OpenCL Host Runtime and MMD Libraries (link-config or linkflags) on page 85
   To display a list of flags necessary for linking a host application with OpenCL host runtime and MMD libraries, invoke the `link-config` or `linkflags` utility command.

1.7.6.2.1 Displaying Flags for Compiling Host Application (compile-config)

To display a list of flags necessary for compiling a host application, invoke the `compile-config` utility command.

1. At a command prompt, invoke the `aocl compile-config` utility command. The software displays the path to the folder or directory in which the OpenCL API header files reside. For example:

   - For Windows systems, the path is `-I%ALTERAOCLSDKROOT%/host/include`

   - For Linux systems, the path is `-I$ALTERAOCLSDKROOT/host/include`

   where *ALTERAOCLSDKROOT* points to the location of the software installation.

2. Add this path to your C preprocessor.

Attention: In your host source, include the `opencl.h` OpenCL header file, located in the `ALTERAOCLSDKROOT/host/include/CL` folder or directory.

1.7.6.2.2 Displaying Paths to OpenCL Host Runtime and MMD Libraries (ldflags)

To display the paths necessary for linking a host application to the OpenCL host runtime and MMD libraries, invoke the `ldflags` utility command.

- At a command prompt, invoke the `aocl ldflags` utility command.
  The software displays the paths for linking your host application with the following libraries:

  1. The OpenCL host runtime libraries that provide OpenCL platform and runtime APIs. The OpenCL host runtime libraries are available in the `ALTERAOCLSDKROOT/host/<OS_platform>/lib` directory.

  2. The path to the Custom Platform-specific MMD libraries. The MMD libraries are available in the `<board_family_name>/<OS_platform>/lib` directory of your Custom Platform.

1.7.6.2.3 Listing OpenCL Host Runtime and MMD Libraries (ldlibs)

To display the names of the OpenCL host runtime and MMD libraries necessary for linking a host application, invoke the `ldlibs` utility command.

- At a command prompt, invoke the `aocl ldlibs` utility command.
  The software lists the OpenCL host runtime libraries residing in the *ALTERAOCLSDKROOT*/host/*<OS_platform>*/lib directory. It also lists the Custom Platform-specific MMD libraries residing in the /*<board_family_name>*/*<OS_platform>*/lib directory of your Custom Platform.

  — For Windows systems, the output might resemble the following example:

  ```
  alterahalmmd.lib
  <board_vendor_name>_<board_family_name>_mmd.[lib|so|a|dll]
  alteracl.lib
  acl_emulator_kernel_rt.lib
  pkg_editor.lib
  libelf.lib
  acl_hostxml.lib
  ```

  — For Linux systems, the output might resemble the following example:

  ```
  -lalteracl
  -ldl
  -lacl_emulator_kernel_rt
  -lalterahalmmd
  -l<board_vendor_name>_<board_family_name>_mmd
  -lelf
  -lrt
  -lstdc++
  ```

1.7.6.2.4 Displaying Information on OpenCL Host Runtime and MMD Libraries (link-config or linkflags)

To display a list of flags necessary for linking a host application with OpenCL host runtime and MMD libraries, invoke the `link-config` or `linkflags` utility command.

This utility command combines the functions of the `ldflags` and `ldlibs` utility commands.

1. At a command prompt, invoke the `aocl link-config` or `aocl linkflags` command.
   The software displays the link options for linking your host application with the following libraries:

   a. The path to and the names of OpenCL host runtime libraries that provide OpenCL platform and runtime APIs. The OpenCL host runtime libraries are available in the *ALTERAOCLSDKROOT*/host/*<OS_platform>*/lib directory .

   b. The path to and the names of the Custom Platform-specific MMD libraries. The MMD libraries are available in the *<board_family_name>*/*<OS_platform>*/lib directory of your Custom Platform.

- For Windows systems, the link options might resemble the following example output:

```
/libpath:%ALTERAOCLSDKROOT%/board/<board_name>/windows64/lib
/libpath:%ALTERAOCLSDKROOT%/host/windows64/lib
alterahalmmd.lib
<board_vendor_name>_<board_family_name>_mmd.[lib|so|a|dll]
alteracl.lib
acl_emulator_kernel_rt.lib
pkg_editor.lib
libelf.lib
acl_hostxml.lib
```

- For Linux systems, the link options might resemble the following example output:

```
-L/$ALTERAOCLSDKROOT/board/<board_name>/linux64/lib
-L/$ALTERAOCLSDKROOT/host/linux64/lib
-lalterac
-ldl
-lacl_emulator_kernel_rt
-lalterahalmmd
-l<board_vendor_name>_<board_family_name>_mmd
-lelf
-lrt
-lstdc++
```

1.7.6.3 Linking Your Host Application to the Khronos ICD Loader Library

The Intel FPGA SDK for OpenCL supports the OpenCL ICD extension from the Khronos Group. The OpenCL ICD extension allows you to have multiple OpenCL implementations on your system. With the OpenCL ICD Loader Library, you may choose from a list of installed platforms and execute OpenCL API calls that are specific to your OpenCL implementation of choice.

In addition to the SDK's host runtime libraries, Intel supplies a version of the ICD Loader Library that supports the OpenCL Specification version 1.0 and the implemented APIs from the OpenCL Specification versions 1.1, 1.2, and 2.0. To use an ICD library from another vendor, consult the vendor's documentation on how to link to their ICD library.

Linking to the ICD Loader Library on Windows on page 86
To link your Windows OpenCL host application to the ICD Loader Library, modify the `Makefile` and set up the FPGA Client Driver (FCD), formerly the Altera® Client Driver (ACD).

Linking to the ICD Loader Library on Linux on page 87
To link your Linux OpenCL host application to the ICD Loader Library, modify the `Makefile`. For Cyclone V SoC boards, you also have to create an `Altera.icd` file.

1.7.6.3.1 Linking to the ICD Loader Library on Windows

To link your Windows OpenCL host application to the ICD Loader Library, modify the `Makefile` and set up the FPGA Client Driver (FCD), formerly the Altera® Client Driver (ACD).

Attention: For Windows systems, you must use the ICD in conjunction with the FCD. If the custom platform from your board vendor does not currently support FCD, you can set it up manually.

1. Prior to linking your host application to any Intel FPGA SDK for OpenCL host runtime libraries, link it to the OpenCL library by modifying the `Makefile`.

   A modified `Makefile` might include the following lines:

   ```
   AOCL_COMPILE_CONFIG=$(shell aocl compile-config)
   AOCL_LDFLAGS=$(shell aocl ldflags)
   AOCL_LDLIBS=$(shell aocl ldlibs)

   host_prog.exe : host_prog.obj
       link -nologo /OUT:host_prog.exe host_prog.obj $(AOCL_ LDFLAGS) OpenCL.lib

   host_prog.obj : host_prog.cpp
       cl /MD /Fohost_prog.obj -c host_prog.cpp $(AOCL_COMPILE_CONFIG)
   ```

2. If you need to manually set up FCD support for your Custom Platform, perform the following tasks:

   a. Consult with your board vendor to identify the libraries that the FCD requires. Alternatively, you can invoke the `aocl ldlibs` command and identify the libraries that your OpenCL application requires.

   b. Specify the libraries in the registry key **HKEY_LOCAL_MACHINE \SOFTWARE\Altera\OpenCL\Boards**. Specify the value name to be the path to the library, and specify the data to be a **DWORD** that is set to 0.

      Attention: If your board vendor provides multiple libraries, you might need to load them in a particular order. Consult with your board vendor to determine the correct order to load the libraries. List the libraries in the registry in their loading order.

To enumerate board vendor-specific ICDs, the ICD Loader scans the values in the **HKEY_LOCAL_MACHINE\SOFTWARE\Altera\OpenCL\Boards** registry key. For each **DWORD** value that is set to 0, the ACD Loader opens the corresponding DLL that is specified in the value name.

Consider the following registry key value:

[HKEY_LOCAL_MACHINE\SOFTWARE\Altera\OpenCL\Boards] "c: \board_vendor a\my_board_mmd.dll"=dword:00000000

The ICD Loader scans this value, and then the ACD Loader opens the library `my_board_mmd.dll` from the `board_vendor a` folder.

Attention: If your host application fails to run while it is linking to the ICD, ensure that the **HKEY_LOCAL_MACHINE\SOFTWARE\Khronos\OpenCL\Vendors** registry key contains the following value:

[HKEY_LOCAL_MACHINE\SOFTWARE\Khronos\OpenCL\Vendors] "alteracl_icd.dll"=dword:00000000

1.7.6.3.2 Linking to the ICD Loader Library on Linux

To link your Linux OpenCL host application to the ICD Loader Library, modify the `Makefile`. For Cyclone V SoC boards, you also have to create an `Altera.icd` file.

1. Prior to linking your host application to any Intel FPGA SDK for OpenCL host runtime libraries, link it to the OpenCL library by modifying the `Makefile`.

   A modified `Makefile` might include the following lines:

   ```
   AOCL_LDFLAGS=$(shell aocl ldflags)
   AOCL_LDLIBS=$(shell aocl ldlibs)

   host_prog : host_prog.o
   g++ -o host_prog host_prog.o $(AOCL_LDFLAGS) -lOpenCL $(AOCL_LDLIBS)
   ```

2. For Cyclone V SoC boards, when you build the SD flash card image for your Custom Platform, create an `Altera.icd` file containing the text `libalteracl.so`. Store the `Altera.icd` file in the `/etc/OpenCL/vendors` directory of your Custom Platform.

   Refer to *Building an SD Flash Card Image* section of the *Intel FPGA SDK for OpenCL Cyclone V SoC Development Kit Reference Platform Porting Guide* for more information.

Attention: If your host application fails to run while linking to the ICD, ensure that the file `/etc/OpenCL/vendors/Altera.icd` matches the file found in the directory that *ALTERAOCLSDKROOT* specifies. The environment variable *ALTERAOCLSDKROOT* points to the location of the SDK installation. If the files do not match, or if it is missing from `/etc/OpenCL/vendors`, copy the `Altera.icd` file from *ALTERAOCLSDKROOT* to `/etc/OpenCL/vendors`.

Related Links
Building an SD Flash Card Image

1.7.6.4 Programming an FPGA via the Host

The Intel FPGA SDK for OpenCL Offline Compiler is an offline compiler that compiles kernels independently of the host application. To load the kernels into the OpenCL runtime, include the `clCreateProgramWithBinary` function in your host application.

Caution: If your host system consists of multiple processors, only one processor can access the FPGA at a given time. Consider an example where there are two host applications, corresponding to two processors, attempting to launch kernels onto the same FPGA at the same time. The second host application wil receive an error message indicating that the device is busy. The second host application cannot run until the first host application releases the OpenCL context.

1. Compile your OpenCL kernel with the offline compiler to create the `.aocx` file.

2. Include the `clCreateProgramWithBinary` function in your host application to create the `cl_program` OpenCL program objects from the `.aocx` file.

3. Include the `clBuildProgram` function in your host application to create the program executable for the specified device.

Below is an example host code on using `clCreateProgramWithBinary` to program an FPGA device:

```
size_t lengths[1];
unsigned char* binaries[1] ={NULL};
cl_int status[1];
cl_int error;
cl_program program;
const char options[] = "";

FILE *fp = fopen("program.aocx","rb");
fseek(fp,0,SEEK_END);
lengths[0] = ftell(fp);
binaries[0] = (unsigned char*)malloc(sizeof(unsigned char)*lengths[0]);
rewind(fp);
fread(binaries[0],lengths[0],1,fp);
fclose(fp);

program = clCreateProgramWithBinary(context,
                                    1,
                                    device_list,
                                    lengths,
                                    (const unsigned char **)binaries,
                                    status,
                                    &error);
clBuildProgram(program,1,device_list,options,NULL,NULL);
```

If the `clBuildProgram` function executes successfully, it returns `CL_SUCCESS`.

4. Create kernel objects from the program executable using the `clCreateKernelsInProgram` or `clCreateKernel` function.

5. Include the kernel execution function to instruct the host runtime to execute the scheduled kernel(s) on the FPGA.

   — To enqueue a command to execute an NDRange kernel, use `clEnqueueNDRangeKernel`.

   — To enqueue a single work-item kernel, use `clEnqueueTask`.

   Attention: Intel recommends that you release an event object when it is not in use. The SDK keeps an event object live until you explicitly instruct it to release the event object. Keeping an unused event object live causes unnecessary memory usage.

   To release an event object, call the `clReleaseEvent` function.

You can load multiple FPGA programs into memory, which the host then uses to reprogram the FPGA as required.

For more information on these OpenCL host runtime API calls, refer to the *OpenCL Specification version 1.0*.

Related Links

OpenCL Specification version 1.0

1.7.6.4.1 Programming Multiple FPGA Devices

If you install multiple FPGA devices in your system, you can direct the host runtime to program a specific FPGA device by modifying your host code.

Important: By default, you may only program multiple FPGA devices from the *same* Custom Platform because the *AOCL_BOARD_PACKAGE_ROOT* environment variable points to the location of a single Custom Platform.

Linking your host application to ACD allows you to target multiple FPGA devices from different Custom Platforms. However, this feature has limited support for Custom Platforms that are compatible with SDK versions prior to 16.1.

You can present up to 32 FPGA devices to your system in the following manner:

- Multiple FPGA accelerator boards, each consisting of a single FPGA.

- Multiple FPGAs on a single accelerator board that connects to the host system via a PCIe switch.

- Combinations of the above.

The host runtime can load kernels onto each and every one of the FPGA devices. The FPGA devices can then operate in a parallel fashion.

Related Links

Accessing Custom Platform-Specific Functions on page 80
> To reference Custom Platform-specific user-accessible functions while linking to the FCD, include the `clGetBoardExtensionFunctionAddressIntel` extension in your host application.

Probing the OpenCL FPGA Devices

The host must identify the number of OpenCL FPGA devices installed into the system.

1. To query a list of FPGA devices installed in your machine, invoke the `aocl diagnose` command.

2. To direct the host to identify the number of OpenCL FPGA devices, add the following lines of code to your host application:

```
//Get the platform
ciErrNum = clGetPlatformID(&cpPlatform);

//Get the devices
ciErrNum = clGetDeviceIDs(cpPlatform,
                          CL_DEVICE_TYPE_ALL,
                          0,
                          NULL,
                          &ciDeviceCount);
cdDevices = (cl_device_id * )malloc(ciDeviceCount * sizeof(cl_device_id));
ciErrNum = clGetDeviceIDs(cpPlatform,
                          CL_DEVICE_TYPE_ALL,
                          ciDeviceCount,
                          cdDevices,
                          NULL);
```

For example, on a system with two OpenCL FPGA devices, `ciDeviceCount` has a value of 2, and `cdDevices` contains a list of two device IDs (`cl_device_id`).

Related Links

Querying the Device Name of Your FPGA Board (diagnose) on page 17
> When you query a list of accelerator boards, the OpenCL software produces a list of installed devices on your machine in the order of their device names.

Querying Device Information

You can direct the host to query information on your OpenCL FPGA devices.

- To direct the host to output a list of OpenCL FPGA devices installed into your system, add the following lines of code to your host application:

```
char buf[1024];
for (unsigned i = 0; i < ciDeviceCount; i++);
{
    clGetDeviceInfo(cdDevices[i], CL_DEVICE_NAME, 1023, buf, 0);
    printf("Device %d: '%s'\n", i, buf);
}
```

When you query the device information, the host will list your FPGA devices in the following manner: Device *<N>*: *<board_name>*: *<name_of_FPGA_board>*

Where:

- *<N>* is the device number.

- *<board_name>* is the board designation you use to target your FPGA device when you invoke the `aoc` command.

- *<name_of_FPGA_board>* is the advertised name of the FPGA board.

For example, if you have two identical FPGA boards on your system, the host generates an output that resembles the following:

```
Device 0: board_1: Stratix V FPGA Board
Device 1: board_1: Stratix V FPGA Board
```

Note: The `clGetDeviceInfo` function returns the board type (for example, `board_1`) that the Intel FPGA SDK for OpenCL Offline Compiler lists on-screen when you invoke the `aoc --list-boards` command. If your accelerator board contains more than one FPGA, each device is treated as a "board" and is given a unique name.

Related Links

Loading Kernels for Multiple FPGA Devices

If your system contains multiple FPGA devices, you can create specific `cl_program` objects for each FPGA and load them into the OpenCL runtime.

The following host code demonstrates the usage of the `clCreateProgramWithBinary` and `createMultiDeviceProgram` functions to program multiple FPGA devices:

```
cl_program createMultiDeviceProgram(cl_context context,
                                    const cl_device_id *device_list,
                                    cl_uint num_devices,
                                    const char *aocx_name);

// Utility function for loading file into Binary String
//
unsigned char* load_file(const char* filename, size_t *size_ret)
{
```

```
    FILE *fp = fopen(aocx_name,"rb");
    fseek(fp,0,SEEK_END);
    size_t len = ftell(fp);
    char *result = (unsigned char*)malloc(sizeof(unsigned char)*len);
    rewind(fp);
    fread(result,len,1,fp);
    fclose(fp);
    *size_ret = len;
    return result;
}

//Create a Program that is compiled for the devices in the "device_list"
//
cl_program createMultiDeviceProgram(cl_context context,
                                    const cl_device_id *device_list,
                                    cl_uint num_devices,
                                    const char *aocx_name)
{
    printf("creating multi device program %s for %d devices\n",
            aocx_name, num_devices);
    const unsigned char **binaries =
        (const unsigned char**)malloc(num_devices*sizeof(unsigned char*));
    size_t *lengths=(size_t*)malloc(num_devices*sizeof(size_t));
    cl_int err;

    for(cl_uint i=0; i<num_devices; i++)
    {
        binaries[i] = load_file(aocx_name,&lengths[i]);
        if (!binaries[i])
        {
            printf("couldn't load %s\n", aocx_name);
            exit(-1);
        }
    }

    cl_program p = clCreateProgramWithBinary(context,
                                             num_devices,
                                             device_list,
                                             lengths,
                                             binaries,
                                             NULL,
                                             &err);
    free(lengths);
    free(binaries);

    if (err != CL_SUCCESS)
    {
        printf("Program Create Error\n");
    }
    return p;
}


// main program

main ()
{
    // Normal OpenCL setup
}
program = createMultiDeviceProgram(context,
                                   device_list,
                                   num_devices,
                                   "program.aocx");
clBuildProgram(program,num_devices,device_list,options,NULL,NULL);
```

1.7.6.5 Termination of the Runtime Environment and Error Recovery

In the event that the host application terminates unexpectedly, you must restart the runtime environment and reprogram the FPGA.

The runtime environment is a library that is compiled as part of the host application. When the host application terminates, the runtime environment will also terminate along with any tracking activity that it performs. If you restart the host application, a new runtime environment and its associated tracking activities will reinitialize. The initialization functions reset the kernel's hardware state.

In same cases, unexpected termination of the host application causes the configuration of certain hardware (for example, PCIe hard IP) to be incomplete. To restore the configuration of these hardware, the host needs to reprogram the FPGA.

If you use a Custom Platform that implements customized hardware blocks, be aware that restarting the host application and resetting these blocks might have design implications:

- When the host application calls the `clGetPlatformIDs` function, all kernels and channels will be reset for all available devices.

- When the host application calls the `clGetPlatformIDs` function, it resets FIFO buffers and channels as it resets the device.

- The host application initializes memory buffers via the `clCreateBuffer` and `clEnqueueWriteBuffer` function calls. You cannot access the contents of buffers from a previous host execution within a new host execution.

1.7.7 Allocating Shared Memory for OpenCL Kernels Targeting SoCs

Intel recommends that OpenCL kernels that run on Intel SoCs access shared memory instead of the FPGA DDR memory. FPGA DDR memory is accessible to kernels with very high bandwidths. However, read and write operations from the ARM CPU to FPGA DDR memory are very slow because they do not use direct memory access (DMA). Reserve FPGA DDR memory only for passing temporary data between kernels or within a single kernel for testing purposes.

Note:
- Mark the shared buffers between kernels as volatile to ensure that buffer modification by one kernel is visible to the other kernel.

- To access shared memory, you only need to modify the host code. Modifications to the kernel code are unnecessary.

- You cannot use the library function `malloc` or the operator `new` to allocate physically shared memory. Also, the `CL_MEM_USE_HOST_PTR` flag does not work with shared memory.

  In DDR memory, shared memory must be physically contiguous. The FPGA cannot consume virtually contiguous memory without a scatter-gather direct memory access (SG-DMA) controller core. The `malloc` function and the `new` operator are for accessing memory that is virtually contiguous.

- CPU caching is disabled for the shared memory.

- When you use shared memory, one copy of the data is used for both the host and the kernel. When this memory is used, OpenCL memory calls are done as zero-copy transfers for buffer reads, buffer writers, maps, and unmaps.

The ARM CPU and the FPGA can access the shared memory simultaneously. You do not need to include the `clEnqueueReadBuffer` and `clEnqueueWriteBuffer` calls in your host code to make data visible to either the FPGA or the CPU.

- To allocate and access shared memory, structure your host code in a similar manner as the following example:

```
cl_mem src = clCreateBuffer(…, CL_MEM_ALLOC_HOST_PTR, size, …);
int *src_ptr  = (int*)clEnqueueMapBuffer  (…, src, size, …);
*src_ptr = input_value; //host writes to ptr directly
clSetKernelArg (…, src);
clEnqueueNDRangeKernel(…);
clFinish();
printf ("Result = %d\n", *dst_ptr); //result is available immediately
clEnqueueUnmapMemObject(…, src, src_ptr, …);
clReleaseMemObject(src); // actually frees physical memory
```
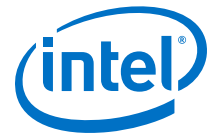
You can include the `CONFIG_CMA_SIZE_MBYTES` kernel configuration option to control the maximum total amount of shared memory available for allocation. In practice, the total amount of allocated shared memory is smaller than the value of `CONFIG_CMA_SIZE_MBYTES`.

Important: 1. If your target board has multiple DDR memory banks, the `clCreateBuffer(..., CL_MEM_READ_WRITE,  ...)` function allocates memory to the nonshared DDR memory banks. However, if the FPGA has access to a single DDR bank that is shared memory, then `clCreateBuffer(..., CL_MEM_READ_WRITE,  ...)` allocates to shared memory, similar to using the `CL_MEM_ALLOC_HOST_PTR` flag.

2. The shared memory that you request with the `clCreateBuffer(..., CL_MEM_ALLOC_HOST_PTR, size,  ...)` function is allocated in the Linux OpenCL kernel driver, and it relies on the contiguous memory allocator (CMA) feature of the Linux kernel. For detailed information on enabling and configuring the CMA, refer to the *Recompiling the Linux Kernel and the OpenCL Linux Kernel Driver* section of the *Intel FPGA SDK for OpenCL Cyclone V SoC Development Kit Reference Platform Porting Guide*.

- To transfer data from shared hard processor system (HPS) DDR to FPGA DDR efficiently, include a kernel that performs the `memcpy` function, as shown below.

```
__attribute__((num_simd_work_items(8)))
mem_stream(__global uint * src, __global uint * dst)
{
    size_t gid = get_global_id(0);
    dst[gid] = src[gid];
}
```

Attention: Allocate the `src` pointer in the HPS DDR as shared memory using the `CL_MEM_ALLOC_HOST_PTR` flag.

- If the host allocates constant memory to shared HPS DDR system and then modifies it after kernel execution, the modifications might not take effect. As a result, subsequent kernel executions might use outdated data. To prevent kernel execution from using outdated constant memory, perform one of the following tasks:

  a. Do not modify constant memory after its initialization.

  b. Create multiple constant memory buffers if you require multiple `__constant` data sets.

  c. If available, allocate constant memory to the FPGA DDR on your accelerator board.

Related Links

Recompiling the Linux Kernel and the OpenCL Linux Kernel Driver

1.8 Compiling Your OpenCL Kernel

The Intel FPGA SDK for OpenCL offers a list of compiler options that allows you to customize the kernel compilation process. An Intel FPGA SDK for OpenCL Offline Compiler command consists of the `aoc` command, compiler option(s) and settings, and kernel filenames. You can invoke an `aoc` command to direct the compiler to target a specific FPGA board, generate reports, or implement optimization techniques.

Before you compile an OpenCL kernel, ensure that the environment variable *AOCL_BOARD_PACKAGE_ROOT* points to the location of the appropriate Custom Platform. Also, verify that the *QUARTUS_ROOTDIR_OVERRIDE* environment variable points to the correct edition of the Quartus Prime software.

If these environment variables do not have the correct settings, follow the instructions in the *Setting the Intel FPGA SDK for OpenCL User Environment Variables* section of the *Intel FPGA SDK for OpenCL Getting Started Guide* to modify the settings.

Compiling Encrypted Source: When you compile an encrypted `.cl` file provided to you, you can compile only that file with the `aoc` command . You cannot compile multiple encyrpted `.cl` files at the same time with the `aoc` command. You cannot compile source that you encrypt yourself.

Attention: If you use the Stratix V Network Reference Platform, you must acquire and install the PLDA QuickUDP intellectual property (IP) core license. Refer to the PLDA website for more information. If you use a Custom Platform that includes the QuickUDP IP core, refer to your board vendor's documentation for more information on the acquisition and installation of the QuickUDP IP license.

Caution: Improper installation of the QuickUDP IP license causes kernel compilation to fail with the following error message:

```
Error (292014): Can't find valid feature line for core PLDA
QUICKTCP (73E1_AE12) in current license.
```

Note that the error has no actual dependency on the TCP Hardware Stack QuickTCP IP from PLDA.

Related Links

- Setting the Intel FPGA SDK for OpenCL User Environment Variables (Windows)
- Setting the Intel FPGA SDK for OpenCL User Environment Variables (Linux)

1.8.1 Compiling Your Kernel to Create Hardware Configuration File

You can compile an OpenCL kernel and create the hardware configuration file (that is, the `.aocx` file) in a single step.

Intel recommends that you use this one-step compilation strategy under the following circumstances:

- After you optimize your kernel via the Intel FPGA SDK for OpenCL design flow, and you are now ready to create the `.aocx` file for deployment onto the FPGA.

- You have one or more simple kernels that do not require any optimization.

- To compile the kernel and generate the `.aocx` file in one step, invoke the `aoc` *<your_kernel_filename1>*`.cl [`*<your_kernel_filename2>*`.cl ...]` command.

  Where [*<your_kernel_filename2>*`.cl ...]` are the optional space-delimited file names of kernels that you can compile in addition to *<your_kernel_filename1>*`.cl`.

The Intel FPGA SDK for OpenCL Offline Compiler groups the `.cl` files into a temporary file. It then compiles this file to generate the `.aocx` file. You must specify the order of the kernels in this temporary file on the command line.

1.8.2 Compiling Your Kernel without Building Hardware (-c)

To direct the Intel FPGA SDK for OpenCL Offline Compiler to compile your OpenCL kernel and generate a Quartus Prime hardware design project without creating a hardware configuration file, include the `-c` option in your `aoc` command.

- At a command prompt, invoke the `aoc -c` *<your_kernel_filename1>*`.cl [`*<your_kernel_filename2>*`.cl ...]` command.

  Where [*<your_kernel_filename2>*`.cl ...]` are the optional space-delimited file names of kernels that you can compile in addition to *<your_kernel_filename1>*`.cl`.

When you invoke the `aoc` command with the `-c` flag, the offline compiler compiles the kernel and creates the following files and directories:

  — The `.aoco` file. The offline compiler creates the `.aoco` file in a matter of seconds to minutes. If you compile multiple kernels, their information in the `.aoco` file appears in the order in which you list them on the command line.

  — A *<your_kernel_filename>* folder or subdirectory. It contains intermediate files that the SDK uses to build the hardware configuration file necessary for FPGA programming.

1.8.3 Specifying the Location of Header Files (-I <directory>)

To add a directory to the list of directories that the Intel FPGA SDK for OpenCL Offline Compiler searches for header files during kernel compilation, include the `-I` *<directory>* option in your `aoc` command.

If the header files are in the same directory as your kernel, you do not need to include the `-I` *<directory>* option in your `aoc` command. The offline compiler automatically searches the current folder or directory for header files.

- At a command prompt, invoke the `aoc -I` *`<directory>`* *`<your_kernel_filename>`*`.cl` command.

  Caution: For Windows systems, ensure that your include path does not contain any trailing slashes. The offline compiler considers a trailing forward slash (/) or backward slash (\) as illegal.

  The offline compiler generates an error message if you invoke the `aoc` command in the following manner:

  `aoc -I` *`<drive>`*`\`*`<folder>`*`\ ... \`*`<subfolder>`*`\` *`<your_kernel_filename>`*`.cl`

  or

  `aoc -I` *`<drive>`*`/`*`<folder>`*`/ ... /`*`<subfolder>`*`/` *`<your_kernel_filename>`*`.cl`

  The correct way to specify the include path is as follows:

  `aoc -I` *`<drive>`*`\`*`<folder>`*`\ ... \`*`<subfolder>`* *`<your_kernel_filename>`*`.cl`

  or

  `aoc -I` *`<drive>`*`/`*`<folder>`*`/ ... /`*`<subfolder>`* *`<your_kernel_filename>`*`.cl`

1.8.4 Specifying the Name of an Intel FPGA SDK for OpenCL Offline Compiler Output File (-o <filename>)

To specify the name of a `.aoco` file or a `.aocx` file, include the `-o` *`<filename>`* option in your `aoc` command.

- If you implement the multistep compilation flow, specify the names of the output files in the following manner:

  a. To specify the name of the `.aoco` file that the offline compiler creates during an intermediate compilation step, invoke the `aoc -c -o` *`<your_object_filename>`*`.aoco` *`<your kernel_filename>`*`.cl` command.

  b. To specify the name of the `.aocx` file that the offline compiler creates during the final compilation step, invoke the `aoc -o` *`<your_executable_filename>`*`.aocx` *`<your_object_filename>`*`.aoco` command.

- If you implement the one-step compilation flow, specify the name of the `.aocx` file by invoking the `aoc -o` *`<your_executable_filename>`*`.aocx` *`<your_kernel_filename>`*`.cl` command.

1.8.5 Compiling a Kernel for a Specific FPGA Board (--board <board_name>)

To compile your OpenCL kernel for a specific FPGA board, include the `--board` *`<board_name>`* option in the `aoc` command.

To compile a kernel for a specific board in your Custom Platform, you must first set the environment variable *AOCL_BOARD_PACKAGE_ROOT* to point to the location of your Custom Platform.

Attention: If you want to program multiple FPGA devices, you may select board types that are available in the same Custom Platform because *AOCL_BOARD_PACKAGE_ROOT* only points to the location of one Custom Platform.

When you compile your kernel by including the `--board` *<board_name>* option in the `aoc` command, the Intel FPGA SDK for OpenCL Offline Compiler defines the preprocessor macro `AOCL_BOARD_`*<board_name>* to be 1, which allows you to compile device-optimized code in your kernel.

1. To obtain the names of the available FPGA boards in your Custom Platform, invoke the `aoc --list-boards` command.

   For example, the offline compiler generates the following output:

   ```
   Board List:
   FPGA_board_1
   ```

   where `FPGA_board_1` is the *<board_name>*.

2. To compile your OpenCL kernel for FPGA_board_1, invoke the `aoc --board FPGA_board_1` *<your_kernel_filename>*`.cl` command.
   The offline compiler defines the preprocessor macro `AOCL_BOARD_FPGA_board_1` to be 1 and compiles kernel code that targets FPGA_board_1.

Tip: To readily identify compiled kernel files that target a specific FPGA board, Intel recommends that you rename the kernel binaries by including the `-o` option in the `aoc` command.

To target your kernel to FPGA_board_1 in the one-step compilation flow, invoke the following command:

`aoc --board FPGA_board_1` *<your_kernel_filename>*`.cl -o` *<your_executable_filename>*`_FPGA_board_1.aocx`

To target your kernel to FPGA_board_1 in the multistep compilation flow, perform the following tasks:

1. Invoke the following command to generate the `.aoco` file:

   `aoc -c --board FPGA_board_1` *<your_kernel_filename>*`.cl -o` *<my_object_filename>*`_FPGA_board_1.aoco`

2. Invoke the following command to generate the `.aocx` file:

   `aoc --board FPGA_board_1` *<your_object_filename>*`_FPGA_board_1.aoco -o` *<your_executable_filename>*`_FPGA_board_1.aocx`

If you have an accelerator board consisting of two FPGAs, each FPGA device has an equivalent "board" name (for example, board_fpga_1 and board_fpga_2). To target a `kernel_1.cl` to board_fpga_1 and a `kernel_2.cl` to board_fpga_2, invoke the following commands:

```
aoc --board board_fpga1 kernel_1.cl

aoc --board board_fpga2 kernel_2.cl
```

Related Links

[Specifying the Name of an Intel FPGA SDK for OpenCL Offline Compiler Output File (-o <filename>) on page 97](navigation)

To specify the name of a `.aoco` file or a `.aocx` file, include the `-o <filename>` option in your `aoc` command.

1.8.6 Resolving Hardware Generation Fitting Errors during Kernel Compilation (--high-effort)

Sometimes, OpenCL kernel compilation fails during the hardware generation stage because the design fails to meet fitting constraints. In this case, recompile the kernel using the `--high-effort` option of the `aoc` command.

When kernel compilation fails because of a fitting constraint problem, the Intel FPGA SDK for OpenCL Offline Compiler displays the following error message:

```
Error: Kernel fit error, recommend using --high-effort.
Error: Cannot fit kernel(s) on device
```

- To overcome this problem, recompile your kernel by invoking the following command:

  `aoc --high-effort <your_kernel_filename>.cl`

  After you invoke the command, the offline compiler displays the following message:

  ```
  High-effort hardware generation selected, compile time may increase
  significantly.
  ```

The offline compiler will make three attempts to recompile your kernel and generate hardware. Modify your kernel if compilation still fails after the `--high-effort` attempt.

1.8.7 Defining Preprocessor Macros to Specify Kernel Parameters (-D <macro_name>)

The Intel FPGA SDK for OpenCL Offline Compiler supports preprocessor macros that allow you to pass macro definitions and compile code on a conditional basis.

- To pass a preprocessor macro definition to the offline compiler, invoke the `aoc -D <macro_name> <kernel_filename>.cl` command.

- To override the existing value of a defined preprocessor macro, invoke the `aoc -D <marco_name>=<value> <kernel_filename>.cl` command.

  Consider the following code snippet for the kernel `sum`:

```
#ifndef UNROLL_FACTOR
  #define UNROLL_FACTOR 1
#endif

__kernel void sum (__global const int * restrict x,
                   __global int * restrict sum)
{
  int accum = 0;

  #pragma unroll UNROLL_FACTOR
  for(size_t i = 0; i < 4; i++)
  {
    accum += x[i + get_global_id(0) * 4];
  }
  sum[get_global_id(0)] = accum;
}
```

  To override the `UNROLL_FACTOR` of 1 and set it to 4, invoke the `aoc -D UNROLL_FACTOR=4 sum.cl` command. Invoking this command is equivalent to replacing the line `#define UNROLL_FACTOR 1` with `#define UNROLL_FACTOR 4` in the `sum` kernel source code.

- To use preprocessor macros to control how the offline compiler optimizes your kernel without modifying your kernel source code, invoke the `aoc -o <hardware_filename>.aocx -D <macro_name>=<value> <kernel_filename>.cl`

  Where:

  `-o` is the offline compiler option you use to specify the name of the `.aocx` file that the offline compiler generates.

  <hardware_filename> is the name of the `.aocx` file that the offline compiler generates using the preprocessor macro value you specify.

  Tip: To preserve the results from both compilations on your file system, compile your kernels as separate binaries by using the `-o` flag of the `aoc` command.

  For example, if you want to compile the same kernel multiple times with required work-group sizes of 64 and 128, you can define a `WORK_GROUP_SIZE` preprocessor macro for the kernel attribute `reqd_work_group_size`, as shown below:

```
__attribute__((reqd_work_group_size(WORK_GROUP_SIZE,1,1)))
__kernel void myKernel(...)
for (size_t i = 0; i < 1024; i++)
{
      // statements
}
```

  Compile the kernel multiple times by typing the following commands:

  `aoc -o myKernel_64.aocx -D WORK_GROUP_SIZE=64 myKernel.cl`

  `aoc -o myKernel_128.aocx -D WORK_GROUP_SIZE=128 myKernel.cl`

1.8.8 Generating Compilation Progress Report (-v)

To direct the Intel FPGA SDK for OpenCL Offline Compiler to report on the progress of a compilation, include the `-v` option in your `aoc` command.

- To direct the offline compiler to report on the progress of a full compilation, invoke the `aoc -v <your_kernel_filename>.cl` command.

  The offline compiler generates a compilation progress report similar to the following example:

  ```
  aoc: Environment checks are completed successfully.
  You are now compiling the full flow!!
  aoc: Selected target board s5_net
  aoc: Running OpenCL parser....
  aoc: OpenCL parser completed successfully.
  aoc: Compiling....
  aoc: Linking with IP library ...
  aoc: First stage compilation completed successfully.
  aoc: Setting up project for CvP revision flow....
  aoc: Hardware generation completed successfully.
  ```

- To direct the offline compiler to report on the progress of an intermediate compilation step that does not build hardware, invoke the `aoc -c -v <your_kernel_filename>.cl` command.

  The offline compiler generates a compilation progress report similar to the following example:

  ```
  aoc: Environment checks are completed successfully.
  aoc: Selected target board s5_net
  aoc: Running OpenCL parser....
  aoc: OpenCL parser completed successfully.
  aoc: Compiling....
  aoc: Linking with IP library ...
  aoc: First stage compilation completed successfully.
  aoc: To compile this project, run "aoc <your_kernel_filename>.aoco"
  ```

- To direct the offline compiler to report on the progress of a compilation for emulation, invoke the `aoc -march=emulator -v <your_kernel_filename>.cl` command.

  The offline compiler generates a compilation progress report similar to the following example:

  ```
  aoc: Environment checks are completed successfully.
  You are now compiling the full flow!!
  aoc: Selected target board s5_net
  aoc: Running OpenCL parser....ex
  aoc: OpenCL parser completed successfully.
  aoc: Compiling for Emulation ....
  aoc: Emulator Compilation completed successfully.
  Emulator flow is successful.
  ```

Related Links

- [Compiling Your Kernel without Building Hardware (-c)](#) on page 96
  To direct the Intel FPGA SDK for OpenCL Offline Compiler to compile your OpenCL kernel and generate a Quartus Prime hardware design project without creating a hardware configuration file, include the `-c` option in your `aoc` command.

- • [Emulating and Debugging Your OpenCL Kernel](#) on page 105
  The Intel FPGA SDK for OpenCL Emulator assesses the functionality of your kernel.

1.8.9 Displaying the Estimated Resource Usage Summary On-Screen (--report)

By default, the Intel FPGA SDK for OpenCL Offline Compiler estimates hardware resource usage during compilation . The offline compiler factors in the usage of external interfaces such as PCIe, memory controller, and DMA engine in its calculations. During kernel compilation, the offline compiler generates an estimated resource usage summary in the `<your_kernel_filename>.log` file within the `<your_kernel_filename>` directory. To review the estimated resource usage summary on-screen, include the `--report` option in the `aoc` command.

You can review the estimated resource usage summary without performing a full compilation. To review the summary on-screen prior to generating the hardware configuration file, include the `-c` option in your `aoc` command.

- • At a command prompt, invoke the `aoc -c <your_kernel_filename>.cl --report` command.
  The offline compiler generates an output similar to the following example:

```
+-----------------------------------------------------------------------+
; Estimated Resource Usage Summary                                      ;
+---------------------------------------+-------------------------------+
; Resource                              + Usage                         ;
+---------------------------------------+-------------------------------+
; Logic utilization                     ;    35%                        ;
; ALUTs                                 ;    22%                        ;
; Dedicated logic registers             ;    15%                        ;
; Memory blocks                         ;    29%                        ;
; DSP blocks                            ;     0%                        ;
+---------------------------------------+-------------------------------;
```

Related Links

[Compiling Your Kernel without Building Hardware (-c)](#) on page 96
To direct the Intel FPGA SDK for OpenCL Offline Compiler to compile your OpenCL kernel and generate a Quartus Prime hardware design project without creating a hardware configuration file, include the `-c` option in your `aoc` command.

1.8.10 Suppressing Warning Messages from the Intel FPGA SDK for OpenCL Offline Compiler (-W)

To suppress all warning messages, include the `-W` option in your `aoc` command.

- • At a command prompt, invoke the `aoc -W <your_kernel_filename>.cl` command.

1.8.11 Converting Warning Messages from the Intel FPGA SDK for OpenCL Offline Compiler into Error Messages (-Werror)

To convert all warning messages into error messages, include the `-Werror` option in your `aoc` command.

- At a command prompt, invoke the `aoc -Werror`
  `<your_kernel_filename>.cl` command.

1.8.12 Removing Debug Data from Compiler Reports and Source Code from the .aocx File (-g0)

By default, the Intel FPGA SDK for OpenCL Offline Compiler includes source information in compiler reports and embeds the source code into the `.aocx` binary when it compiles the `.cl` or `.aoco` file. Include the `-g0` option in the `aoc` command to remove source information from the compiler reports and to remove source code and customer IP information from the `.aocx` file.

- To remove source information in reports and remove source code and customer IP information from the `.aocx` file, invoke the `aoc -g0`
  `<your_kernel_filename>.cl` command.

1.8.13 Disabling Burst-Interleaving of Global Memory (--no-interleaving <global_memory_type>)

The Intel FPGA SDK for OpenCL Offline Compiler cannot burst-interleave global memory across different memory types. You can disable burst-interleaving for all global memory banks of the same type and manage them manually by including the `--no-interleaving <global_memory_type>` option in your `aoc` command. Manual partitioning of memory buffers overrides the default burst-interleaved configuration of global memory.

Caution: The `--no-interleaving` option requires a global memory type parameter. If you do not specify a memory type, the offline compiler issues an error message.

- To direct the offline compiler to disable burst-interleaving for the default global memory, invoke the `aoc <your_kernel_filename>.cl --no-interleaving default` command.

  Your accelerator board might include multiple global memory types. To identify the default global memory type, refer to board vendor's documentation for your Custom Platform.

- For a heterogeneous memory system, to direct the offline compiler to disable burst-interleaving of a specific global memory type, perform the following tasks:

  a. Consult the `board_spec.xml` file of your Custom Platform for the names of the available global memory types (for example, DDR and quad data rate (QDR)).

  b. To disable burst-interleaving for one of the memory types (for example, DDR), invoke the `aoc <your_kernel_filename>.cl --no-interleaving DDR` command.
     The offline compiler enables manual partitioning for the DDR memory bank, and configures the other memory bank in a burst-interleaved fashion.

  c. To disable burst-interleaving for more than one type of global memory buffers, include a `--no-interleaving <global_memory_type>` option for each global memory type.

     For example, to disable burst-interleaving for both DDR and QDR, invoke the `aoc <your_kernel_filename>.cl --no-interleaving DDR --no-interleaving QDR` command.

Caution: Do not pass a buffer as kernel arguments that associate it with multiple memory technologies.

1.8.14 Configuring Constant Memory Cache Size (--const-cache-bytes <N>)

Include the `--const-cache-bytes <N>` flag in your `aoc` command to direct the Intel FPGA SDK for OpenCL Offline Compiler to configure the constant memory cache size (rounded up to the closest power of 2).

The default constant cache size is 16 kB.

- To configure the constant memory cache size, invoke the `aoc --const-cache-bytes <N> <your_kernel_filename>.cl` command, where *<N>* is the cache size in bytes.

  For example, to configure a 32 kB cache during compilation of the OpenCL kernel `myKernel.cl`, invoke the `aoc --const-cache-bytes 32768 myKernel.cl` command.

  Note: This argument has no effect if none of the kernels uses the `__constant` address space.

1.8.15 Relaxing the Order of Floating-Point Operations (--fp-relaxed)

Include the `--fp-relaxed` option in your `aoc` command to direct the Intel FPGA SDK for OpenCL Offline Compiler to relax the order of arithmetic floating-point operations using a balanced tree hardware implementation.

Implementing a balanced tree structure leads to more efficient hardware at the expense of numerical variation in results.

Caution: To implement this optimization control, your program must be able to tolerate small variations in the floating-point results.

- To direct the offline compiler to execute a balanced tree hardware implementation, invoke the `aoc --fp-relaxed` *`<your_kernel_filename>`*`.cl` command.

1.8.16 Reducing Floating-Point Rounding Operations (--fpc)

Include the `--fpc` option in your `aoc` command to direct the Intel FPGA SDK for OpenCL Offline Compiler to remove intermediary floating-point rounding operations and conversions whenever possible, and to carry additional bits to maintain precision.

Implementing this optimization control also changes the rounding mode. It rounds towards zero only at the end of a chain of floating-point arithmetic operations (that is, multiplications, additions, and subtractions).

- To direct the offline compiler to reduce the number of rounding operations, invoke the `aoc --fpc` *`<your_kernel_filename>`*`.cl` command.

1.9 Emulating and Debugging Your OpenCL Kernel

The Intel FPGA SDK for OpenCL Emulator assesses the functionality of your kernel.

The Intel FPGA SDK for OpenCL Emulator generates a `.aocx` file that executes on x86-64 Windows or Linux host. This feature allows you to emulate the functionality of your kernel and iterate on your design without executing it on the actual FPGA each time. For Linux platform, you can also use the Emulator to perform functional debug.

Caution: Emulation does not support cross-compilation to ARM processor. To run emulation on a design that targets an SoC, emulate on a non-SoC board (for example, `/board/s5_ref`). When you are satisfied with the emulation results, you may target your design on an SoC board for subsequent optimization steps.

1. Modifying Channels Kernel Code for Emulation on page 106
   To emulate applications with a channel that reads or writes to an I/O channel, modify your kernel to add a read or write channel that replaces the I/O channel, and make the source code that uses it is conditional.

2. Compiling a Kernel for Emulation (-march=emulator) on page 107
   To compile an OpenCL kernel for emulation, include the `-march=emulator` option in your `aoc` command.

3. Emulating Your OpenCL Kernel on page 109
   To emulate your OpenCL kernel, run the emulation `.aocx` file on the platform on which you build your kernel.

4. Debugging Your OpenCL Kernel on Linux on page 110
   For Linux systems, you can direct the Intel FPGA SDK for OpenCL Emulator to run your OpenCL kernel in the debugger and debug it functionally as part of the host application.

5. Limitations of the Intel FPGA SDK for OpenCL Emulator on page 111
   The Intel FPGA SDK for OpenCL Emulator feature has some limitations.

1.9.1 Modifying Channels Kernel Code for Emulation

The Emulator emulates kernel-to-kernel channels. It does not support the emulation of I/O channels that interface with input or output features of your FPGA board. To emulate applications with a channel that reads or writes to an I/O channel, modify your kernel to add a read or write channel that replaces the I/O channel, and make the source code that uses it is conditional.

The Intel FPGA SDK for OpenCL does not set the `EMULATOR` macro definition. You must set it manually either from the command line or in the source code.

Consider the following kernel example:

```
channel unlong4 inchannel __attribute__((io("eth0_in")));

__kernel void send (int size)
{
    for (unsigned i=0; i < size; i++)
    {
        ulong4 data = read_channel_intel(inchannel);
        //statements
    }
}
```

To enable the Emulator to emulate a kernel with a channel that interfaces with an I/O channel, perform the following tasks:

1. Modify the kernel code in one of the following manner:

   — Add a matching `write_channel_intel` call such as the one shown below.

   ```
   #ifdef EMULATOR

   __kernel void io_in (__global char * restrict arr, int size)
   {
       for (unsigned i=0; i<size; i++)
       {
           ulong4 data = arr[i]; //arr[i] being an alternate data source
           write_channel_intel(inchannel, data);
       }
   }
   #endif
   ```

   — Replace the I/O channel access with a memory access, as shown below:

   ```
   __kernel void send (int size)
   {
       for (unsigned i=0; i < size; i++)
       {
           #ifndef EMULATOR

               ulong4 data = read_channel_intel(inchannel);

           #else
               ulong4 data = arr[i]; //arr[i] being an alternate data source

           #endif
           //statements
       }
   }
   ```

2. Modify the host application to create and start this conditional kernel during emulation.

Related Links

> Include an `io` attribute in your channel declaration to declare a special I/O channel to interface with input or output features of an FPGA board.

1.9.1.1 Emulating a Kernel that Passes Pipes or Channels by Reference

The Intel FPGA SDK for OpenCL Emulator supports a kernel that passes pipes or channels by reference.

For example, you may emulate a kernel that has the following structure:

```
void my_function (pipe uint * pipe_ref,
                  __global uint * dst, int i)
{
    read_pipe (*pipe_ref, &dst[i]);
}

__kernel void
consumer (__global uint * restrict dst,
          read_only pipe uint __attribute__((blocking)) c0)
{
    for (int i=0;i<5;i++)
    {
        my_function( &c0, dst, i );
    }
}
```

1.9.1.2 Emulating Channel Depth

When you compile your OpenCL kernel for emulation, the default channel depth is different from the default channel depth generated when your kernel is compiled for hardware. You can change this behavior when you compile your kernel for emulation with the `--emulator-channel-depth-model` option.

The the `--emulator-channel-depth-model` compiler option can take the following values:

default	Channels with an explicit depth attribute have their specified depth. Channels without a specified depth are given a default channel depth that is chosen to provide the fastest execution time for your kernel emulation.
strict	All channel depths in the emulation are given a depth that matches the depth given for the FPGA compilation.
ignore-depth	All channels are given a channel depth chosen to provide the fastest execution time for your kernel emulation. Any explicitly set channel depth attribute is ignored.

1.9.2 Compiling a Kernel for Emulation (-march=emulator)

To compile an OpenCL kernel for emulation, include the `-march=emulator` option in your `aoc` command.

- Before you perform kernel emulation, perform the following tasks:
  - — Install a Custom Platform from your board vendor for your FPGA accelerator boards.
  - — Verify that the environment variable *AOCL_BOARD_PACKAGE_ROOT* points to the location of the Custom Platform. Alternatively, if your kernel targets a board from an Intel FPGA SDK for OpenCL Reference Platform, set *AOCL_BOARD_PACKAGE_ROOT* to the path of the Reference Platform (for example, `ALTERAOCLSDKROOT/board/<Reference_Platform_name>`).
  - — Verify that the environment variable *QUARTUS_ROOTDIR_OVERRIDE* points to the correct edition of the Quartus Prime software.
    - For non-Arria 10 devices, *QUARTUS_ROOTDIR_OVERRIDE* points to the installation directory of the Quartus Prime Standard Edition software.
    - For Arria 10 devices, *QUARTUS_ROOTDIR_OVERRIDE* points to the installation directory of the Quartus Prime Pro Edition software.
- To emulate your kernels on Windows systems, you need the Microsoft linker and additional compilation time libraries. Verify that the *PATH* environment variable setting includes all the paths described in the *Setting the Intel FPGA SDK for OpenCL User Environment Variables* section of the *Intel FPGA SDK for OpenCL Getting Started Guide*.

  The *PATH* environment variable setting must include the path to the `LINK.EXE` file in Microsoft Visual Studio.

- Ensure that your *LIB* environment variable setting includes the path to the Microsoft compilation time libraries.

  The compilation time libraries are available with Microsoft Visual Studio.

- Verify that the *LD_LIBRARY_PATH* environment variable setting includes all the paths described in the *Setting the Intel FPGA SDK for OpenCL User Environment Variables* section in the *Intel FPGA SDK for OpenCL Getting Started Guide*.

- To create kernel programs that are executable on x86-64 host systems, invoke the `aoc -march=emulator <your_kernel_filename>.cl` command.

- To compile a kernel for emulation that targets a specific board, invoke the `aoc -march=emulator --board <board_name> <your_kernel_filename>.cl` command.

- For Linux systems, the Intel FPGA SDK for OpenCL Offline Compiler offers symbolic debug support for the debugger.

  The offline compiler's debug support allows you to pinpoint the origins of functional errors in your kernel source code.

Related Links

- [Compiling a Kernel for a Specific FPGA Board (--board <board_name>)](#) on page 97

  To compile your OpenCL kernel for a specific FPGA board, include the `--board <board_name>` option in the `aoc` command.

- [Setting the Intel FPGA SDK for OpenCL User Environment Variables (Windows)](#)
- [Setting the Intel FPGA SDK for OpenCL User Environment Variables (Linux)](#)

1.9.3 Emulating Your OpenCL Kernel

To emulate your OpenCL kernel, run the emulation `.aocx` file on the platform on which you build your kernel.

To emulate your kernel, perform the following steps:

1. Run the utility command `aocl linkflags` to find out which libraries are necessary for building a host application. The software lists the libraries for both emulation and regular kernel compilation flows.

2. Build a host application and link it to the libraries from Step 1.

   Attention: To emulate multiple devices alongside other OpenCL SDKs, link your host application to the Khronos ICD Loader Library *before* linking it to the host runtime libraries. Link the host application to the ICD Loader Library by modifiying the `Makefile` for the host application. Refer to *Linking Your Host Application to the Khronos ICD Loader Library* for more information.

3. If necessary, move the `<your_kernel_filename>`.aocx file to a location where the host can find easily, preferably the current working directory.

4. To run the host application for emulation:

   − For Windows, first define the number of emulated devices by invoking the `set` `CL_CONTEXT_EMULATOR_DEVICE_ALTERA=<number_of_devices>` command and then run the host application.

   After you run the host application, invoke `set` `CL_CONTEXT_EMULATOR_DEVICE_ALTERA=` to unset the variable.

   − For Linux, invoke the `env` `CL_CONTEXT_EMULATOR_DEVICE_ALTERA=<number_of_devices>` `<host_application_filename>` command.

   This command specifies the number of identical emulation devices that the Emulator needs to provide.

5. If you change your host or kernel program and you want to test it, only recompile the modified host or kernel program and then rerun emulation.

Each invocation of the emulated kernel creates a shared library copy called `<process_ID>`-libkernel.so in a default temporary directory, where `<process_ID>` is a unique numerical value assigned to each emulation run. You may override the default directory by setting the *TMP* or *TEMP* environment variable on Windows, or setting *TMPDIR* on Linux.

Related Links

- Displaying Information on OpenCL Host Runtime and MMD Libraries (link-config or linkflags) on page 85
  To display a list of flags necessary for linking a host application with OpenCL host runtime and MMD libraries, invoke the `link-config` or `linkflags` utility command.

- Linking Your Host Application to the Khronos ICD Loader Library on page 86
  The Intel FPGA SDK for OpenCL supports the OpenCL ICD extension from the Khronos Group.

1.9.4 Debugging Your OpenCL Kernel on Linux

For Linux systems, you can direct the Intel FPGA SDK for OpenCL Emulator to run your OpenCL kernel in the debugger and debug it functionally as part of the host application. The debugging feature allows you to debug the host and the kernel seamlessly. You can step through your code, set breakpoints, and examine and set variables.

Prior to debugging your kernel, you must perform the following tasks:

1. During program execution, the debugger cannot step from the host code to the kernel code. You must set a breakpoint before the actual kernel invocation by adding these lines:

   a. `break <your_kernel>`

      This line sets a breakpoint before the kernel.

   b. `continue`

      If you have not begun debugging your host, then type `start` instead.

2. The kernel is loaded as a shared library immediately before the host loads the kernels. The debugger does not recognize the kernel names until the host actually loads the kernel functions. As a result, the debugger will generate the following warning for the breakpoint you set before the execution of the first kernel:

   `Function "<your_kernel>" not defined.`

   `Make breakpoint pending on future shared library load? (y or [n])`

   Answer `y`. After initial program execution, the debugger will recognize the function and variable names, and line number references for the duration of the session.

Caution: The Emulator uses the OpenCL runtime to report some error details. For emulation, the runtime uses a default print out callback when you initialize a context via the `clCreateContext` function.

Note: Kernel debugging is independent of host debugging. Debug your host code in existing tools such as Microsoft Visual Studio Debugger for Windows and GDB for Linux.

To compile your OpenCL kernel for debugging, perform the following steps:

1. To generate a `.aocx` file for debugging that targets a specific accelerator board, invoke the `aoc -march=emulator <your_kernel_filename>.cl --board <board_name>` command.

> **Attention:** Specify the name of your FPGA board when you run your host
> application. To verify the name of the target board for which you
> compile your kernel, invoke the `aoc -march=emulator -v`
> `<your_kernel_filename>.cl` command. The Intel FPGA SDK for
> OpenCL Offline Compiler will display the name of the target FPGA
> board.

2. Run the utility command `aocl linkflags` to find out the additional libraries
   necessary to build a host application that supports kernel debugging.

3. Build a host application and link it to the libraries from Step 2.

4. Ensure that the `<your_kernel_filename>.aocx` file is in a location where the
   host can find it, preferably the current working directory.

5. To run the application, invoke the command `env`
   `CL_CONTEXT_EMULATOR_DEVICE_ALTERA=<number_of_devices>` gdb --
   `args <your_host_program_name>`, where `<number_of_devices>` is the
   number of identical emulation devices that the Emulator needs to provide.

6. If you change your host or kernel program and you want to test it, only recompile
   the modified host or kernel program and then rerun the debugger.

Related Links

- Compiling a Kernel for a Specific FPGA Board (--board <board_name>) on page 97
  To compile your OpenCL kernel for a specific FPGA board, include the `--`
  `board <board_name>` option in the `aoc` command.

- Generating Compilation Progress Report (-v) on page 101
  To direct the Intel FPGA SDK for OpenCL Offline Compiler to report on the
  progress of a compilation, include the `-v` option in your `aoc` command.

- Displaying Information on OpenCL Host Runtime and MMD Libraries (link-config or
  linkflags) on page 85
  To display a list of flags necessary for linking a host application with OpenCL
  host runtime and MMD libraries, invoke the `link-config` or `linkflags`
  utility command.

1.9.5 Limitations of the Intel FPGA SDK for OpenCL Emulator

The Intel FPGA SDK for OpenCL Emulator feature has some limitations.

1. Execution model

   The Emulator supports the same compilation modes as the FPGA variant. As a
   result, you must call the `clCreateProgramBinary` function to create
   `cl_program` objects for emulation.

2. Concurrent execution

   Modeling of concurrent kernel executions has limitations. During execution, the
   Emulator does not actually run interacting work-items in parallel. Therefore, some
   concurrent execution behaviors, such as different kernels accessing global
   memory without a barrier for synchronization, might generate inconsistent
   emulation results between executions.

3. Kernel performance

The `.aocx` file that you generate for emulation does not include any optimizations. Therefore, it might execute at a significantly slower speed than what an optimized kernel might achieve. In addition, because the Emulator does not implement actual parallel execution, the execution time multiplies with the number of work-items that the kernel executes.

4. The Emulator executes the host runtime and the kernels in the same address space. Certain pointer or array usages in your host application might cause the kernel program to fail, and vice versa. Example usages include indexing external allocated memory and writing to random pointers. You may use memory leak detection tools such as Valgrind to analyze your program. However, the host might encounter a fatal error caused by out-of-bounds write operations in your kernel, and vice versa.

5. Emulation of channel behavior has limitations, especially for conditional channel operations where the kernel does not call the channel operation in every loop iteration. In these cases, the Emulator might execute channel operations in a different order than on the hardware.

1.10 Reviewing Your Kernel's report.html File

Attention: The `analyze-area` Intel FPGA SDK for OpenCL utility option has been deprecated. To view your kernel's estimated area usage, refer to the `report.html` file.

For reference information on the deprecated area report, refer to the *Review Your Kernel's Area Report to Identify Inefficiencies in Resource Usage* section in version 16.0 of the *Altera SDK for OpenCL Best Practices Guide*.

After compiling your OpenCL kernel, the Intel FPGA SDK for OpenCL Offline Compiler automatically generates an HTML report that analyzes various aspects of your kernel, such as area, loop structure, memory usage, and kernel pipeline.
To launch the HTML report, open the `report.html` file in the `<your_kernel_filename>/reports` directory.

For more information on the HTML report, refer to the *Review Your Kernel's report.html File* section in the *Intel FPGA SDK for OpenCL Best Practices Guide*.

Related Links

- Review Your Kernel's report.html File
- Altera SDK for OpenCL Best Practices Guide version 16.0

1.11 Profiling Your OpenCL Kernel

The Intel FPGA SDK for OpenCL Profiler measures and reports performance data collected during OpenCL kernel execution on the FPGA. The SDK's Profiler relies on performance counters to gather kernel performance data. You can then review performance data in the Profiler GUI.

1. Instrumenting the Kernel Pipeline with Performance Counters (--profile) on page 113
   To instrument the OpenCL kernel pipeline with performance counters, include the `--profile` option of the `aoc` command when you compile your kernel.

2. Launching the Intel FPGA SDK for OpenCL Profiler GUI (report) on page 113

You can use the Intel FPGA SDK for OpenCL Profiler `report` utility command to launch the Profiler GUI.

1.11.1 Instrumenting the Kernel Pipeline with Performance Counters (--profile)

To instrument the OpenCL kernel pipeline with performance counters, include the `--profile` option of the `aoc` command when you compile your kernel.

Attention: Instrumenting the Verilog code with performance counters increases hardware resource utilization (that is, increases FPGA area usage) and typically decreases performance.

- To instrument the Verilog code in the *`<your_kernel_filename>`*`.aocx` file with performance counters, invoke the `aoc --profile` *`<your_kernel_filename>`*`.cl` command.

    Attention: When profiling multiple, different kernels, do not use the same kernel names across different `.aocx` files. If the kernel names are the same, the profile data will be wrong for these kernels.

- Run your host application from a local disk to execute the *`<your_kernel_filename>`*`.aocx` file on your FPGA. During kernel execution, the performance counters throughout the kernel pipeline collects profile information. The host saves the information in a `profile.mon` monitor description file in your current working directory.

    Caution: Because of slow network disk accesses, running the host application from a networked directory might introduce delays between kernel executions. These delays might increase the overall execution time of the host application. In addition, they might introduce delays between kernel launches while the runtime stores profile output data to disk.

1.11.2 Launching the Intel FPGA SDK for OpenCL Profiler GUI (report)

You can use the Intel FPGA SDK for OpenCL Profiler `report` utility command to launch the Profiler GUI. The Profiler GUI allows you to view kernel performance data statistics that the SDK's Profiler collects during kernel execution.

The SDK's Profiler stores performance data in a `profile.mon` file in your current working directory.

- To launch the Profiler GUI, invoke the `aocl report` *`<your_kernel_filename>`*`.aocx profile.mon` [*`<your_kernel_filename>`*`.source`] utility command.

    Important: If you do not specify the `.source` file in the command, the Profiler GUI will not have a source code tab.

1.12 Conclusion

You have now familiarized yourself with the Intel FPGA SDK for OpenCL design flow and the tools available to help you achieve your design goals. For more information on the support statuses of the OpenCL APIs and programming language, refer to *Appendix A: Support Statuses of OpenCL Features*.

For in-depth information on optimizing your OpenCL kernel to maximize performance, refer to the *Intel FPGA SDK for OpenCL Best Practices Guide*.

Related Links

Intel FPGA SDK for OpenCL Best Practices Guide

1.13 Document Revision History

Table 3. Document Revision History of the Intel FPGA SDK for OpenCL Programming Guide

Date	Version	Changes
May 2017	2017.05.08	• Rebranded some functions in code examples as follows: — Rebranded `read_channel_altera` to `read_channel_intel`. — Rebranded `write_channel_altera` to `write_channel_intel`. — Rebranded `read_channel_nb_altera` to `read_channel_nb_intel`. — Rebranded `write_channel_nb_altera` to `write_channel_nb_intel`. — Rebranded `clGetBoardExtensionFuncationAddressAltera` to `clGetBoardExensionFunctionAdressIntelFPGA`. • Added Emulating I/O Channels on page 37. • Added Implementing Arbitrary Precision Integers on page 61. • Added Coalescing Nested Loops on page 21. • Added Specifying a loop initiation interval (II) on page 23. • Added Emulating Channel Depth on page 107. • Added Avalon Streaming (Avalon-ST) Interface on page 124. • Removed all references to `#pragma OPENCL EXTENSION cl_altera_channels : enable` because this pragma is not required to implement channels. • Reorganized information related to heterogeneous memory as follows: — Merged *Specifying Pointer Size in Memory* content into Programming Strategies for Optimizing Local Memory Efficiency on page 25. — Restructured *Allocating OpenCL Buffer for Manual Partitioning of Global Memory* into three topics: • Allocating OpenCL Buffers for Manual Partitioning of Global Memory on page 74 • Partitioning Buffers Across Multiple Interfaces of the Same Memory Type on page 74 • Partitioning Buffers Across Different Memory Types (Heterogeneous Memory) on page 76 — Moved *Specifying Buffer Location in Global Memory* (previously under *Programming Strategies for Optimizing Access Efficiency*) content into Partitioning Buffers Across Different Memory Types (Heterogeneous Memory) on page 76. • Updated Collecting Profile Data During Kernel Execution on page 78 with a warning about the affect of collecting profile data on kernel launch times. • Updated Compiling Your OpenCL Kernel on page 95 with restrictions on compiling an encrypted `.cl` file. • Updated Restrictions and Limitations in RTL Support for the Intel FPGA SDK for OpenCL Library Feature on page 133 to indicate that an RTL module must use a single-input Avalon-ST interface to control inputs.

continued...

Date	Version	Changes
		• Updated topics affected by changes to the OpenCL profiler as follows: — Updated Launching the Intel FPGA SDK for OpenCL Profiler GUI (report) on page 113 with new command options. — Updated Figure 4 on page 10 in Multistep Intel FPGA SDK for OpenCL Design Flow on page 9 to reflect the new command options. • Corrected code example in Implementing Nonblocking Channel Read Extensions on page 34. • Corrected code example in *Channel Execution in Loop with Multiple Work-Items* section of Work-Item Serial Execution of Channels on page 29.
October 2016	2016.10.31	• Rebranded the Altera SDK for OpenCL to Intel FPGA SDK for OpenCL. • Rebranded the Altera Offline Compiler to Intel FPGA SDK for OpenCL Offline Compiler. • Deprecated and removed support for big-endian system, resulting in the following documentation changes: — Removed the topic *Compiling a Kernel for a Big-Endian System (--big-endian)*. — Removed big-endian (64-bit) from the list of architectures that the host application can target. • Added the topic *Displaying the Compilation Environment of an OpenCL Binary* to introduce the `aoc env` command. • Removed *Adding Source References to Optimization Reports (-g)* because the offline compiler automatically includes source information in the compiler reports and enables symbolic debug during emulation on an x86 Linux machine. • Added the topic *Removing Debug Data from Compiler Reports and Source Code from the .aocx File (-g0)* to introduce the `-g0 aoc` command option. • In *Limitations of the Intel FPGA SDK for OpenCL Emulator*, removed the limitation "The Emulator does not support half data type". • In *Linking Your Host Application to the Khronos ICD Loader Library*, provided an update that the Intel-supplied ICD Loader Library supports OpenCL Specification version 1.0 as well as implemented APIs from the OpenCL Specification versions 1.1, 1.2, and 2.0. • In *Managing an FPGA Board*, provided the following updates: — Noted that the SDK supports installation of multiple Custom Platforms. To use the SDK utilities on each board in a multi-board installation, the *AOCL_BOARD_PACKAGE_ROOT* environment variable setting must correspond to the Custom Platform subdirectory of the associated board. — Noted that in a system with multiple Custom Platforms, the host program should use ACD to discover the boards instead of directly linking to the MMD libraries. • Added the topic *Reviewing Your Kernel's report.html File* and included deprecation notice for the `analyze-area` utility option. As a result of introducing the HTML report, removed the following topics: — *Reviewing Your Kernel's Resource Usage Information in the Area Report* — *Accessing the Area Report* — *Layout of the Area Report* • In *Multistep Design Flow*, updated the design steps and the figure *The Multistep Intel FPGA SDK for OpenCL Design Flow* to replace area report with the HTML report, and remove information on enabling `-g`. • In *Inferring a Register*, corrected the text following the code snippet that explained how the offline compiler decide on the implementation of the array in hardware. • In *Linking to the ICD Loader Library* on Windows, updated the text to improve clarity.

<div align="right">***continued...***</div>

Date	Version	Changes

May 2016	2016.05.02	• Added a schematic diagram of the AOCL programming model in the *Altera SDK for OpenCL FPGA Programming Flow* section.
		• Moved the figure *The AOCL FPGA Programming Flow* to the *Altera Offline Compiler Kernel Compilation Flows* section.
		• Updated the figure *The Multistep AOCL Design Flow* and associated text to include the Review Area Report step.
		• Added information on the single-cycle floating-point accumulator feature for single work-item kernels. Refer to the *Single-Cycle Floating-Point Accumulator for Single Work-Item Kernels* section for more information.
		• Added information in the *Emulating Your OpenCL Kernel* section on multi-device support for emulation alongside other OpenCL SDKs using ICD.
		• Included information on the enhanced area report feature:
		  — Added the option to invoke the `analyze-area` AOCL utility command to generate an HTML area report.
		  — Included a topic that describes the layout of the HTML area report.
		• In *Linking to the ICD Loader Library on Windows*, removed `$(AOCL_LDLIBS)` from the code example for the modified `Makefile`.
		• In the *Multiple Work-Item Ordering* sections for channels and pipes, modified the characteristics that the AOCL uses to check whether the channel or pipe call is work-item invariant.

November 2015	2015.11.02	• Added the option to invoke the `aoc` command with no argument to access the Altera Offline Compiler help menu.
		• Updated the *Mutliple Host Threads* section to specify that the OpenCL host runtime is thread-safe.
		• Updated the following figure and sections to reflect multiple kernel source file support:
		  — The figure *The AOCL FPGA Programming Flow* in the *AOCL FPGA Programming Flow* section
		  — The *Compiling Your Kernel to Create Hardware Configuration File* section
		  — The *Compiling Your Kernel without Building Hardware (-c)* section
		• In *Multiple Work-Item Ordering for Channels*, removed misleading text.
		• Updated the *Overview of Channels Implementation* figure.
		• Updated the the following sections on OpenCL pipes:
		  — *Overview of a Pipe Network Implementation* figure in *Overview of the OpenCL Pipe Functions*
		  — Emulation support in *Restrictions in OpenCL Pipes Implementation* section
		  — Replaced erroneous code with the correct syntax
		  — Added link to *Implementing I/O Pipes Using the io Attribute* in *Declaring the Pipe Handle*
		• Added a reminder in *Programming an FPGA via the Host* that you should release an event object after use to prevent excessive memory usage.

May 2015	15.0.0	• In *Guidelines for Naming the Kernel*, added entry that advised against naming an OpenCL kernel `kernel.cl`. • In *Instrumenting the Kernel Pipeline with Performance Counters (--profile)*, specified that you should run the host application from a local disk to avoid potential delays caused by slow network disk accesses. • In *Emulating and Debugging Your OpenCL Kernel*, modified Caution note to indicate that you must emulate a design targeting an SoC on a non-SoC board. • In *Emulating Your OpenCL Kernel*, updated command to run the host application and added instruction for overriding default temporary directory containing `<process_ID>-libkernel.so`. • Introduced the `--high-effort` aoc command flag in *Resolving Hardware Generation Fitting Errors during Kernel Compilation*.

continued...

| Date	Version	Changes
		• In *Enabling Double Precision Floating-Point Operations*, introduced the `OPENCL EXTENSION` pragma for enabling double precision floating-point operations. • Introduced OpenCL pipes support. Refer to *Implementing OpenCL Pipes* (and subsequent subtopics) and *Creating a Pipe Object in Your Host Application* for more information. • In *AOCL Channels Extension: Restrictions*, added code examples to demonstrate how to statically index into arrays of channel IDs. • In *Multiple Host Threads*, added recommendation for synchronizing OpenCL host function calls in a multi-threaded host application. • Introduced ICD and ACD support. Refer to *Linking Your Host Application to the Khronos ICD Loader Library* for more information. • Introduced `clGetBoardExtenstionFunctionAddressAltera` for referencing user-accessible functions. Refer to *Accessing Custom Platform-Specific Functions* for more information.
December 2014	14.1.0	• Reorganized information flow. Information is now presented based on the tasks you might perform using the Altera SDK for OpenCL (AOCL) or the Altera RTE for OpenCL. • Removed information pertaining to the `--util <N>` and `-O3` Altera Offline Compiler (AOC) options. • Added the following information on PLDA QuickUDP IP core licensing in *Compiling Your OpenCL Kernel*: 1. A PLDA QuickUDP IP core license is required for the Stratix V Network Reference Platform or a Custom Platform that uses the QuickUDP IP core. 2. Improper installation of the QuickUDP IP core licence causes compilation to fail with an error message that refers to the QuickTCP IP core. • Added reminder that conditionally shifting a large shift register is not recommended. • Removed the *Emulating Systems with Multiple Devices* section. A new env `CL_CONTEXT_EMULATOR_DEVICE_ALTERA=<number_of_devices>` command is now available for emulating multiple devices. • Removed language support limitation from the *Limitations of the AOCL Emulator* section.
June 2014	14.0.0	• Removed the `--estimate-throughput` and `--sw-dimm-partition` AOC options • Added the `-march=emulator`, `-g`, `--big-endian`, and `--profile` AOC options • `--no-interleaving` needs `<global_memory_type>` argument • `-fp-relaxed=true` is now `--fp-relaxed` • `-fpc=true` is now `--fpc` • For non-SoC devices, `aocl diagnostic` is now `aocl diagnose` and `aocl diagnose <device_name>` • `program` and `flash` need `<device_name>` arguments • Added *Identifying the Device Name of Your FPGA Board* • Added *AOCL Profiler Utility* • Added *AOCL Channels Extension* and associated subsections • Added *Attributes for Channels* • Added *Match Data Layouts of Host and Kernel Structure Data Types* • Added *Register Inference* and *Shift Register Inference* • Added *Channels and Multiple Command Queues* • Added *Shared Memory Accesses for OpenCL Kernels Running on SoCs* • Added *Collecting Profile Data During Kernel Execution* • Added *Emulate and Debug Your OpenCL Kernel* and associated subsections

continued...

Date	Version	Changes
		• Updated *AOC Kernel Compilation Flows* • Updated *-v* • Updated *Host Binary Requirement* • Combined *Partitioning Global Memory Accesses* and *Partitioning Heterogeneous Global Memory Accesses* into the section *Partitioning Global Memory Accesses* • Updated *AOC Allocation Limits* in *Appendix A* • Removed `max_unroll_loops`, `max_share_resources`, `num_share_resources`, and `task` kernel attributes • Added `packed`, and `aligned(<N>)` kernel attributes
December 2013	13.1.1	• Removed the section *-W and -Werror*, and replaced it with two sections: *-W* and *-Werror*. • Updated the following contents to reflect multiple devices support: — The figure *The AOCL FPGA Programming Flow*. — *--list-boards* section. — *-board <board_name>* section. — *AOCL Utilities for Managing an FPGA Board* section. — Added the subsection *Programming Multiple FPGA Devices* under *FPGA Programming*. • The following contents were added to reflect heterogeneous global memory support: — *--no-interleaving* section. — `buffer_location` kernel attribute under *Kernel Pragmas and Attributes*. — *Partitioning Heterogeneous Global Memory Accesses* section. • Modified support status designations in *Appendix: Support Statuses of OpenCL Features*. • Removed information on OpenCL programming language restrictions from the section *OpenCL Programming Language Implementation*, and presented the information in a new section titled *OpenCL Programming Language Restrictions*.
November 2013	13.1.0	• Reorganized information flow. • Updated and renamed *Intel FPGA SDK for OpenCL Compilation Flow* to *AOCL FPGA Programming Flow*. • Added figures *One-Step AOC Compilation Flow* and *Two-Step AOC Compilation Flow*. • Updated the section *Contents of the AOCL Version 13.1*. • Removed the following sections: — *OpenCL Kernel Source File Compilation*. — *Using the Altera Offline Kernel Compiler*. — *Setting Up Your FPGA Board*. — *Targeting a Specific FPGA Board*. — *Running Your OpenCL Application*. — *Consolidating Your Kernel Source Files*. — *Aligned Memory Allocation*. — *Programming the FPGA Hardware*. — *Programming the Flash Memory of an FPGA*. • Updated and renamed *Compiling the OpenCL Kernel Source File* to *AOC Compilation Flows*. • Renamed *Passing File Scope Structures to OpenCL Kernels* to *Use Structure Arguments in OpenCL Kernels*. • Updated and renamed *Augmenting Your OpenCL Kernel by Specifying Kernel Attributes and Pragmas* to *Kernel Pragmas and Attributes*. • Renamed *Loading Kernels onto an FPGA* to *FPGA Programming*.

continued...

| Date	Version	Changes
		• Consolidated *Compiling and Linking Your Host Program*, *Host Program Compilation Settings*, and *Library Paths and Links* into a single section. • Inserted the section *Preprocessor Macros*. • Renamed *Optimizing Global Memory Accesses* to *Partitioning Global Memory Accesses*.

June 2013	13.0 SP1.0	• Added the section *Setting Up Your FPGA Board*.
		• Removed the subsection *Specifying a Target FPGA Board* under *Kernel Programming Considerations*.
		• Inserted the subsections *Targeting a Specific FPGA Board* and *Generating Compilation Reports* under *Compiling the OpenCL Kernel Source File*.
		• Renamed *File Scope __constant Address Space Qualifier* to *__constant Address Space Qualifiers*, and inserted the following subsections:
		  — *Function Scope __constant Variables*.
		  — *File Scope __constant Variables*.
		  — *Points to __constant Parameters from the Host*.
		• Inserted the subsection *Passing File Scope Structures to OpenCL Kernels* under *Kernel Programming Considerations*.
		• Renamed *Modifying Your OpenCL Kernel by Specifying Kernel Attributes and Pragmas* to *Augmenting Your OpenCL Kernel by Specifying Kernel Attributes and Pragmas*.
		• Updated content for the `unroll` pragma directive in the section *Augmenting Your OpenCL Kernel by Specifying Kernel Attributes and Pragmas*.
		• Inserted the subsections *Out-of-Order Command Queues* and *Modifying Host Program for Structure Parameter Conversion* under *Host Programming Considerations*.
		• Updated the sections *Loading Kernels onto an FPGA Using clClreateProgramWithBinary* and *Aligned Memory Allocation*.
		• Updated flash programming instructions.
		• Renamed *Optional Extensions* in *Appendix B* to *Atomic Functions*, and updated its content.
		• Removed *Platform Layer and Runtime Implementation* from *Appendix B*.

May 2013	13.0.1	• Explicit memory fence functions are now supported; the entry is removed from the table OpenCL Programming Language Implementation. • Updated the section Programming the Flash Memory of an FPGA. • Added the section *Modifying Your OpenCL Kernel by Specifying Kernel Attributes and Pragmas* to introduce kernel attributes and pragmas that can be implemented to optimize kernel performance. • Added the section Optimizing Global Memory Accesses to discuss data partitioning. • Removed the section *Programming the FPGA with the aocl program Command* from Appendix A.
May 2013	13.0.0	• Updated compilation flow. • Updated kernel compiler commands. • Included Altera SDK for OpenCL Utility commands. • Added the section *OpenCL Programming Considerations*. • Updated flash programming procedure and moved it to *Appendix A*. • Included a new `clCreateProgramWithBinary` FPGA hardware programming flow. • Moved the hostless `clCreateProgramWithBinary` hardware programming flow to *Appendix A* under the title *Programming the FPGA with the aocl program Command*. • Moved updated information on allocation limits and OpenCL language support to *Appendix B*.
November 2012	12.1.0	Initial release.

2 Intel FPGA SDK for OpenCL Advanced Features

The Intel FPGA SDK for OpenCL provides advanced features you can use to control certain aspects of the design architecture and the Intel FPGA SDK for OpenCL Offline Compiler's behavior.

2.1 OpenCL Library

An OpenCL library is a single file that contains multiple functions. Each function is comprised of data processing logic that works at any clock frequency. You can create an OpenCL library in OpenCL or RTL. You can then include this library file and use the functions inside your OpenCL kernels.

Figure 12.   Overview of Intel FPGA SDK for OpenCL's Library Support

You may use a previously-created library or create your own library. To use an OpenCL library, you do not require in-depth knowledge in hardware design or in the implementation of library components. To create an OpenCL library, you need to create the following files and components:

Table 4.   Necessary Files and Components for Creating an OpenCL Library

File or Component	Description
RTL Components	
RTL source files	Verilog, System Verilog, or VHDL files that define the RTL component.
	Additional files such as Quartus Prime IP File (`.qip`), Synopsys Design Constraints File (`.sdc`), and Tcl Script File (`.tcl`) are not allowed.
	continued...

File or Component	Description
eXtensible Markup Language File (`.xml`)	Describes the properties of the RTL component. The Intel FPGA SDK for OpenCL Offline Compiler uses these properties to integrate the RTL component into the OpenCL pipeline.
Header file (`.h`)	A C-style header file that declares the signatures of function(s) that are implement by the RTL component.
OpenCL emulation model file (`.cl`)	Provides C model for the RTL component that is used only for emulation. Full hardware compilations use the RTL source files.
OpenCL Functions	
OpenCL source files (`.cl`)	Contains definitions of the OpenCL functions. These functions are used during emulation and full hardware compilations.
Header file (`.h`)	A C-style header file that declares the signatures of function(s) that are defined in the OpenCL source files.

Understanding RTL Modules and the OpenCL Pipeline on page 122
   This section provides an overview of how the Intel FPGA SDK for OpenCL Offline Compiler integrates RTL modules into the Intel FPGA SDK for OpenCL pipeline architecture.

Packaging an OpenCL Helper Function File for an OpenCL Library on page 131
   Before creating an OpenCL library file, package each OpenCL source file with helper functions into a `.aoco` file.

Packaging an RTL Component for an OpenCL Library on page 132
   Before creating an OpenCL library file, package each RTL component into a `.aoco` file.

Verifying the RTL Modules on page 135
   The creator of an OpenCL library is responsible for verifying the RTL modules within the library, both as stand-alone entities and as part of an OpenCL system.

Packaging Multiple Object Files into a Library File on page 135
   After creating the `.aoco` files that you want to include into an OpenCL library, package them into a library file by invoking the Intel FPGA SDK for OpenCL `library` utility command option.

Specifying an OpenCL Library when Compiling an OpenCL Kernel on page 135
   To use an OpenCL library in an OpenCL kernel, specify the library file name and directory when you compile the kernel.

Using an OpenCL Library that Works with Simple Functions (Example 1) on page 136
   Intel provides an OpenCL library design example of a simple kernel that uses a library containing RTL implementations of three double-precision functions: `sqrt`, `rsqrt`, and `divide`.

Using an OpenCL Library that Works with External Memory (Example 2) on page 137
   Intel provides an OpenCL library design example of a simple kernel that uses a library containing two RTL modules that communicate with global memory.

OpenCL Library Command-Line Options on page 138
   Both the Intel FPGA SDK for OpenCL Offline Compiler's set of commands and the SDK utility include options you can invoke to perform OpenCL library-related tasks.

Related Links

OpenCL Library Command-Line Options on page 138

Both the Intel FPGA SDK for OpenCL Offline Compiler's set of commands and the SDK utility include options you can invoke to perform OpenCL library-related tasks.

2.1.1 Understanding RTL Modules and the OpenCL Pipeline

The OpenCL library feature allows you to use RTL modules, written in Verilog, SystemVerilog, or VHDL, inside OpenCL kernels. This section provides an overview of how the Intel FPGA SDK for OpenCL Offline Compiler integrates RTL modules into the Intel FPGA SDK for OpenCL pipeline architecture.

You might want to use RTL modules under the following circumstances:

- You want to use optimized and verified RTL modules in OpenCL kernels without rewriting the modules as OpenCL functions.

- You want to implement OpenCL kernel functionality that you cannot express effectively in OpenCL.

Overview: Intel FPGA SDK for OpenCL Pipeline Approach

The following figure depicts the architecture of an Intel FPGA SDK for OpenCL pipeline:

Figure 13.   Parallel Execution Model of AOCL Pipeline Stages

The operations on the right represent the SDK's pipeline implementation of the OpenCL kernel code on the left. Each yellow box is an operation or data value found in the pipeline. The number associated with each operation represents the number of threads in the pipeline.

Assume each level of operation is one stage in the pipeline. At each stage, the Intel FPGA SDK for OpenCL Offline Compiler executes all operations in parallel by the thread existing at that stage. For example, thread 2 executes Load A, Load B, and copies the current global ID (via `gid`) to the next pipeline stage. Similar to the pipelined execution on instructions in reduced instruction set computing (RISC) processors, the SDK's pipeline stages also execute in parallel. The threads will advance to the next pipeline stage only after all the stages have completed execution.
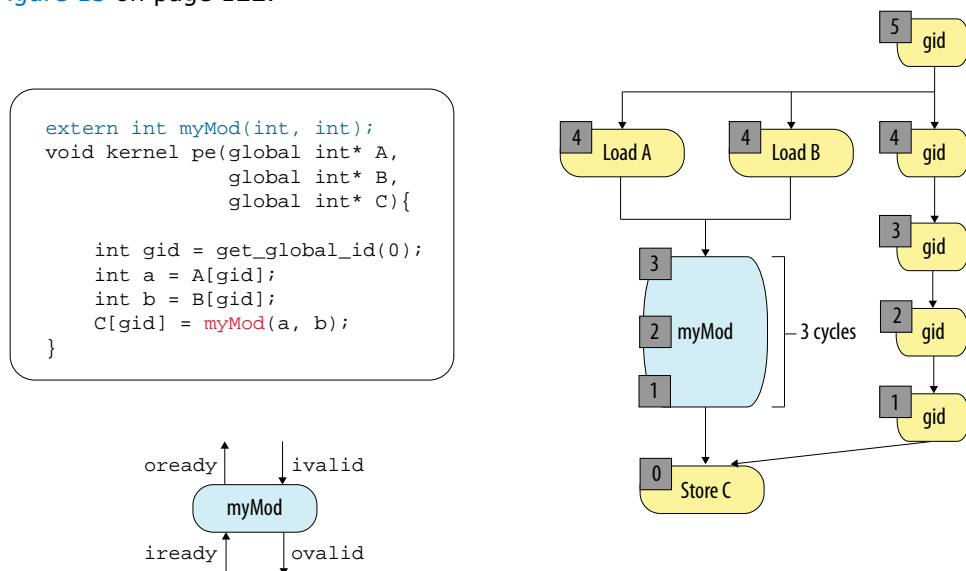
Some operations are capable of stalling the Intel FPGA SDK for OpenCL pipeline. Examples of such operations include variable latency operations like memory load and store operations. To support stalls, ready and valid signals need to propagate throughout the pipeline so that the offline compiler can schedule the pipeline stages. However, ready signals are not necessary if all operations have fixed latency. In these cases, the offline compiler optimizes the pipeline to statically schedule the operations, which significantly reduces the logic necessary for pipeline implementation.

Integration of an RTL Module into the Intel FPGA SDK for OpenCL Pipeline

When you specify an OpenCL library during kernel compilation, the offline compiler integrates the RTL module within the library into the overall pipeline.

Figure 14. **Integration of an RTL Module into an Intel FPGA SDK for OpenCL Pipeline**

This figure depicts the integration of the RTL module `myMod` into the pipeline depicted in Figure 13 on page 122.

The RTL module depicted in Figure 14 on page 123 has a balanced latency where the threads of the RTL module match the number of pipeline stages. A balanced latency allows the threads of the RTL module to execute without stalling the SDK's pipeline.

Setting the latency of the RTL module in the RTL specification file allows the offline compiler to balance the pipeline latency. RTL supports Avalon® Streaming (Avalon-ST) interfaces; therefore, the latency of the RTL module can be variable (that is, not fixed). However, the variability in the latency should be small in order to maximize performance. In addition, specify the latency in the *<RTL module description file name>.xml* specification file so that the RTL module experiences a good approximation of the actual latency in steady state.

2.1.1.1 Stall-Free RTL

The Intel FPGA SDK for OpenCL Offline Compiler can optimize hardware resource usage and performance by removing stall logic around an RTL module with fixed latency.

An RTL module that has a variable latency and uses Avalon-ST input and output signals can wait until input data is ready. Conversely, the Intel FPGA SDK for OpenCL pipeline can stall until it receives valid output data from the RTL module. For an RTL module with a fixed latency, you can remove an RTL stall by modifying the *<RTL module description file name>*.xml specification file, as described below.

To instruct the offline compiler to remove stall logic around the RTL module, if appropriate, set the IS_STALL_FREE attribute under the FUNCTION element to "yes". This modification informs the offline compiler that the RTL module produces valid data every EXPECTED_LATENCY cycle(s). EXPECTED_LATENCY is an attribute you specify in the .xml file under the FUNCTION element. Specify a value for EXPECTED_LATENCY such that the latency equals the number of pipeline stages in the module. An inaccurate EXPECTED_LATENCY value will cause the RTL module to be out of sync with the rest of the pipeline.

A stall-free RTL module might receive an invalid input signal (that is, ivalid is low). In this case, the module ignores the input and produces invalid data on the output. For a stall-free RTL module without an internal state, it might be easier to propagate the invalid input through the module. However, for an RTL module with an internal state, you must handle an ivalid=0 input carefully.

2.1.1.2 Avalon Streaming (Avalon-ST) Interface

The offline compiler expects the RTL module to support Avalon-ST interface with readyLatency = 0, at both input and output.
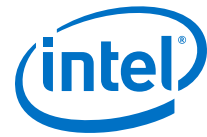
As shown in Figure 14 on page 123, the RTL module must have 4 ports:

- ivalid and iready, as the input Avalon-ST interface
- ovalid and oready, as the output Avalon-ST interface

The following figure shows the timing diagram for input data transfer with back pressure. For more information about Avalon-ST interfaces, see the "Avalon Streaming Interfaces" section in *Avalon Interface Specifications*.

For an RTL module with a fixed latency, the output signals (ovalid and oready) can have constant high values, and the input ready signal (iready) can be ignored.

A stall-free RTL module might receive an invalid input signal (`ivalid` is low). In this case, the module ignores the input and produces invalid data on the output. For a stall-free RTL module without an internal state, it might be easier to propagate the invalid input through the module. However, for an RTL module with an internal state, you must handle an `ivalid = 0` input carefully.

2.1.1.3 RTL Reset and Clock Signals

Resets and clocks of RTL modules are connected to the same clock and reset drivers as the rest of the OpenCL pipeline.

Because of the common clock and reset drivers, an RTL module runs in the same clock domain as the OpenCL kernel. The module is reset only when the OpenCL kernel is first loaded onto the FPGA, either via Intel FPGA SDK for OpenCL `program` utility or the `clCreateProgramwithBinary` host function. In particular, if the host restarts a kernel via successive `clEnqueueNDRangeKernel` or `clEnqueueTask` invocations, the associated RTL modules will not reset in between these restarts.

The following steps outline the process of setting the kernel clock frequency:

1. The Quartus Prime software's Fitter applies an aggressive constraint on the kernel clock.

2. The Quartus Prime software's TimeQuest Timing Analyzer performs static timing analysis to determine the frequency that the Fitter actually achieves.

3. The phase-locked loop (PLL) that drives the kernel clock sets the frequency determined in Step 2 to be the kernel clock frequency.

2.1.1.4 XML Syntax of an RTL Module

This section provides the syntax of a simple XML specification file for an RTL module that implements double-precision square root function. The RTL module is implemented in VHDL with a Verilog wrapper.

The following XML specification file is for an RTL module named `my_fp_sqrt_double` (line 2.5) that implements an OpenCL helper function named `my_sqrtfd` (line 2).

```
  1: <RTL_SPEC>
  2:   <FUNCTION name="my_sqrtfd"
2.5:          module="my_fp_sqrt_double">
  3:     <ATTRIBUTES>
3.5:       <PARAMETER name="WIDTH" value="32"/>
  4:       <IS_STALL_FREE value="yes"/>
  5:       <IS_FIXED_LATENCY value="yes"/>
  6:       <EXPECTED_LATENCY value="31"/>
  7:       <CAPACITY value="1"/>
  8:       <HAS_SIDE_EFFECTS value="no"/>
  9:       <ALLOW_MERGING value="yes"/>
 10:     </ATTRIBUTES>
 11:     <INTERFACE>
 12:       <AVALON port="clock" type="clock"/>
 13:       <AVALON port="resetn" type="resetn"/>
 14:       <AVALON port="ivalid" type="ivalid"/>
 15:       <AVALON port="iready" type="iready"/>
 16:       <AVALON port="ovalid" type="ovalid"/>
 17:       <AVALON port="oready" type="oready"/>
 18:       <INPUT port="datain" width="64"/>
 19:       <OUTPUT port="dataout" width="64"/>
 20:     </INTERFACE>
 21:     <C_MODEL>
 22:       <FILE name="c_model.cl" />
```

```
23:      </C_MODEL>
24:      <REQUIREMENTS>
25:        <FILE name="my_fp_sqrt_double_s5.v" />
26:        <FILE name="fp_sqrt_double_s5.vhd" />
27:      </REQUIREMENTS>
28:    </FUNCTION>
29: </RTL_SPEC>
```

Table 5.   Elements and Attributes in the XML Specification File

XML Element	Description
RTL_SPEC	Top-level element in the XML specification file. There can only be one such top-level element in the file. In this example, the name RTL_SPEC is historic and carries no file-specific meaning.
FUNCTION	Element that defines the OpenCL function that the RTL module implements. The name attribute within the FUNCTION element specifies the function's name. You may have multiple FUNCTION elements, each declaring a different function that you can call from the OpenCL kernel. The same RTL module can implement multiple functions by specifying different parameters.
ATTRIBUTES	Element containing other XML elements that describe various characteristics (for example, latency) of the RTL module. The example RTL module takes one PARAMETER setting named WIDTH, which has a value of 32. Refer to Table 6 on page 127 for more details other ATTRIBUTES-specific elements. *Note:* If you create multiple OpenCL helper functions for different modules, or use the same RTL module with different PARAMETER settings, you must create a separate FUNCTION element for each function.
INTERFACE	Element containing other XML elements that describe the RTL module's interface. The example XML specification file shows the Avalon-ST interface signals that every RTL module must provide (that is, clock, resetn, ivalid, iready, ovalid, and oready). The signal names must match the ones specified in the .xml file. An error will occur during library creation if a signal name is inconsistent.
C_MODEL	Element specifying one or more files that implement OpenCL C model for the function. The model is used only during emulation. However, the C_MODEL element and the associated file(s) must be present when you create the library file.
REQUIREMENTS	Element specifying one or more RTL resource files (that is, .v, .sv, .vhd, .hex, and .mif). The specified paths to these files are relative to the location of the XML specification file. Each RTL resource file becomes part of the associated Qsys component that corresponds to the entire OpenCL system. *Note:* The Intel FPGA SDK for OpenCL library feature does not support .qip files. An Intel FPGA SDK for OpenCL Offline Compiler error will occur if you compile an OpenCL kernel while using a library that includes an unsupported resource file type.

2.1.1.4.1 XML Elements for ATTRIBUTES

In the XML specification file of the RTL module within an Intel FPGA SDK for OpenCL library, there are XML elements under ATTRIBUTES that you can specify to set the module's characteristics.

Table 6. **XML Elements Associated with the ATTRIBUTES Element in the XML Specification File of an RTL Module**

> *Attention:* Except for `IS_STALL_FREE` and `EXPECTED_LATENCY`, all elements have safe values. If you are unsure which value you should specify for an attribute, set it to the safe value. Compiling your kernel with a library that uses safe values will result in functional hardware. However, the hardware might be larger than the actual size.

XML Element	Description
`IS_STALL_FREE`	Instructs the Intel FPGA SDK for OpenCL Offline Compiler to remove all stall logic around the RTL module. Set `IS_STALL_FREE` to `"yes"` to indicate that the module neither generates stalls internally nor can it properly handle incoming stalls. The module simply ignores its stall input. If you set `IS_STALL_FREE` to `"no"`, the module must properly handle all stall and valid signals. *Note:* If you set `IS_STALL_FREE` to `"yes"`, you must also set `IS_FIXED_LATENCY` to `"yes"`. Also, if the RTL module has an internal state, it must properly handle `ivalid=0` inputs. An incorrect `IS_STALL_FREE` setting will lead to incorrect results in hardware.
`IS_FIXED_LATENCY`	Indicates whether the RTL module has a fixed latency. Set `IS_FIXED_LATENCY` to `"yes"` if the RTL module always takes known a number of clock cycles to compute its output. The value you assign to the `EXPECTED_LATENCY` element specifies the number of clock cycles. The safe value for `IS_FIXED_LATENCY` is `"no"`. *Note:* For a given module, you may set `IS_FIXED_LATENCY` to `"yes"` and `IS_STALL_FREE` to `"no"`. Such a module produces its output in a fixed number of clock cycles and handles stall signals properly.
`EXPECTED_LATENCY`	Specifies the expected latency of the RTL module. If you set `IS_FIXED_LATENCY` to `"yes"`, the `EXPECTED_LATENCY` value indicates the number of pipeline stages inside the module. In this case, you must set this value to be the exact latency of the module. Otherwise, the offline compiler will generate incorrect hardware. For a module with variable latency, the offline compiler balances the pipeline around this module to the `EXPECTED_LATENCY` value that you specify. The specified value and the actual latency might differ, which might affect the number of stalls inside the pipeline. However, the resulting hardware will be correct.
`CAPACITY`	Specifies the number of multiple inputs that this module can process simultaneously. You must specify a value for `CAPACITY` if you also set `IS_STALL_FREE="no"` and `IS_FIXED_LATENCY="no"`. Otherwise, you do not need to specify a value for `CAPACITY`. If `CAPACITY` is strictly less than `EXPECTED_LATENCY`, the offline compiler will automatically insert capacity-balancing FIFO buffers after this module when necessary. The safe value for `CAPACITY` is `1`.
`HAS_SIDE_EFFECTS`	Indicates whether the RTL module has side effects. Modules that have internal states or communicate with external memories are examples of modules with side effects. Set `HAS_SIDE_EFFECTS` to `"yes"` to indicate that the module has side effects. Specifying `HAS_SIDE_EFFECTS` to `"yes"` ensures that optimization efforts do not remove calls to modules with side effects. Stall-free modules with side effects (that is, `IS_STALL_FREE="yes"` and `HAS_SIDE_EFFECTS="yes"`) must properly handle `ivalid=0` input cases because the module might receive invalid data occasionally. The safe value for `HAS_SIDE_EFFECTS` is `"yes"`.
`ALLOW_MERGING`	Instructs the offline compiler to merge multiple instances of the RTL module.

continued...

XML Element	Description
	Set `ALLOW_MERGING` to `"yes"` to allow merging of multiple instances of the module. Intel recommends setting `ALLOW_MERGING` to `"yes"`. The safe value for `ALLOW_MERGING` is `"no"`. *Note:* Marking the module with `HAS_SIDE_EFFECTS="yes"` does not prevent merging.

2.1.1.4.2 XML Elements for INTERFACE

In the XML specification file of the RTL module within an Intel FPGA SDK for OpenCL library, there are XML elements under `INTERFACE` that you can define to specify aspects of the RTL module's interface (for example, Avalon-ST interface).

Table 7. **Mandatory XML Elements Associated with the INTERFACE Eement in the XML Specification File of an RTL Module**

XML Element	Description

`INPUT`	Specifies the input parameter of the RTL module.
	`INPUT` attributes:
	• `port`—Specifies the port name of the RTL module.
	• `width`—Specifies the width of the port in bits.
	  AOCL only supports widths that correspond to OpenCL data types (that is, 8 (`uchar`), 16, 32, 64, 128, 256, 512, and 1024 bits (`long16`)).
	  Note: Size of a *type*3 vector is 4 x sizeof(*type*), giving the impression that valid sizes of 24, 48, 96, and 192 bits are unsupported.
	The input parameters are concatenated to form the input stream.
	Aggregate data structures such as structs and arrays are not supported as input parameters.

`OUTPUT`	Specifies the output parameter of the RTL module.
	`OUTPUT` attributes:
	• `port`—Specifies the port name of the RTL module.
	• `width`—Specifies the width of the port in bits.
	  The SDK only supports widths that correspond to OpenCL data types (that is, 8 (`uchar`), 16, 32, 64, 128, 256, 512, and 1024 bits (`long16`)).
	  Note: Size of *type*3 vector is 4 x sizeof(*type*), giving the impression that valid sizes of 24, 48, 96, and 192 bits are unsupported.
	The return value from the input stream is sent out via the output parameter on the output stream.
	Aggregate data structures such as structs and arrays are not supported as input parameters.

If your RTL module communicates with external memory, you need to include additional XML elements:

```
<MEM_INPUT port="m_input_A" access="readonly"/>
<MEM_INPUT port="m_input_sum" access ="readwrite"/>
<AVALON_MEM port="avm_port0" width="512" burstwidth="5" optype="read"
buffer_location=""/>
```

Table 8.   Additional XML Elements to Support External Memory Access

XML Element	Description
MEM_INPUT	Describes a pointer input to the RTL module. MEM_INPUT attributes: • port—Specifies the name of the pointer input. • access—Specifies to the Intel FPGA SDK for OpenCL Offline Compiler how the RTL module will use this pointer. Valid access values are readonly and readwrite. If the RTL module only writes with this pointer, assign readwrite to access. Because all pointers to external memory must be 64 bits, there is no width attribute associated with MEM_INPUT.
AVALON_MEM	Declares the Avalon-MM interface for your RTL module. AVALON_MEM attributes: • port—Specifies the root of the corresponding port names in the RTL module. For example, if port has a value of avm_port0_, the names of all Avalon-MM interface ports for the RTL module will start with avm_port0_. • width—Specifies the data width, which must match the corresponding width value in the accelerator board's board_spec.xml file. Within the board_spec.xml file, the width value is specified in the interface element under global_mem. For more information, refer to the *global_mem* section under *XML Elements, Attributes, and Parameters in the board_spec.xml File* in the *Intel FPGA SDK for OpenCL Custom Platform Toolkit User Guide* . • burstwidth—Specifies the number of bits required to represent burst size. Use burstwidth = log(maxburst) +1 to calculate the burst size, where maxburst is the corresponding maximum burst size specified in the board_spec.xml file. For example, if maxburst=16, burstwidth=5. • optype—Specifies either the Avalon-MM port is reading (read) or writing (write) from external memory. You can only assign either read or write to optype. • buffer_location—Supports heterogeneous memory. Leave this attribute blank because the heterogeneous memory compilation flow is currently untested.

For the AVALON_MEM element defined in the code example above, the corresponding RTL module ports are as follows:

```
output         avm_port0_enable,
input  [511:0] avm_port0_readdata,
input          avm_port0_readdatavalid,
input          avm_port0_waitrequest,
output [31:0]  avm_port0_address,
output         avm_port0_read,
output         avm_port0_write,
input          avm_port0_writeack,
output [511:0] avm_port0_writedata,
output [63:0]  avm_port0_byteenable,
output [4:0]   avm_port0_burstcount,
```

There is no assumed correspondence between pointers that you specify with MEM_INPUT and the Avalon-MM interfaces that you specify with AVALON_MEM. An RTL module can use a single pointer to address zero to multiple Avalon-MM interfaces.

Related Links

XML Elements, Attributes, and Parameters in the board_spec.xml File: global_mem

2.1.1.5 Interaction between RTL Module and External Memory

Allow your RTL module to interact with external memory only if the interaction is necessary and unavoidable.

Important: The preferred method for having your RTL module interact with external memory is to have the OpenCL kernel access global memory and then feed that memory content to the RTL module. For operations like reading from and writing to external memory on every kernel invocation, instruct the OpenCL kernel to perform the operation. To do so, you can create an OpenCL helper function for the OpenCL kernel in the same Intel FPGA SDK for OpenCL library as the RTL module.

The following examples demonstrate how to structure code in an RTL module for easy integration into an OpenCL library:

Table 9. **Example Code in an RTL Module that Interacts with External Memory**

Complex RTL Module	Simplified RTL Module

```// my_rtl_fn does:```
```// out_ptr[idx] = fn(in_ptr[idx])```
```my_rtl_fn (in_ptr, out_ptr,idx);```
	```int in_value = in_ptr[idx];```
	```// my_rtl_fn now does: out = fn(in)```
	```int out_value = my_rtl_fn (in_value);```
	```out_ptr[idx] = out_value;```

The complex RTL module on the left reads a value from external memory, performs a scalar function on the value, and then writes the value back to global memory. Such an RTL module is difficult to describe when you integrate it into an OpenCL library. In addition, this RTL module is harder to verify and causes very conservative pointer analysis in the Intel FPGA SDK for OpenCL Offline Compiler.

The simplified RTL module on the right provides the same overall functionality as the complex RTL module. However, the simplified RTL module only performs a scalar-to-scalar calculation without connecting to global memory. Integrating this simplified RTL module into the OpenCL library makes it much easier for the offline compiler to analyze the resulting OpenCL kernel.

There are times when an RTL module requires an Avalon-MM port to communicate with external memory. This Avalon-MM port connects to the same arbitration network to which all other global load and store units in the OpenCL kernels connect.

If an RTL module receives a memory pointer as an argument, the offline compiler enforces the following memory model:

- If an RTL module writes to a pointer, nothing else in the OpenCL kernel can read from or write to this pointer.

- If an RTL module reads from a pointer, the rest of the OpenCL kernel and other RTL modules may also read from this pointer.

- You may set the access field of the MEM_INPUT attribute to specify how the RTL module uses the memory pointer. Ensure that you set the value for access correctly because there is no way to verify the value.

2.1.1.6 Order of Threads Entering an RTL Module

Do not assume that threads entering an RTL module follow a defined order. In addition, an RTL module can reorder threads. As a result, thread 0 does not necessarily enter the module before thread 1.

2.1.1.7 OpenCL C Model of an RTL Module

Each RTL module within an OpenCL library must have an OpenCL C model. The OpenCL C model verifies the overall OpenCL system during emulation.

Example OpenCL C model file for a square root function:

```
double my_sqrtfd (double a)
{
    return sqrt(a);
}
```

Intel recommends that you emulate your OpenCL system. If you decide not to emulate your OpenCL system, create an empty function with a name that matches the function name you specified in the XML specification file.

Related Links

XML Syntax of an RTL Module on page 125
    This section provides the syntax of a simple XML specification file for an RTL module that implements double-precision square root function.

2.1.1.8 Potential Incompatibility between RTL Modules and Partial Reconfiguration

When creating an OpenCL library using RTL modules, you might encounter Partial Reconfiguration-related issues.

Consider a situation where you create and verify your library on a device that does not support Partial Reconfiguration (PR). If a library user then uses the library's RTL module inside a PR region, the module might not function correctly after PR.

To ensure that the RTL modules function correctly on a device that uses PR:

- The RTL modules do not use memory logic array blocks (MLABs) with initialized content.

- The RTL modules do not make any assumptions regarding the power-up values of any logic.

For complete PR coding guidelines, refer to Creating a Partial Reconfiguration Design in volume 1 of *Quartus Prime Pro Edition Handbook*.

2.1.2 Packaging an OpenCL Helper Function File for an OpenCL Library

Before creating an OpenCL library file, package each OpenCL source file with helper functions into a `.aoco` file. Unlike RTL modules, you do not need to create an XML specification file.

In general, you do not need to create a library to share helper functions written in OpenCL. You can distribute a helper function in source form (for example, *<shared_file>*`.cl`) and then insert the line `#include "`*<shared_file>*`.cl"` in the OpenCL kernel source code.

Consider creating a library under the following circumstances:

- The helper functions are in multiple files and you want to simplify distribution.

- You do not want to expose the helper functions' source code.

  The helper functions are stored as LLVM IR, an assembly-like language, without comments inside the associated library.

Hardware generation is not necessary for the creation of a `.aoco` file. Compile the OpenCL source file using the `-c` offline compiler command option.

Note: A library can only include OpenCL helper functions. The Intel FPGA SDK for OpenCL Offline Compiler will issue an error message if the library contains OpenCL kernels.

- To package an OpenCL source file into a `.aoco` file, invoke the following command: `aoc -c -shared <OpenCL_source_file_name>.cl -o <OpenCL_object_file_name>.aoco`

  where the `-shared` offline compiler command option instructs the compiler to create a `.aoco` file that is suitable for inclusion into an OpenCL library.

Related Links

- Packaging Multiple Object Files into a Library File on page 135
     After creating the `.aoco` files that you want to include into an OpenCL library, package them into a library file by invoking the Intel FPGA SDK for OpenCL `library` utility command option.

- Specifying an OpenCL Library when Compiling an OpenCL Kernel on page 135
     To use an OpenCL library in an OpenCL kernel, specify the library file name and directory when you compile the kernel.

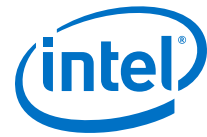## 2.1.3 Packaging an RTL Component for an OpenCL Library

Before creating an OpenCL library file, package each RTL component into a `.aoco` file.

Hardware generation is not necessary for the creation of a `.aoco` file. Compile the OpenCL source file using the `-c` Intel FPGA SDK for OpenCL Offline Compiler command option.

- To package an RTL component into a `.aoco` file, invoke the following command: `aoc -c <RTL component description file name>.xml -o <RTL object file name>.aoco`

Related Links

- Packaging Multiple Object Files into a Library File on page 135
     After creating the `.aoco` files that you want to include into an OpenCL library, package them into a library file by invoking the Intel FPGA SDK for OpenCL `library` utility command option.

- Verifying the RTL Modules on page 135
     The creator of an OpenCL library is responsible for verifying the RTL modules within the library, both as stand-alone entities and as part of an OpenCL system.

- Specifying an OpenCL Library when Compiling an OpenCL Kernel on page 135

To use an OpenCL library in an OpenCL kernel, specify the library file name and directory when you compile the kernel.

2.1.3.1 Restrictions and Limitations in RTL Support for the Intel FPGA SDK for OpenCL Library Feature

The Intel FPGA SDK for OpenCL supports the use of RTL modules in an OpenCL library with some restrictions and limitations.

When creating your RTL module, ensure that it operates within the following restrictions:

- An RTL module must use a single input Avalon-ST interface. That is, a single pair of `ready` and `valid` logic must control all the inputs.

  You have the option to provide the necessary Avalon-ST ports but declare the RTL module as stall-free. In this case, you do not have to implement proper stall behavior because the Intel FPGA SDK for OpenCL Offline Compiler creates a wrapper for your module. Refer to *XML Syntax of an RTL Module* and *Using an OpenCL Library that Works with Simple Functions (Example 1)* for more syntax and usage information, respectively.

  Note: You must handle `ivalid` signals properly if your RTL module has an internal state. Refer to *Stall-Free RTL* for more information.

- The RTL module must work correctly with exactly one clock, regardless of clock frequency.

- Data input and output sizes must match valid OpenCL data types, from 8 bits for `char` to 1024 bits for `long16`.

  For example, if you work with 24-bit values inside an RTL module, declare inputs to be 32 bits and declare function signature in the SDK's library header file to accept the `uint` data type. Then, inside the RTL module, accept the 32-bit input but discard the top 8 bits.

- RTL modules cannot connect to external I/O signals. All input and output signals must come from an OpenCL kernel.

- An RTL module must have a `clock` port, a `resetn` port, and Avalon-ST input and output ports (that is, `ivalid`, `ovalid`, `iready`, `oready`). Name the ports as specified here.

- RTL modules that communicate with external memory must have Avalon Memory-Mapped (Avalon-MM) port parameters that match the corresponding Custom Platform parameters. The offline compiler does not perform any width or burst adaptation.

- RTL modules that communicate with external memory must behave as follows:

  — They cannot burst across the burst boundary.

  — They cannot make requests every clock cycle and stall the hardware by monopolizing the arbitration logic. An RTL module must pause its requests regularly to allow other load or store units to execute their operations.

- RTL modules cannot act as stand-alone OpenCL kernels. RTL modules can only be helper functions and be integrated into an OpenCL kernel during kernel compilation.

- Every function call that corresponds to RTL module instantiation is completely independent of other instantiations. There is no hardware sharing.
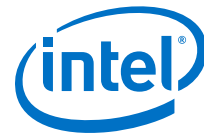
- Do not incorporate kernel code (that is, functions marked as `kernel`) into a `.aoclib` library file. Incorporating kernel code into the library file causes the offline compiler to issue an error message. You may incorporate helper functions into the library file.

- An RTL component must receive all its inputs at the same time. A single `ivalid` input signifies that all inputs contain valid data.

- The SDK does not support I/O RTL modules.

- You can only set RTL module parameters in the `<RTL module description file name>.xml` specification file, not the OpenCL kernel source file. To use the same RTL module with multiple parameters, create a separate `FUNCTION` tag for each parameter combination.

Currently, the SDK's RTL module support for the library feature has the following limitations:

- You can only pass data inputs to an RTL module by value via the OpenCL kernel code. Do not pass data inputs to an RTL module via pass-by reference, structs, or channels. In the case of channel data, extract the data from the channel first and then pass the extracted the scalar data to the RTL module.

  Note: Passing data inputs to an RTL module via pass-by reference or structs will cause a fatal error to occur in the offline compiler.

- The debugger (for example, GDB for Linux) cannot step into a library function during emulation. In addition, optimization and area reports will not include code line numbers beside the library functions.

- Names of RTL module source files cannot conflict with the file names of Intel FPGA SDK for OpenCL Offline Compiler IP. Both the RTL module source files and the offline compiler IP files are stored in the `<kernel file name>/system/synthesis/submodules` directory. Naming conflicts will cause existing offline compiler IP files in the directory to be overwritten by the RTL module source files.

- The SDK does not support `.qip` files. You must manually parse nested `.qip` files to create a flat list of RTL files.

- It is very difficult to debug an RTL module that works correctly on its own but works incorrectly as part of an OpenCL kernel. Double check all parameters under the `ATTRIBUTES` element in the `<RTL module description file name>.xml` file.

- All offline compiler area estimation tools assume that RTL module area is 0. The SDK does not currently support the capability of specifying an area model for RTL modules.

- RTL modules cannot access a 2x clock that is in-phase with the kernel clock and at twice the kernel clock frequency.

Related Links

- XML Syntax of an RTL Module on page 125
  This section provides the syntax of a simple XML specification file for an RTL module that implements double-precision square root function.

- Using an OpenCL Library that Works with Simple Functions (Example 1) on page 136
  Intel provides an OpenCL library design example of a simple kernel that uses a library containing RTL implementations of three double-precision functions: `sqrt`, `rsqrt`, and `divide`.

- Stall-Free RTL on page 123
  The Intel FPGA SDK for OpenCL Offline Compiler can optimize hardware
  resource usage and performance by removing stall logic around an RTL module
  with fixed latency.

2.1.4 Verifying the RTL Modules

The creator of an OpenCL library is responsible for verifying the RTL modules within
the library, both as stand-alone entities and as part of an OpenCL system.

1. Verify each RTL module using standard hardware verification methods.

2. Modify one of Intel FPGA SDK for OpenCL library design examples to test your RTL
   modules inside the overall OpenCL system.

   This testing step is critical to prevent library users from encountering hardware
   problems.

   It is crucial that you set the values for the `ATTRIBUTES` elements in the XML
   specification file correctly. Because you cannot simulate the entire OpenCL system,
   you will likely not discover problems caused by interface-level errors until
   hardware runs.

3. *Note:* The Intel FPGA SDK for OpenCL `library` utility performs consistency
   checks on the XML specification file and source files, with some limitations.

   Invoke the `aocl library [<command option>]` command.

   - For a list of supported *<command options>*, invoke the `aocl library`
     command.

   - The `library` utility does not detect errors in values assigned to elements
     within the `ATTRIBUTES`, `MEM_INPUT`, and `AVALON_MEM` elements in the XML
     specification file.

   - The `library` utility does not detect RTL syntax errors. You must check the
     `<your_kernel_filename>/quartus_sh_compile.log` file for RTL syntax
     errors. However, parsing the errors might be time consuming.

2.1.5 Packaging Multiple Object Files into a Library File

After creating the `.aoco` files that you want to include into an OpenCL library,
package them into a library file by invoking the Intel FPGA SDK for OpenCL `library`
utility command option.

- To package multiple object files into a single library file, invoke the following
  command: `aocl library create -o <library file name>.aoclib`
  `<object file 1>.aoco [<object file 2>.aoco ... <object file`
  `N>.aoco]`

  The `aocl library` utility command creates a `<library file name>.aoclib`
  library file, which includes the `.aoco` object files you specify in the command. A
  library file may contain both RTL-based object files and OpenCL-based object files.

2.1.6 Specifying an OpenCL Library when Compiling an OpenCL Kernel

To use an OpenCL library in an OpenCL kernel, specify the library file name and
directory when you compile the kernel.

Important: Using a library does not reduce kernel compilation time.

- To specify an OpenCL library to the Intel FPGA SDK for OpenCL Offline Compiler, invoke the following command: `aoc -l <library_file_name>.aoclib [-L <library directory>] <kernel file name>.cl`

    where the `-l <library_file_name>.aoclib` command option specifies the library file name, and the `-L <library directory>` command option specifies the directory containing the library files.

    You may include multiple instances of `-l <library file name>` and `-L <library directory>` in the offline compiler command.

For example, if you create a library that includes the functions `my_div_fd()`, `my_sqrtfd()`, and `myrsqrtfd()`, the OpenCL kernel code might resemble the following:

```
#include "lib_header.h"

kernel void test_lib (
    global double * restrict in,
    global double * restrict out,
    int N) {
        int i = get_global_id(0);
        for (int k =0; k < N; k++) {
            double x = in[i*N + k];
            out[i*N + k] = my_divfd
                (my_rsqrtfd(x),
                my_sqrtfd(my_rsqrtfd (x)));
        }
}
```

*Note:*   Library-related lines are highlighted in bold.

The corresponding `lib_header.h` file might resemble the following:

```
double my_sqrtfd (double x);
double my_rsqrtfd(double x);
double my_divfd(double a, double b);
```

2.1.7 Using an OpenCL Library that Works with Simple Functions (Example 1)

Intel provides an OpenCL library design example of a simple kernel that uses a library containing RTL implementations of three double-precision functions: `sqrt`, `rsqrt`, and `divide`.

The `example1.tgz` tar ball includes a library, a kernel, and a host system. The `example1.cl` kernel source file includes two kernels. The kernel `test_lib` uses library functions; the kernel `test_builtin` uses built-in functions. The host runs both kernels and then compares their outputs and runtimes. Intel recommends that you use the same strategy to verify your own library functions.

To compile this design example, perform the following tasks:

1. Obtain `example1.tgz` from the OpenCL Design Examples web page and store it in a directory that you own.

2. At the Intel FPGA SDK for OpenCL command prompt, navigate to the location of the design example.

3. Type `perl make_lib.pl` to create the `lib1/double_lib.aoclib` library file.

4. Type `aoc -l double_lib.aoclib -L lib1 -I lib1 example1.cl` to compile the OpenCL kernel while including the `double_lib.aoclib` library file.

5. Type `cd host` to navigate to the host directory.

6. Type `gmake -f Makefile` to build the host program.

7. Type `host/example1` to run the host program.
   The SDK generates the following messages after the host program runs successfully:

```
Loading example1.aocx ...
Create buffers
Generate random data for conversion...
Enqueueing both library and builtin in kernels 4 times with global size 65536
Kernel computation using library function took 5.35333 seconds
Kernel computation using built-in function took 5.39949 seconds
Reading results to buffers...
Checking results...
Library function throughput is within 5% of builtin throughput.
PASSED
```

Related Links

OpenCL Design Examples page

2.1.8 Using an OpenCL Library that Works with External Memory (Example 2)

Intel provides an OpenCL library design example of a simple kernel that uses a library containing two RTL modules that communicate with global memory.

The `example1.tgz` tar ball includes a library, a kernel, and a host system. In this example, the RTL code that communicates with global memory is Custom Platform- or Reference Platform-dependent. Ensure that the compilation targets the board that corresponds to the Stratix V Network Reference Platform.

Intel generated the RTL modules `copyElement()` and `sumOfElements()` using the Intel FPGA SDK for OpenCL Offline Compiler, which explains the extra inputs in the code.

The `example2.cl` kernel source file includes two kernels. The kernel `test6` is an NDRange kernel that calls the `copyElement()` RTL function, which copies data from `B[]` to `A[]` and then stores `global_id+100` in `C[]`. The kernel `test11` is a single work-item kernel that uses an RTL function . The `sumOfElements()` RTL function determines the sum of the elements of `A[]` in range `[i, N]` and then adds the rest to `C[i]`.

Note: First invocations of `sumOfElements(i=0)` will take more time to execute than later invocations.

To compile this design example, perform the following tasks:

1. Obtain `example2.tgz` from the OpenCL Design Examples web page and store it in a directory that you own.

2. At the Intel FPGA SDK for OpenCL command prompt, navigate to the location of the design example.

3. Type `perl make_lib.pl` to create the `lib1/mem_users.aoclib` library file.

4. Type `aoc -l lib/mem_users.aoclib -I lib example2.cl` to compile the OpenCL kernel while including the `mem_users.aoclib` library file.

5. Type `cd host` to navigate to the host directory.

6. Type `gmake -f Makefile` to build the host program.

7. Type `host/example2` to run the host program.
   The SDK generates the following messages after the host program runs successfully:

```
Loading example2.aocx ...
Running test6
Launching the kernel test6 with globalsize=128 localSize=16
Loading example2.aocx ...
Running test11
Launching the kernel test11 with globalsize=1 localSize=1
PASSED
```

Related Links

- OpenCL Design Examples page

- Compiling a Kernel for a Specific FPGA Board (--board <board_name>) on page 97

    To compile your OpenCL kernel for a specific FPGA board, include the `--board <board_name>` option in the `aoc` command.

- Intel FPGA SDK for OpenCL Stratix V Network Reference Platform Porting Guide

2.1.9 OpenCL Library Command-Line Options

Both the Intel FPGA SDK for OpenCL Offline Compiler's set of commands and the SDK utility include options you can invoke to perform OpenCL library-related tasks.

Table 10.   Library-Related Intel FPGA SDK for OpenCL Offline Compiler Command Options

Command Option	Description
`-shared`	In conjunction with the `-c` command option, compiles an OpenCL source file into an object file (`.aoco`) that you can then include into a library. `aoc -c -shared <OpenCL source file name>.cl -o <OpenCL object file name>.aoco`
`-I <library_directory>`	Adds *<library directory>* to the header file search path. `aocl -I <library_header_file_directory> -l <library_file_name>.aoclib <kernel_file_name>.cl`
`-L <library directory>`	Adds *<library directory>* to the OpenCL library search path. Space after "-L" is optional.

continued...

Command Option	Description
	`aoc -l <library_file_name>.aoclib [-L <library directory>] <kernel file name>.cl`
`-l <library_file_name>.aoclib`	Specifies the OpenCL library file (`<library_file_name>.aoclib`). Space after `-l` is optional. `aoc -l <library_file_name>.aoclib [-L <library directory>] <kernel file name>.cl`
`--library-debug`	Generates debug output that relates to libraries. Part of the additional output appears in stdout, the other part appears in the `<kernel_file_name>/<kernel_file_name>.log` file. `aoc -l <library_file_name>.aoclib --library-debug <kernel_file_name>.cl`

Table 11. **Intel FPGA SDK for OpenCL Library Utility (aocl library) Command Options**

Command Option	Description
`hdl-comp-pkg <XML_specification_ file>.xml`	Packages a single HDL component into a `.aoco` file that you then include into a library. Invoking this command option is similar to invoking `aoc -c <XML_specification_file>.xml`. However, the processing time is faster because the `aocl` utility will not perform any environment checks. `aocl library hdl-comp-pkg <XML_specification_ file>.xml -o <output_file>.aoco`
`-c <XML_specification_ file>.xml`	Same function as `hdl-comp-pkg <XML_specification_ file>.xml`. `aocl library -c <XML_specification_ file>.xml`
`create`	Creates a library file from the `.aoco` files that you created by invoking the `hdl-comp-pkg` utility option or the `aoc -shared` command, and any other `.aoclib` libraries. `aocl library create [-name <library_name>] [-vendor <library_vendor>] [-version <library_version>] [-o <output_file>.aoclib] [.aoco...] [.aoclib...]` where `-name`, `-vendor`, and `-version` are optional information strings you can specify and add to the library.
`list <library_name>`	Lists all the RTL components in the library. Currently, this option is not available for use to list OpenCL functions. `aocl library list <library_name>`
`help`	Prints the list of Intel FPGA SDK for OpenCL library utility options and their descriptions on screen. `aocl library help`

2.2 Kernel Attributes for Configuring Local Memory System

The Intel FPGA SDK for OpenCL includes kernel attributes that you can include in a kernel to customize the geometry of the local memory system.

Attention: Only apply these local memory kernel attributes to local variables.

**Table 12.** **OpenCL Kernel Attributes for Configuring Local Memory**

Kernel Attribute	Description
register	Specifies that the local variable must be implemented in a register.
memory	Specifies that the local variable must be implemented in a memory system. Including the memory kernel attribute is equivalent to declaring the local variable with the __local qualifier.
numbanks(N) N is an integer value.	Specifies that the memory system implementing the local variable must have N banks, where N is a power-of-2 integer value greater than zero.
bankwidth(N) N is an integer value.	Specifies that the memory system implementing the local variable must have banks that are N bytes wide, where N is a power-of-2 integer value greater than zero.
singlepump	Specifies that the memory system implementing the local variable must be single pumped.
doublepump	Specifies that the memory system implementing the local variable must be double pumped.
numreadports(N) N is an integer value.	Specifies that the memory system implementing the local variable must have N read ports, where N is an integer value greater than zero.
numwriteports(N) N is an integer value.	Specifies that the memory system implementing the local variable must have N read ports, where N is an integer value greater than zero.
merge("label", "direction")	Forces two or more variables to be implemented in the same memory system. *label* is an arbitrary string. Assign the same label to all variables that you wish to merge. Specify *direction* as either width or depth to identify whether the memories should be merged width-wise or depth-wise, respectively.
bank_bits(b_0, b_1, ..., b_n)	Forces the memory system to split into 2^n banks, with {b_0, b_1, ..., b_n} forming the bank-select bits. *Important:* b_0, b_1, ..., b_n must be consecutive, positive integers. If you specify the numbanks(n) attribute without the bank_bits attribute, the bank-select bits default to the least significant bits (that is, 0, 1, ..., log2(numbanks)-1).

Table 13. **Code Examples for Local Memory Kernel Attributes**

Example Use Case	Syntax
Implements a local variable in a register	```int __attribute__((register)) a[12];```
Implements a local memory system with eight banks, each with a width of 8 bytes	```int __attribute__((memory,
numbanks(8),
bankwidth(8)) b[16];```
Implements a double-pumped local memory system with one 128-byte wide bank, one write port, and four read ports	```int __attribute__((memory,
numbanks(1),
bankwidth(128),
doublepump,
numwriteports(1),
numreadports(4)) c[32];```

Related Links

- Improve Kernel Performance by Banking the Local Memory

- Optimize Accesses to Local Memory by Controlling the Memory Replication Factor

2.2.1 Restrictions on the Usage of Local Variable-Specific Kernel Attributes

The Intel FPGA SDK for OpenCL Offline Compiler will error out or issue warnings if it detects unsupported usages of the local variable-specific kernel attributes or incorrect memory configurations.

Unsupported usages of local variable-specific kernel attributes that cause compilation errors:

- You use the kernel attributes in declarations other than local variable declarations (for example, declarations for function parameters, global variable declarations, or function declarations).

- You use the `register` attribute in conjunction with any of the other local variable-specific kernel attributes.

- You specify the `numbanks` kernel attribute but not the `bankwidth` kernel attribute in the same local variable declaration, and vice versa.

- You include both the `singlepump` and `doublepump` kernel attributes in the same local variable declaration.

- You specify the `numreadports` and `numwriteports` kernel attributes without also including the `singlepump` or `doublepump` kernel attribute in the same local variable declaration.

- You specify the `numreadports` kernel attribute but not the `numwriteports` kernel attribute in the same local variable declaration, and vice versa.

- You specify any of the following kernel attributes without also specifying the `numbanks` and `bankwidth` kernel attributes in the same local variable declaration:
  - `numreadports`
  - `numwriteports`
  - `singlepump`
  - `doublepump`

Incorrect memory configurations that cause the offline compiler to issue warnings during compilation:

- The memory configuration that is defined by the local variable-specific kernel attributes exceeds the available storage size (for example, specifying eight banks of local memory for an integer variable).

Incorrect memory configurations that cause compilation errors:

- The bank width is smaller than the data storage size (for example, bank width is 2 bytes for an array of 4-byte integers).

- You specify memory configurations for the local variables. However, because of compiler restrictions or coding style, the offline compiler implements the variables in the same memory instead of splitting the memory.

- You specify the `register` kernel attribute for a local variable. However, because of compiler restrictions or coding style, the offline compiler cannot implement the variable in a register.

2.3 Kernel Attributes for Reducing the Overhead on Hardware Usage

The Intel FPGA SDK for OpenCL includes kernel attributes that you can include in a single work-item kernel to reduce logic utilization and improve kernel performance.These kernel attributes enables the Intel FPGA SDK for OpenCL Offline Compiler to omit the generation of unnecessary hardware to increase efficiency.

2.3.1 Hardware for Kernel Interface

The Intel FPGA SDK for OpenCL Offline Compiler generates hardware around the kernel pipeline. For some OpenCL applications, these interface hardware components are not necessary.

Hardware around the kernel pipeline is necessary for functions such as the following:

- Dispatching IDs for work-items and work-groups

- Communicating with the host regarding kernel arguments and work-group sizes

Figure 15 on page 143 illustrates the hardware that the offline compiler generates when it compiles the following kernel:

```
__kernel void my_kernel(global int* arg)
{
    …
    int sum = 0;
    for(unsigned i = 0; i < n; i++)
    {
        if(sum < m) sum += val;
    }
    *arg = sum;
    …
}
```
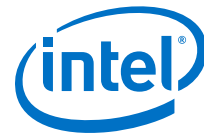
Figure 15. **Intel FPGA SDK for OpenCL Offline Compiler-Generated Interface Hardware around a Kernel Pipeline**

2.3.1.1 Omit Hardware that Generates and Dispatches Kernel IDs

The `max_global_work_dim(0)` kernel attribute instructs the Intel FPGA SDK for OpenCL Offline Compiler to omit logic that generates and dispatches global, local, and group IDs into the compiled kernel.

Semantically, the `max_global_work_dim(0)` kernel attribute specifies that the global work dimension of the kernel is zero. Setting this kernel attribute means that the kernel does not use any global, local, or group IDs. The presence of this attribute in the kernel code serves as a guarantee to the offline compiler that the kernel is a single work-item kernel.

When compiling the following kernel, the offline compiler will generate interface hardware as illustrated in Figure 16 on page 144.

```
channel int chan_in;
channel int chan_out;

__attribute__((max_global_work_dim(0)))
__kernel void plusK (int N, int k) {
    for (int i = 0; i < N; ++i) {
        int data_in = read_channel_intel(chan_in);
        write_channel_intel(chan_out, data_in + k);
    }
}
```

Figure 16. **Intel FPGA SDK for OpenCL Offline Compiler-Generated Interface Hardware for a Kernel with the max_global_work_dim(0) Attribute**

If your current kernel implementation has multiple work-items but does not use global, local, or group IDs, you can use the `max_global_work_dim(0)` kernel attribute if you modify the kernel code accordingly:

1. Wrap the kernel body in a `for` loop that iterates as many times as the number of work-items.

2. Launch the modified kernel with only one work-item.

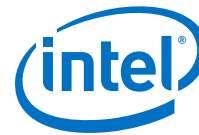## 2.3.1.2 Omit Communication Hardware between the Host and the Kernel

The `autorun` kernel attribute instructs the Intel FPGA SDK for OpenCL Offline Compiler to omit logic that is used for communication between the host and the kernel. A kernel that uses the `autorun` attribute starts executing automatically before any kernel that the host launches explicitly. In addition, this kernel restarts automatically as soon as it finishes its execution.

The `autorun` kernel attribute notifies the offline compiler that the kernel runs on its own and will not be enqueued by any host.

To leverage the `autorun` attribute, a kernel must meet all of the following criteria:

1. Does not use I/O channels

   Note: Kernel-to-kernel channels are supported.

2. Does not have any arguments

3. Has either the `max_global_work_dim(0)` attribute or the `reqd_work_group_size(X,Y,Z)` attribute

   Note: The parameters of the `reqd_work_group_size(X,Y,Z)` attribute must be divisors of 2^{32}.

As mentioned above, kernels with the `autorun` attribute cannot have any arguments and start executing without the host launching them explicitly. As a result, the offline compiler does not need to generate the logic for communication between the host and the kernel. Omitting this logic reduces logic utilization and allows the offline compiler to apply additional performance optimizations.

A typical use case for the `autorun` attribute is a kernel that reads data from one or more kernel-to-kernel channels, processes the data, and then writes the results to one or more channels. When compiling the kernel, the offline compiler will generate hardware as illustrated in .

```
channel int chan_in;
channel int chan_out;

__attribute__((max_global_work_dim(0)))
__attribute__((autorun))
__kernel void plusOne () {
    while(1) {
        int data_in = read_channel_intel(chan_in);
        write_channel_intel(chan_out, data_in + 1);
    }
}
```

Figure 17.   Single Work-Item Kernel with No Interface Hardware

2.4 Kernel Replication Using the num_compute_units(X,Y,Z) Attribute

You can replicate your single work-item OpenCL kernel by including the `num_compute_units(X,Y,Z)` kernel attribute.

As mentioned in *Specifying Number of Compute Units*, including the `num_compute_units(N)` kernel attribute in your kernel instructs the Intel FPGA SDK for OpenCL Offline Compiler to generate multiple compute units to process data. Since the offline compiler processes a single work-item kernel in one compute unit, the `num_compute_unit(N)` attribute instructs the offline compiler to generate *N* identical copies of the kernel in hardware.

Related Links

Specifying Number of Compute Units on page 24
    To increase the data-processing efficiency of an OpenCL kernel, you can instruct the Intel FPGA SDK for OpenCL Offline Compiler to generate multiple kernel compute units. Each compute unit is capable of executing multiple work-groups simultaneously.
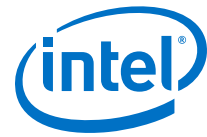
2.4.1 Customization of Replicated Kernels Using the get_compute_id() Function

To create compute units that are slightly different from one another but share a lot of common code, call the `get_compute_id()` intrinsic function in a kernel that also uses the `num_compute_units (X,Y,Z)` attribute.

Attention: You can only use the `get_compute_id()` function in a kernel that also uses the `autorun` and `max_global_work_dim(0)` kernel attributes.

Retrieving compute IDs is a convenient alternative to replicating your kernel in source code and then adding specialized code to each kernel copy. When a kernel uses the `num_compute_units(X,Y,Z)` attribute and calls the `get_compute_id()` function, the Intel FPGA SDK for OpenCL Offline Compiler assigns a unique compute ID to each compute unit. The `get_compute_id()` function then retrieves these unique compute IDs. You can use the compute ID to specify how the associated compute unit should behave differently from the other compute units that are derived from the same kernel source code. For example, you can use the return value of `get_compute_id()` to index into an array of channels to specify which channel each compute unit should read from or write to.

The `num_compute_units` attribute accepts up to three arguments (that is, `num_compute_units(X,Y,Z)`). In conjunction with the `get_compute_id()` function, this attribute allows you to create one-dimensional, two-dimensional, and three-dimensional logical arrays of compute units. An example use case of a 1D array of compute units is a linear pipeline of kernels (also called a daisy chain of kernels). An example use case of a 2D array of compute units is a systolic array of kernels.

Figure 18. **Schematic Diagram of a 4x4 Array of Compute Units**

The following example code specifies `num_compute_units(4,4)` in a single work-item kernel results in a 4x4 array that consists of 4 x 4 = 16 compute units.

```
__attribute__((max_global_work_dim(0)))
__attribute__((autorun))
__attribute__((num_compute_units(4,4)))
__kernel void PE() {

   row = get_compute_id(0);
   col = get_compute_id(1);


   …
}
```


For a 3D array of compute units, you can retrieve the X, Y, and Z coordinates of a compute unit in the logical compute unit array using `get_compute_id(0)`, `get_compute_id(1)`, and `get_compute_id(2)`, respectively. In this case, the API is very similar to the API of the work-item's intrinsic functions (that is, `get_global_id()`, `get_local_id()`, and `get_group_id()`).

Global IDs, local IDs, and group IDs can vary at runtime based on how the host invokes the kernel. However, compute IDs are known at compilation time, allowing the offline compiler to generate optimized hardware for each compute unit.

2.4.2 Using Channels with Kernel Copies

To implement channels within compute units (that is, replicated kernel copies), create an array of channels and then index into that array using the return value of `get_compute_id()`.

The example code below implements channels within multiple compute units.

```
#define N 4
channel int chain_channels[N+1];
```

```
__attribute__((max_global_work_dim(0)))
__kernel void reader(global int *data_in, int size) {
    for (i = 0; i < size; ++i) {
        write_channel_intel(chain_channels[0], data_in[i]);
    }
}

__attribute__((max_global_work_dim(0)))
__attribute__((autorun))
__attribute__((num_compute_units(N)))
__kernel void plusOne() {
    int compute_id = get_compute_id(0);
    int input = read_channel_intel(chain_channels[compute_id]);
    write_channel_intel(chain_channels[compute_id+1], input + 1);
}

__attribute__((max_global_work_dim(0)))
__kernel void writer(global int *data_out, int size) {
    for (i = 0; i < size; ++i) {
        data_out[i] = read_channel_intel(chain_channels[N]);;
    }
}
```

Figure 19.   Example Topology of Kernel Copies that Implement Channels

This figure illustrates the topology of the group of kernels that the OpenCL application code above generates.

Note: The implementation of kernel copies is functionally equivalent to defining four separate kernels in your source code and then hard-coding unique indexes for the accesses to `chain_channels[N]`.

2.5 Document Revision History

Table 14.   Document Revision History of the Advanced Features Chapter of the Intel FPGA SDK for OpenCL Programming Guide

Date	Version	Changes
May 2017	2017.05.08	• Updated Interaction between RTL Module and External Memory on page 130 to indicate preferred method for RTL module and external memory interactions. • Updated Potential Incompatibility between RTL Modules and Partial Reconfiguration on page 131 to include link to the partial reconfiguration guidelines in the *Quartus Prime Pro Edition Handbook*. • Added information about `bankbits` and `merge` kernel attributes to Kernel Attributes for Configuring Local Memory System on page 139. • Rebranded some functions in code examples as follows: — Rebranded `read_channel_altera` to `read_channel_intel`. — Rebranded `write_channel_altera` to `write_channel_intel`.
October 2016	2016.10.31	• Rebranded Altera SDK for OpenCL to Intel FPGA SDK for OpenCL. • Rebranded Altera Offline Compiler to Intel FPGA SDK for OpenCL Offline Compiler.
May 2016	2016.05.02	Initial release.

A Support Statuses of OpenCL Features

TheIntel FPGA SDK for OpenCL host runtime conforms with the OpenCL platform layer and application programming interface (API), with clarifications and exceptions.

A.1 Support Statuses of OpenCL 1.0 Features

The following sections outline the support statuses of the OpenCL features described in the *OpenCL Specification version 1.0*.

A.1.1 OpenCL1.0 C Programming Language Implementation

OpenCL is based on C99 with some limitations. Section 6 of the *OpenCL Specification version 1.0* describes the OpenCL C programming language. The Intel FPGA SDK for OpenCL conforms with the OpenCL C programming language with clarifications and exceptions. The table below summarizes the support statuses of the features in the OpenCL programming language implementation.

Attention: The support status "●" means that the feature is supported, and there might be a clarification for the supported feature in the Notes column. The support status "□" means that the feature is supported with exceptions identified in the Notes column. A feature that is not supported by the SDK is identified with an "X". OpenCL programming language implementations that are supported with no additional clarifications are not shown.

Section	Feature	Support Status	Notes
6.1.1	*Built-in Scalar Data Types*		
	double precision float	□	Preliminary support for all double precision float built-in scalar data type. This feature might not conform with the OpenCL Specification version 1.0.
			Currently, the following double precision floating-point functions conform with the OpenCL Specification version 1.0:
			add / subtract / multiply / divide / ceil / floor / rint / trunc / fabs / fmax / fmin / sqrt / rsqrt / exp / exp2 / exp10 / log / log2 / log10 / sin / cos / asin / acos / sinh / cosh / tanh / asinh / acosh / atanh / pow / pown / powr / tanh / atan / atan2 / ldexp / log1p / sincos

Section	Feature	Support Status	Notes
	half precision float	□	Support for scalar addition, subtraction and multiplication. Support for conversions to and from single-precision floating point. This feature might not conform with the OpenCL Specification version 1.0. This feature is supported in the Emulator.
6.1.2	*Built-in Vector Data Types*	□	Preliminary support for vectors with three elements. Three-element vector support is a supplement to the OpenCL Specification version 1.0.
6.1.3	*Built-in Data Types*	X	
6.1.4	*Reserved Data Types*	X	
6.1.5	*Alignment of Types*	●	All scalar and vector types are aligned as required (vectors with three elements are aligned as if they had four elements).
6.2.1	*Implicit Conversions*	●	Refer to Section 6.2.6: *Usual Arithmetic Conversions* in the *OpenCL Specification version 1.2* for an important clarification of implicit conversions between scalar and vector types.
6.2.2	*Explicit Casts*	●	The SDK allows scalar data casts to a vector with a different element type.
6.5	*Address Space Qualifiers*	□	Function scope`__constant` variables are not supported.
6.6	*Image Access Qualifiers*	X	
6.7	*Function Qualifiers*		
6.7.2	*Optional Attribute Qualifiers*	●	Refer to the *Intel FPGA SDK for OpenCL Best Practices Guide* for tips on using `reqd_work_group_size` to improve kernel performance. The SDK parses but ignores the `vec_type_hint` and `work_group_size_hint` attribute qualifiers.
6.9	*Preprocessor Directives and Macros*		
	`#pragma` directive: `#pragma unroll`	●	The Intel FPGA SDK for OpenCL Offline Compiler supports only `#pragma unroll`. You may assign an integer argument to the unroll directive to control the extent of loop unrolling. For example, `#pragma unroll 4` unrolls four iterations of a loop. By default, an unroll directive with no unroll factor causes the offline compiler to attempt to unroll the loop fully. Refer to the *Intel FPGA SDK for OpenCL Best Practices Guide* for tips on using `#pragma unroll` to improve kernel performance.
	`__ENDIAN_LITTLE__` defined to be value `1`	●	The target FPGA is little-endian.
	`__IMAGE_SUPPORT__`	X	`__IMAGE_SUPPORT__` is undefined; the SDK does not support images.
6.10	*Attribute Qualifiers*—The offline compiler parses attribute qualifiers as follows:		
6.10.2	*Specifying Attributes of Functions —Structure-type kernel arguments*	X	Convert structure arguments to a pointer to a structure in global memory.
6.10.3	*Specifying Attributes of Variables —`endian`*	X	
6.10.4	*Specifying Attributes of Blocks and Control-Flow-Statements*	X	

continued...

Sectio n	Feature	Support Status	Notes
6.10.5	*Extending Attribute Qualifiers*	•	The offline compiler can parse attributes on various syntactic structures. It reserves some attribute names for its own internal use. Refer to the *Intel FPGA SDK for OpenCL Best Practices Guide* for tips on how to optimize kernel performance using these kernel attributes.
6.11.2	*Math Functions*		
	built-in math functions	☐	Preliminary support for built-in math functions for double precision float. These functions might not conform with the OpenCL Specification version 1.0.
	built-in `half_` and `native_` math functions	☐	Preliminary support for built-in `half_` and `native_` math functions for double precision float. These functions might not conform with the OpenCL Specification version 1.0.
6.11.5	*Geometric Functions*	☐	Preliminary support for built-in geometric functions for double precision float. These functions might not conform with the OpenCL Specification version 1.0. Refer to *Argument Types for Built-in Geometric Functions* for a list of built-in geometric functions supported by the SDK.
6.11.8	*Image Read and Write Functions*	X	
6.11.9	*Synchronization Functions—the barrier synchronization function*	☐	Clarifications and exceptions: If a kernel specifies the `reqd_work_group_size` or `max_work_group_size` attribute, barrier supports the corresponding number of work-items. If neither attribute is specified, a barrier is instantiated with a default limit of 256 work-items. The work-item limit is the maximum supported work-group size for the kernel; this limit is enforced by the runtime.
6.11.11	*Async Copies from Global to Local Memory, Local to Global Memory, and Prefetch*	☐	The implementation is naive: Work-item (0,0,0) performs the copy and the `wait_group_events` is implemented as a barrier. If a kernel specifies the `reqd_work_group_size` or `max_work_group_size` attribute, `wait_group_events` supports the corresponding number of work-items. If neither attribute is specified, `wait_group_events` is instantiated with a default limit of 256 work-items.

Related Links

- Intel FPGA SDK for OpenCL Best Practices Guide

- Argument Types for Built-in Geometric Functions on page 152
  The Intel FPGA SDK for OpenCL supports scalar and vector argument built-in geometric functions with certain limitations.

A.1.2 OpenCL C Programming Language Restrictions

The Intel FPGA SDK for OpenCL conforms with the OpenCL Specification restrictions on specific programming language features, as described in section 6.8 of the *OpenCL Specification version 1.0*.

Warning: The Intel FPGA SDK for OpenCL Offline Compiler does not enforce restrictions on certain disallowed programming language features. Ensure that your kernel code does not contain features that the OpenCL Specification version 1.0 does not support.
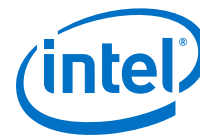
Feature	Support Status	Notes
pointer assignments between address spaces	•	Arguments to `__kernel` functions declared in a program that are pointers must be declared with the `__global`, `__constant`, or `__local` qualifier. The offline compiler enforces the OpenCL restriction against pointer assignments between address spaces.
pointers to functions	X	The offline compiler does not enforce this restriction.
structure-type kernel arguments	X	Convert structure arguments to a pointer to a structure in global memory.
images	X	The SDK does not support images.
bit fields	X	The offline compiler does not enforce this restriction.
variable length arrays and structures	X	
variable macros and functions	X	
C99 headers	X	
`extern`, `static`, `auto`, and `register` storage-class specifiers	X	The offline compiler does not enforce this restriction.
predefined identifiers	•	Use the `-D` option of the `aoc` command to provide preprocessor symbol definitions in your kernel code.
recursion	X	The offline compiler does not enforce this restriction.
irreducible control flow	X	The offline compiler does not enforce this restriction.
writes to memory of built-in types less than 32 bits in size	☐	Store operations less than 32 bits in size might result in lower memory performance.
declaration of arguments to `__kernel` functions of type `event_t`	X	The offline compiler does not enforce this restriction.
elements of a `struct` or a `union` belonging to different address spaces	X	The offline compiler does not enforce this restriction. ***Warning:*** Assigning elements of a `struct` or a `union` to different address spaces might cause a fatal error.

A.1.3 Argument Types for Built-in Geometric Functions

The Intel FPGA SDK for OpenCL supports scalar and vector argument built-in geometric functions with certain limitations.

Function	Argument Type	
	float	**double**
`cross`	•	•
`dot`		•
`distance`		•
`length`		•
`normalize`		•
`fast_distance`		—
`fast_length`		—
`fast_normalize`		—

A.1.4 Numerical Compliance Implementation

Section 7 of the *OpenCL Specification version 1.0* describes features of the C99 and IEEE 754 standards that OpenCL-compliant devices must support. The Intel FPGA SDK for OpenCL operates on 32-bit and 64-bit floating-point values in IEEE Standard 754-2008 format, but not all floating-point operators have been implemented.

The table below summarizes the implementation statuses of the floating-point operators:

Section	Feature	Support Status	Notes
7.1	*Rounding Modes*	□	Conversion between integer and single and half precision floating-point types support all rounding modes. Conversions between integer and double precision floating-point types support all rounding modes on a preliminary basis. This feature might not conform with the OpenCL Specification version 1.0.
7.2	*INF, NaN and Denormalized Numbers*	□	Infinity (INF) and Not a Number (NaN) results for single precision operations are generated in a manner that conforms with the OpenCL Specification version 1.0. Most operations that handle denormalized numbers are flushed prior to and after a floating-point operation. Preliminary support for double precision floating-point operation. This feature might not conform with the OpenCL Specification version 1.0.
7.3	*Floating-Point Exceptions*	X	
7.4	*Relative Error as ULPs*	□	Single precision floating-point operations conform with the numerical accuracy requirements for an embedded profile of the OpenCL Specification version 1.0. Preliminary support for double precision floating-point operation. This feature might not conform with the OpenCL Specification version 1.0.
7.5	*Edge Case Behavior*	●	

A.1.5 Image Addressing and Filtering Implementation

The Intel FPGA SDK for OpenCL does not support image addressing and filtering. The SDK does not support images.

A.1.6 Atomic Functions

Section 9 of the *OpenCL Specification version 1.0* describes a list of optional features that some OpenCL implementations might support. The Intel FPGA SDK for OpenCL supports atomic functions conditionally.

- Section 9.5: *Atomic Functions for 32-bit Integers*—The SDK supports all 32-bit global and local memory atomic functions. The SDK also supports 32-bit atomic functions described in Section 6.11.11 of the *OpenCL Specification version 1.1* and Section 6.12.11 of the *OpenCL Specification version 1.2*.

  — The SDK does not support 64-bit atomic functions described in Section 9.7 of the *OpenCL Specification version 1.0*.

Attention: The use of atomic functions might lower the performance of your design. The operating frequency of the hardware might decrease further if you implement more than one type of atomic functions (for example, `atomic_add` and `atomic_sub`) in the kernel.

A.1.7 Embedded Profile Implementation

Section 10 of the *OpenCL Specification version 1.0* describes the OpenCL embedded profile. The Intel FPGA SDK for OpenCL conforms with the OpenCL embedded profile with clarifications and exceptions.

The table below summarizes the clarifications and exceptions to the OpenCL embedded profile:

Clause	Feature	Support Status	Notes
1	64-bit integers	●	
2	3D images	X	The SDK does not support images.
3	Create 2D and 3D images with `image_channel_data_type` values	X	The SDK does not support images.
4	Samplers	X	
5	Rounding modes	●	The default rounding mode for `CL_DEVICE_SINGLE_FP_CONFIG` is `CL_FP_ROUND_TO_NEAREST`.
6	Restrictions listed for single precision basic floating-point operations	X	
7	half type	X	This clause of the OpenCL Specification version 1.0 does not apply to the SDK.
8	Error bounds listed for conversions from `CL_UNORM_INT8`, `CL_SNORM_INT8`, `CL_UNORM_INT16` and `CL_SNORM_INT16` to float	●	Refer to the table below for a list of allocation limits.

A.2 Support Statuses of OpenCL 1.2 Features

The following sections outline the support statuses of the OpenCL features described in the *OpenCL Specification version 1.2*.

A.2.1 OpenCL 1.2 Runtime Implementation

The Intel FPGA SDK for OpenCL supports the implementation of sub-buffer objects and image objects. For more information on sub-buffer objects and image objects, refer to sections 5.2 and 5.3 of the *OpenCL Specification version 1.2*, respectively.

The SDK also supports the implementation of the following APIs:

- `clSetMemObjectDestructorCallback`
- `clGetKernelArgInfo`
- `clSetEventCallback`

For more information on these APIs, refer to sections 5.4.1, 5.7.3, and 5.9 of the *OpenCL Specification 1.2*, respectively.

Related Links

OpenCL Specification version 1.2

A.2.2 OpenCL 1.2 C Programming Language Implementation

The Intel FPGA SDK for OpenCL supports a number of OpenCL C programming language features that are specified section 6 of the *OpenCL Specification version 1.2*. The SDK conforms with the OpenCL C programming language with clarifications and exceptions.

Attention: The support status "●" means that the feature is supported, and there might be a clarification for the supported feature in the Notes column. The support status "□" means that the feature is supported with exceptions identified in the Notes column.

Table 15. **Support Statuses of OpenCL 1.2 C Programming Language Features**

Section	Feature	Support Status	Notes
6.1.3	Other Built-in Data Types	●	Preliminary support. This feature might not conform with the OpenCL Specification version 1.0.
6.12.12	*Miscellaneous Vector Functions*	●	The SDK supports implementations of the following additional built-in vector functions: • `vec_step` • `shuffle` • `shuffle2`
6.12.13	*printf*	□	Preliminary support. This feature might not conform with the OpenCL Specification version 1.0. See below for details.
The `printf` function in OpenCL has syntax and features similar to the `printf` function in C99, with a few exceptions. For details, refer to the *OpenCL Specification version 1.2*.			
To use a `printf` function, there are no requirements for special compilation steps, buffers, or flags. You can compile kernels that include `printf` instructions with the usual `aoc` command.			
During kernel execution, `printf` data is stored in a global `printf` buffer that the Intel FPGA SDK for OpenCL Offline Compiler allocates automatically. The size of this buffer is 64 kB; the total size of data arguments to a `printf` call should not exceed this size. When kernel execution completes, the contents of the `printf` buffer are printed to standard output.			
Buffer overflows are handled seamlessly; `printf` instructions can be executed an unlimited number of times. However, if the `printf` buffer overflows, kernel pipeline execution stalls until the host reads the buffer and prints the buffer contents.			
Because `printf` functions store their data into a global memory buffer, the performance of your kernel will drop if it includes such functions.			
There are no usage limitations on `printf` functions. You can use `printf` instructions inside `if-then-else` statements, loops, etc. A kernel can contain multiple `printf` instructions executed by multiple work-items.			
Format string arguments and literal string arguments of `printf` calls are transferred to the host system from the FPGA using a special memory region. This memory region can overflow if the total size of the `printf` string arguments is large (3000 characters or less is usually safe in a typical OpenCL application). If there is an overflow, the error message `cannot parse auto-discovery string at byte offset 4096` is printed during host program execution.			
Output from `printf` is never intermixed, even though work-items may execute `printf` functions concurrently. However, the order of concurrent `printf` execution is not guaranteed. In other words, `printf` outputs might not appear in program order if the `printf` instructions are in concurrent datapaths.			

Related Links

OpenCL Specification version 1.2

A.3 Support Statuses of OpenCL 2.0 Features

The following sections outline the support statuses of the OpenCL features described in the *OpenCL Specification version 2.0*.

A.3.1 OpenCL 2.0 Headers

The Intel FPGA SDK for OpenCL provides both the OpenCL 1.0 and OpenCL 2.0 headers by the Khronos Group.

Attention: The SDK currently does not support all OpenCL 2.0 APIs. If you use the OpenCL 2.0 headers and make a call to an unsupported API, the call will return with an error code to indicate that the API is not fully supported.

A.3.2 OpenCL 2.0 Runtime Implementation

The Intel FPGA SDK for OpenCL offers preliminary support for shared virtual memory implementation, as described in section 5.6 of the *OpenCL Specification version 2.0*. For more information on shared virtual memory, refer to section 5.6 of the *OpenCL Specification version 2.0*.

Important: Refer to your board's specifications to verify that your board supports shared virtual memory.

Related Links

OpenCL Specification version 2.0 (API)

## A.3.3 OpenCL 2.0 C Programming Language Restrictions for Pipes

The Intel FPGA SDK for OpenCL offers preliminary support of OpenCL pipes.The following table lists the support statuses of pipe-specific OpenCL C programming language implementations, as described in the *OpenCL Specification version 2.0*

Attention: The support status "●" means that the feature is supported. There might be a clarification for the supported feature in the Notes column. A feature that is not supported by the SDK is identified with an "X".

Table 16. **Support Statuses of Built-in Pipe Read and Write Functions**

Details of the built-in pipe read and write functions are available in section 6.13.16.2 of the *OpenCL Specification version 2.0*.

Function	Support Status
`int read_pipe (pipe gentype p, gentype *ptr)`	●
`int write_pipe (pipe gentype p, const gentype *ptr)`	●
`int read_pipe (pipe gentype p, reserve_id_t reserve_id, uint index, gentype *ptr)`	X
`int write_pipe (pipe gentype p, reserve_id_t reserve_id, uint index, const gentype *ptr)`	X
`reserve_id_t reserve_read_pipe (pipe gentype p, uint num_packets)`	X

continued...

Function	Support Status
`reserve_id_t reserve_write_pipe (pipe gentype p, uint num_packets)`	
`void commit_read_pipe (pipe gentype p, reserve_id_t reserve_id)` `void commit_write_pipe (pipe gentype p, reserve_id_t reserve_id)`	X
`bool is_valid_reserve_id (reserve_id_t reserve_id)`	X

Table 17.   Support Statuses of Built-in Work-Group Pipe Read and Write Functions

Details of the built-in pipe read and write functions are available in section 6.13.16.3 of the *OpenCL Specification version 2.0*.

Function	Support Status
`reserve_id_t work_group_reserve_read_pipe (pipe gentype p, uint num_packets)` `reserve_id_t work_group_reserve_write_pipe (pipe gentype p, uint num_packets)`	X
`void work_group_commit_read_pipe (pipe gentype p, reserve_id_t reserve_id)` `void work_group_commit_write_pipe (pipe gentype p, reserve_id_t reserve_id)`	X

Table 18.   Support Statuses of Built-in Pipe Query Functions

Details of the built-in pipe query functions are available in section 6.13.16.4 of the *OpenCL Specification version 2.0*.

Function	Support Status
`uint get_pipe_num_packets (pipe gentype p)`	X
`uint get_pipe_max_packets (pipe gentype p)`	X

Related Links

OpenCL Specification version 2.0 (C Language)

A.4 Intel FPGA SDK for OpenCL Allocation Limits

Item	Limit
Maximum number of contexts	Limited only by host memory size
Minimum global memory allocation by runtime	The runtime allocates 64 kB of device memory when the context is created. This memory is reserved for program variables in global address space and for static variables inside functions. You may overwrite the amount of this memory, to as low as 0 bytes, by setting the environment variable *CL_CONTEXT_PROGRAM_VARIABLES_TOTAL_SIZE_ALTERA* to the desired value in bytes. If the OpenCL kernel uses the `printf` function, the runtime allocates an additional 64 kB of device memory.
Maximum number of queues	256 **Attention:** Each context uses two queues for system purposes.
Maximum number of program objects per context	20
Maximum number of even objects per context	Limited only by host memory size
Maximum number of dependencies between events within a context	1000

Item	Limit
Maximum number of event dependencies per command	20
Maximum number of concurrently running kernels	The total number of queues
Maximum number of enqueued kernels	1000
Maximum number of kernels per FPGA device	Hardware: no static limit Emulator: 256
Maximum number of arguments per kernel	128
Maximum total size of kernel arguments	256 bytes per kernel

A.5 Document Revision History

Table 19.   Document Revision History of the Intel FPGA SDK for OpenCL Programming Guide Appendix A: Support Statuses of OpenCL Features

Date	Document Version	Changes
May 2017	2017.05.05	• Maintenance release.
October 2016	2016.10.31	• Rebranded the Altera SDK for OpenCL to Intel FPGA SDK for OpenCL. • Rebranded the Altera Offline Compiler to Intel FPGA SDK for OpenCL Offline Compiler. • Modified information in the *Intel FPGA SDK for OpenCL Allocation Limits* section: — Updated information regarding minimum global memory allocation by runtime. — Updated the maximum number of queues from 70 to 256. — Updated the maximum number of kernels per FPGA device from 64 to no static limit when compiling to hardware and 256 when compiling to emulator. • In *OpenCL 1.0 C Programming Language Implementation*, under the Description column for half-precision float, added a note that this feature is supported in the Emulator. In addition, updated the support status of half-precision float from X to □. • Under *Support Statuses of OpenCL 2.0 Features*, added the topic *OpenCL 2.0 Headers* to explain that using the OpenCL 2.0 headers to call unsupported APIs will result in an error.
May 2016	2016.05.02	In *OpenCL 1.2 Runtime Implementation*, noted that AOCL supports the `clSetEventCallback`, `clGetKernelArgInfo`, and `clSetMemObjectDestructorCallback` APIs.
November 2015	2015.11.02	• Categorized feature support statuses and limitations based on OpenCL Specification versions. • Added the following functions to the list of OpenCL-conformant double precision floating-point functions: sinh / cosh / tanh / asinh / acosh / atanh / pow / pown / powr / tanh / atan / atan2 / ldexp / log1p / sincos • In *OpenCL 1.2 Runtime Implementation*, added sub-buffer object support. • In *OpenCL 2.0 Runtime Implementation*, added preliminary shared virtual memory support. • In *Altera SDK for OpenCL Allocation Limits*, added a minimum global memory allocation limit by the runtime.

Date	Document Version	Changes
May 2015	15.0.0	• Listed the double precision floating-point functions that the Altera SDK for OpenCL supports preliminarily. • Added *OpenCL C Programming Language Restrictions for Pipes*.
December 2014	14.1.0	• In *AOCL Allocation Limits*, updated the maximum number of kernels per FPGA device from 32 to 64.
June 2014	14.0.0	• Updated the following AOCL allocation limits: — Maximum number of contexts — Maximum number of queues — Maximum number of even objects per context
December 2013	13.1.1	• Modified support status designations in *Appendix: Support Statuses of OpenCL Features*. • Removed information on OpenCL programming language restrictions from the section *OpenCL Programming Language Implementation*, and presented the information in a new section titled *OpenCL Programming Language Restrictions*.
November 2013	13.1.0	Maintenance release.
June 2013	13.0 SP1.0	• Renamed *Optional Extensions* to *Atomic Functions*, and updated its content. • Removed *Platform Layer and Runtime Implementation*.
May 2013	13.0.1	Maintenance release.
May 2013	13.0.0	• Added updated information on allocation limits and OpenCL language support.
November 2012	12.1.0	Initial release.